

Scheduling → ពីរត្រូវបានដោឡូង ? ឬ

Main Points

- Scheduling policy: what to do next, when there are multiple threads ready to run
 - Or multiple packets to send, or web requests to serve, or ...
- Definitions
 - response time, throughput, predictability
- Uniprocessor policies
 - FIFO, round robin, optimal
 - multilevel feedback as approximation of optimal
- Multiprocessor policies → មិនសាន់
 - Affinity scheduling, gang scheduling
- Queueing theory → មិនចងកំណែ
 - Can you predict/improve a system's response time?

Example

processor = server + req = task, job
Client

- You manage a **web site**, that suddenly becomes wildly popular. Do you?

– Buy more hardware? เทพกรุ

อยู่ติด popular หน้า
ex วันลงทะเบียน เว็บ reg

– Implement a different scheduling policy?

→ต้องให้ software OS หาช่วง
จัดตารางเวลาเพื่อลด load

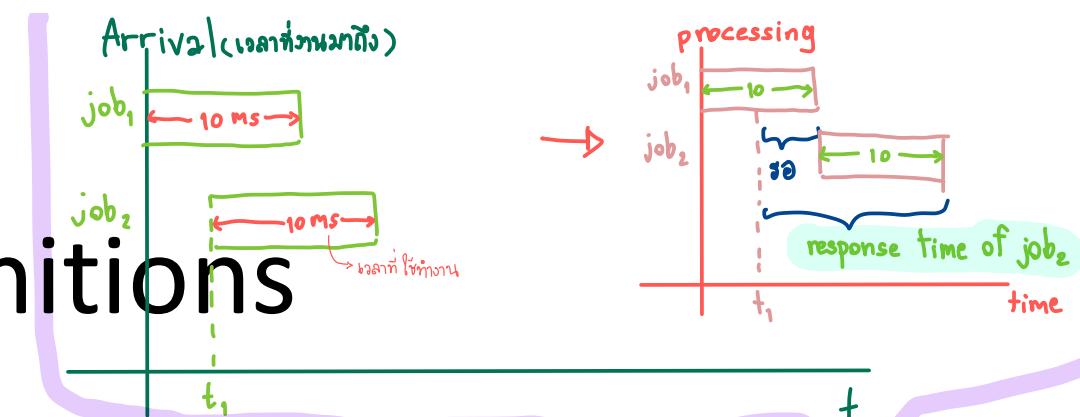
– Turn away some users? Which ones? เพิกเฉย
↑ หรือไม่ก็ปิด server หัวไปเลย

- How much worse will performance get if the web site becomes even more popular?

Definitions

- **Task/Job** งานที่ User ป้อนเข้าไปในระบบ
 - User request: e.g., mouse click, web request, shell command, ...
- **Latency/response time** เวลาที่งานใช้ประมวลผล
 - How long does a task take to complete? \rightarrow ต้องใช้เวลา
- **Throughput** \rightarrow จำนวนงานที่ทำได้ใน 1 unit of time
 - How many tasks can be done per unit of time? \rightarrow user นับได้เลย แต่ต้องทำ
- **Overhead** ผู้ดูแลระบบที่ต้องทำงานเพิ่มเติม
 - How much extra work is done by the scheduler?

\rightarrow overhead หมายความว่า CPU เสียเวลาคิดอยู่ \rightarrow Overhead
- **Fairness** จัดสรรเวลา, ทรัพยากรให้เท่าเทียม
 - How equal is the performance received by different users?
- **Predictability** ค.สามารถในการคำนวณ \rightarrow program นั้น ต้องมีผลลัพธ์ที่สม่ำเสมอ
 - How consistent is the performance over time?



More Definitions

- **Workload** กลุ่มของ job
 - Set of tasks for system to perform
 - **Preemptive scheduler** ก.ทำ scheduler ที่ scheduler ที่ไปปิดจั่งระบบได้
 - If we can take resources away from a running task
 - **Work-conserving** processor จะต้องทำงานตลอดเวลา ที่อยู่ฝั่งหน้า queue
 - Resource is used whenever there is a task to run
 - **Scheduling algorithm** algo ที่ใช้เป็นลักษณะ program ที่ใช้ workload เป็น input
 - takes a workload as input
 - decides which tasks to do first
 - Performance metric (throughput, latency) as output
 - Only preemptive, work-conserving schedulers to be considered
- throughput
 latency ↴
- ต้องดู check
 ณ. อยู่ส่วนใด
- เลือกเลื่อนถูกๆ
- ปรับปรุงผล
- เอาสูงสุด
- โดยใช้ Throughput,
 latency มาเป็นตัวเลือก
- เอาต่ำๆ
- ต้องดูงาน
 ให้ processor ผลิตผล
 (if have job)
- ซึ่งงาน job ที่ใช้งาน resource
 จะถูกมาเมื่อได้รับได้
- ห้ามหัก

ງາຍສຸດ

First In First Out (FIFO)

ເລືອດ job ໃປໃຈ່ນຕາມສໍາດັບຂາດ່ານ → algo ກ່າວ

- Schedule tasks in the order they arrive
 - Continue running them until they complete or give up the processor
- Example: memcached
 - Facebook cache of friend lists, ...
- On what workloads is FIFO particularly bad?

ຫຼຳກຳດົດ

↳ ໃນ performance ອົນທາງກຣນີ້ໄມ່ດີເທິ່ງ

Shortest Job First (SJF)

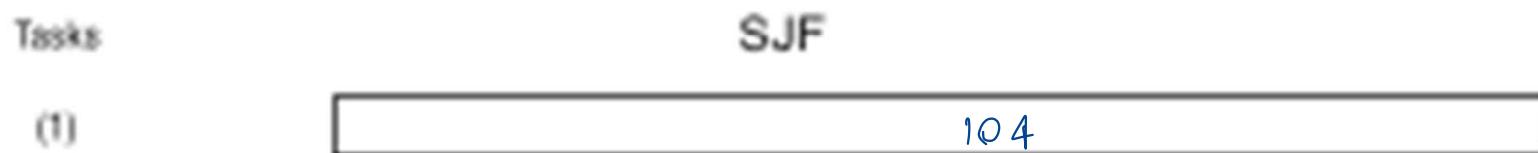
เป็นแบบ first in first out // ต้องการผลิตคิวไว้

- Always do the task that has the shortest remaining amount of work to do
 - Often called Shortest Remaining Time First (SRTF)
// เผื่องที่ใช้เวลาหาน้อยจะผลิตได้
- Suppose we have five tasks arrive one right after each other, but the first one is much longer than the others
 - Which completes first in FIFO? Next?
 - Which completes first in SJF? Next?

FIFO vs. SJF

เลือดเวลาร้อนๆมาก

ทำมันก็ง่ายๆมั้ย?



จัดเรียง
ลำดับ

Time

throughput แรกๆ = 2,3,4,5

$\sqrt{0.01} \text{ FIFO} = \text{SJF}$

ตัวต่อตัวจะดี

throughput เท่ากัน

Question

- Claim: SJF is optimal for average response time
 - Why? ត្រូវបានរួមចំណាំបន្ថែម

- Does SJF have any downsides?
 - ការដោះស្រាយពេលវេលាបានជាការងារបានបញ្ចប់
 - មិនអាចបង្កើតការងារបានបញ្ចប់

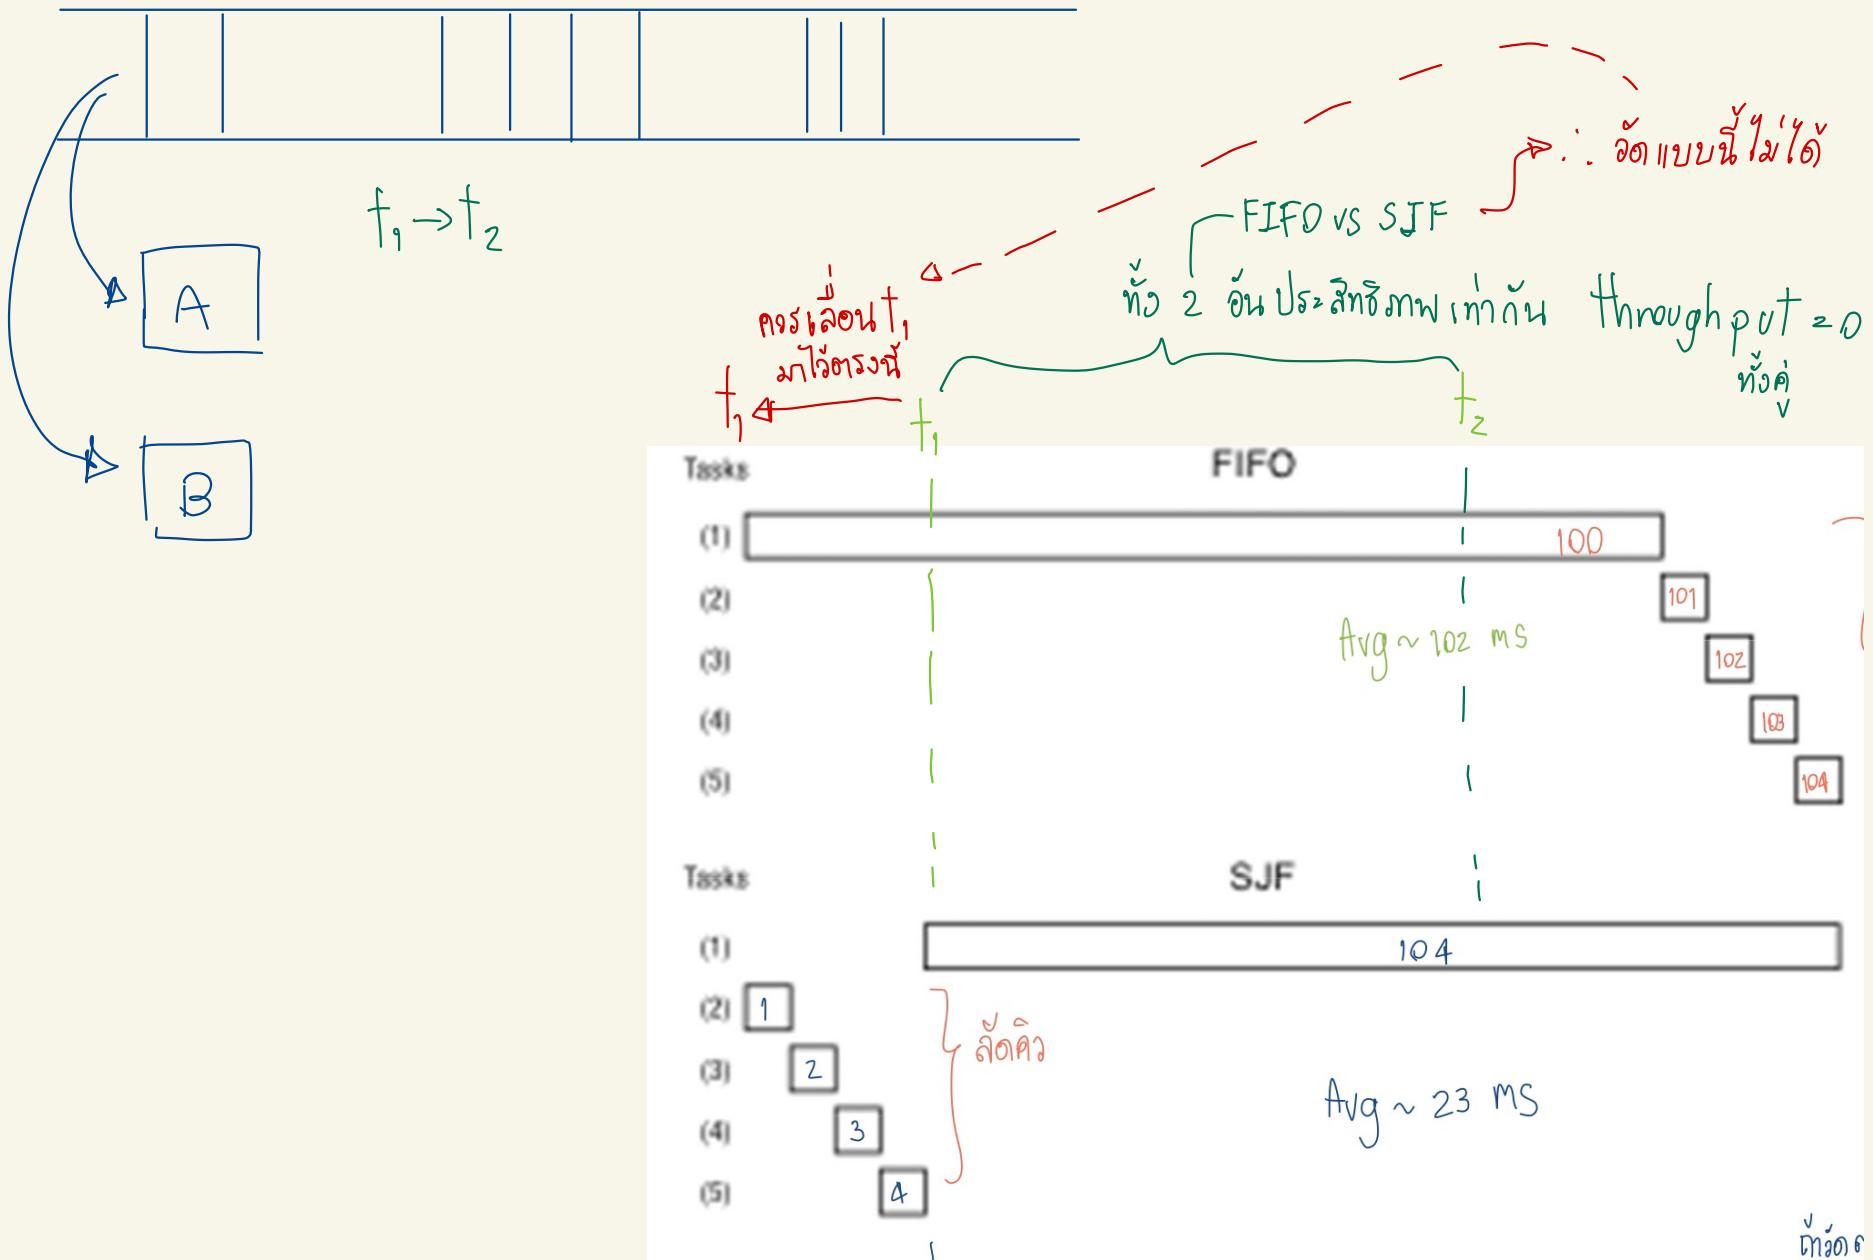
Question

- Is FIFO ever optimal? best case ឱ្យ_FIFO
ការងារដែលមានព័ត៌មានលម្អិត
→ រាងទូទៅ FIFO ក្នុង SJF នឹងឈើនកំណែ
- Pessimal?
ឯកសារ
បញ្ជាផ្ទាល់ព័ត៌មាន មិត្តភីថែលាមាក
អាជីវកម្មប៉ុប្បន្ន

SF
સ્ટારવેશન

Starvation and Sample Bias

- Suppose you want to compare two scheduling algorithms
 - Create some infinite sequence of arriving tasks
 - Start measuring
 - Stop at some point
 - Compute average response time as the average for completed tasks between start and stop
- Is this valid or invalid?



Sample Bias Solutions

- Measure for long enough that # of completed tasks >> # of uncompleted tasks
 - For both systems! បើនឹងព្រមទាំងពារី 2 systems
- Start and stop system in idle periods
 - Idle period: no work to do
 - If algorithms are work-conserving, both will complete the same tasks

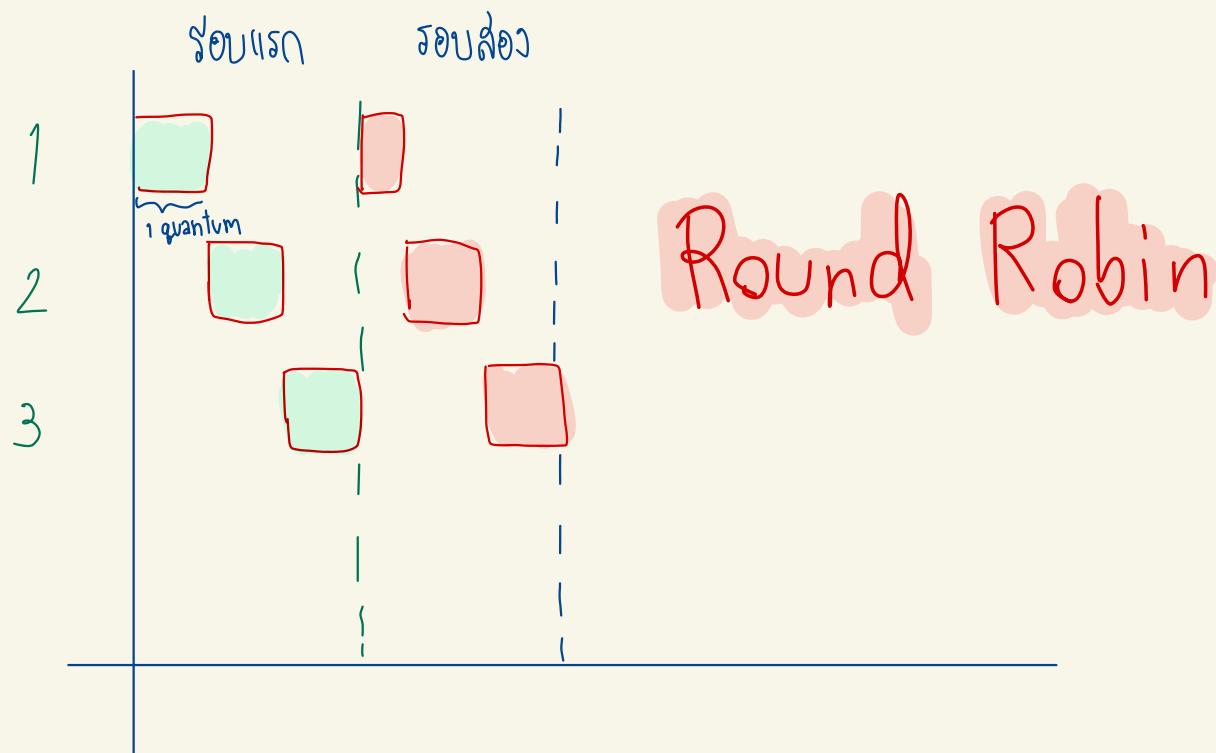
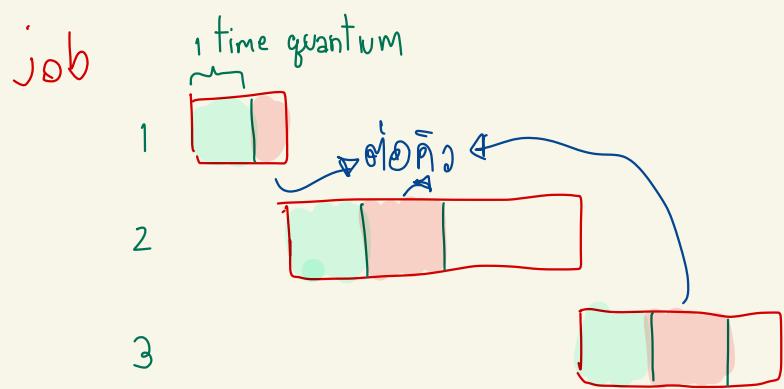
ការ start, stop
នូវការត្រួរបញ្ជី
ក្នុងការសម្រេច
ទៅសម្រេចនៃ workload

Round Robin

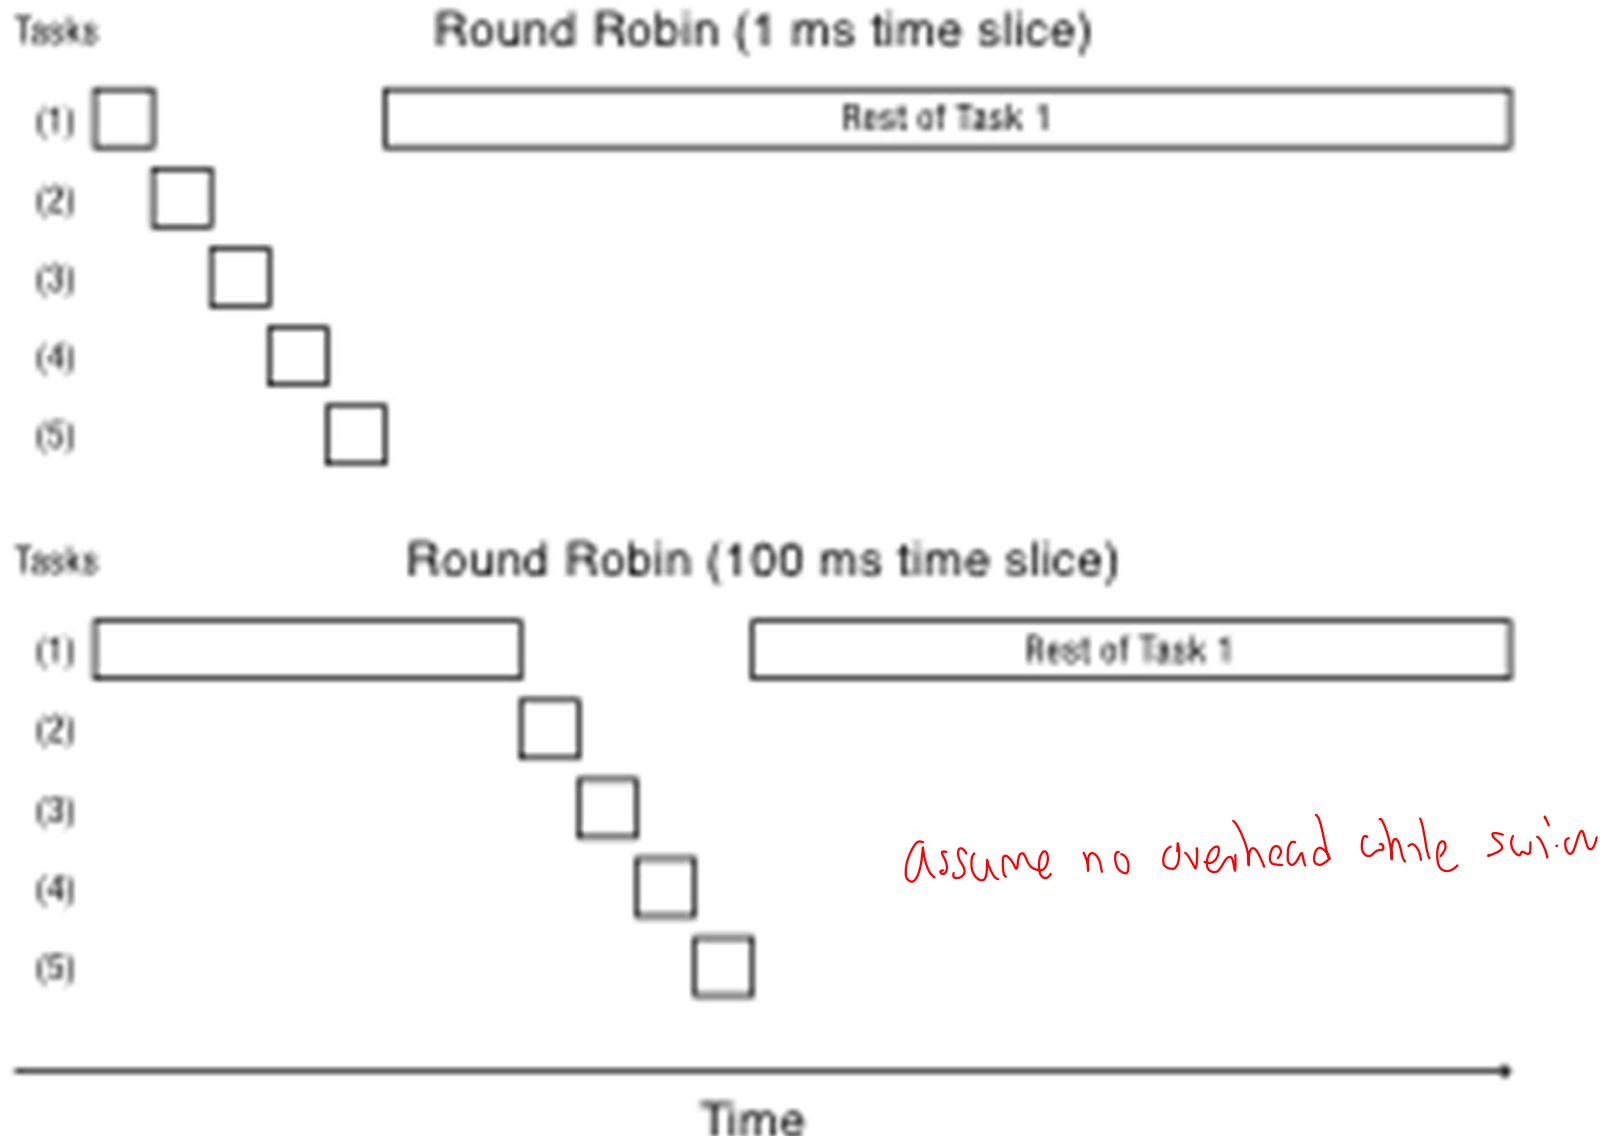
ทำการแบ่ง CPU time ให้กันนุ่มๆ \rightarrow ช่วงเวลาที่ job ได้ = time quantum \rightsquigarrow 1 time quantum
ให้กันนุ่มๆ

↓
ให้กันนุ่มๆ
↓
ทำยังไงสัก
↓
ไปต่อคิวใหม่

- Each task gets resource for a fixed period of time (time quantum)
 - If task doesn't complete, it goes back in line
- Need to pick a time quantum
 - What if time quantum is too long? \rightarrow ต้องรอ \rightarrow ใกล้ FIFO
 - Infinite? \hookrightarrow ถ้ามีขนาดใหญ่มาก $\lim_{n \rightarrow \infty}$ \rightarrow 1 task จะอยู่ใน 1 quantum \rightarrow กลายเป็น FIFO
 - What if time quantum is too short?
 - One instruction? \rightarrow เล็กมากๆ ก็ตามจำนวนที่ต้องการ SJF



Round Robin



Round Robin vs. FIFO

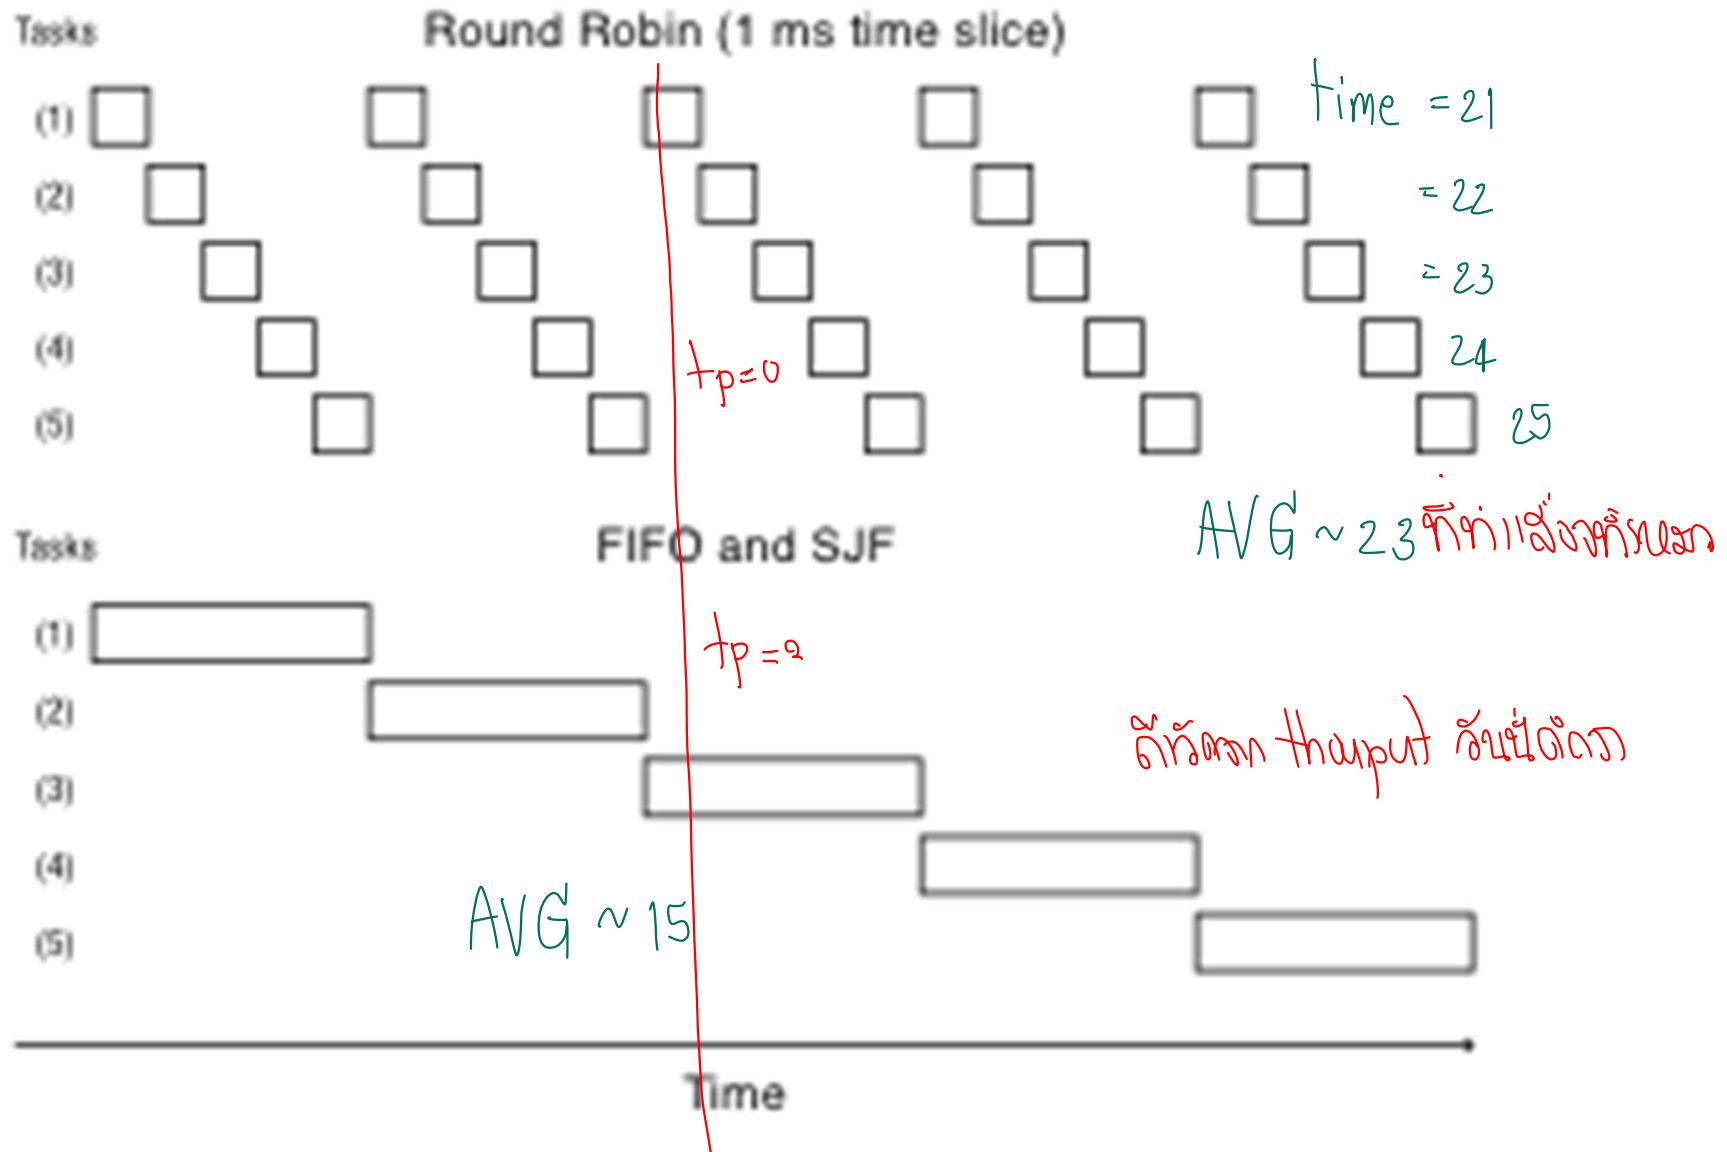
မျှနှုန်းခေါ် slice

- Assuming zero-cost time slice, is Round Robin always better than FIFO?

↳ မျှနှုန်းခေါ်ခြင်းအတွက် ဘက်လဲမှတ်



Round Robin vs. FIFO



↳ ใช่ ถ้าคณิตทาง math

Round Robin = Fairness?

- Is Round Robin always fair?



- What is fair? อะไรที่จะเป็นfair?

- FIFO? มาถึงก่อน?

- Equal share of the CPU? ทุกคนได้สิทธิ์เท่าๆ กับ CPU เท่ากัน?

- What if some tasks don't need their full share? ↳ if last job?
↳ จะมีภาระของหนึ่ง

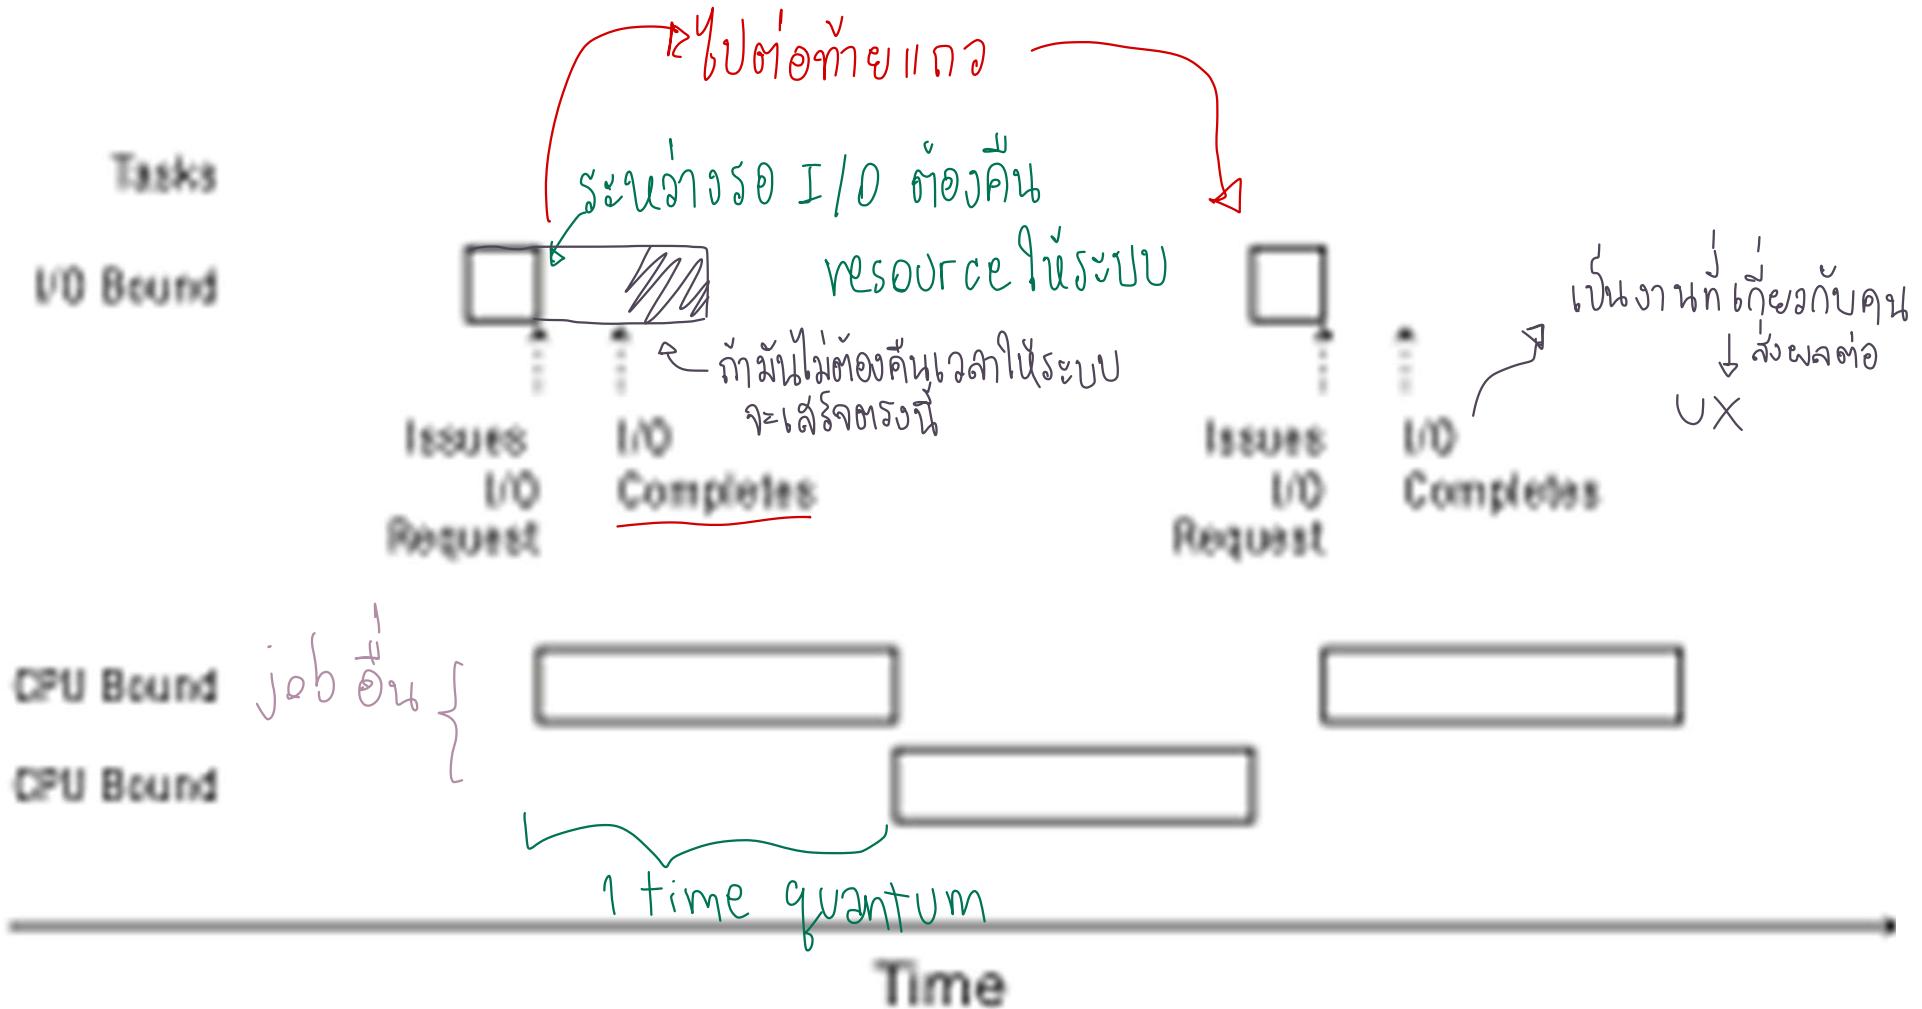
- Minimize worst case divergence? ↳ ไม่ต้องกังวลระบบตลอด ex

- Time task would take if no one else was running

- Time task takes under scheduling algorithm ↳ รอด I/O → วนรอบ yield
(พัฒนาในระบบ)

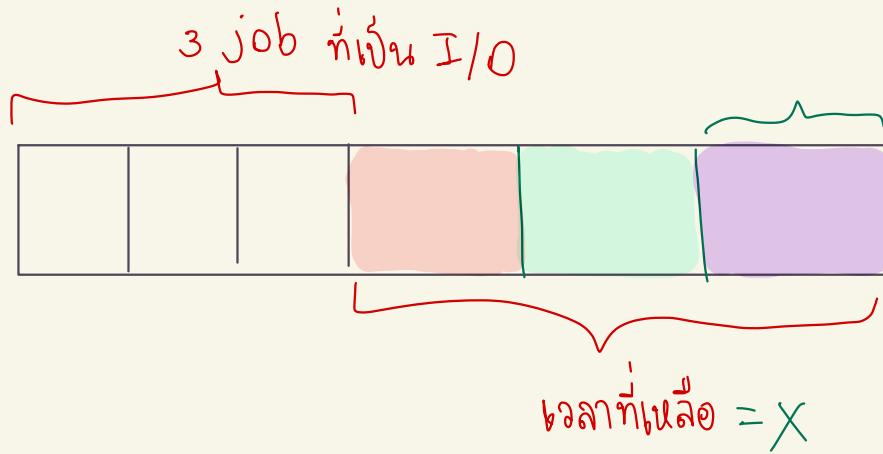
↓
ถ้าเงื่อนไขนี้สุดท้ายชั้งต่อไปลับไปกัน?

Mixed Workload

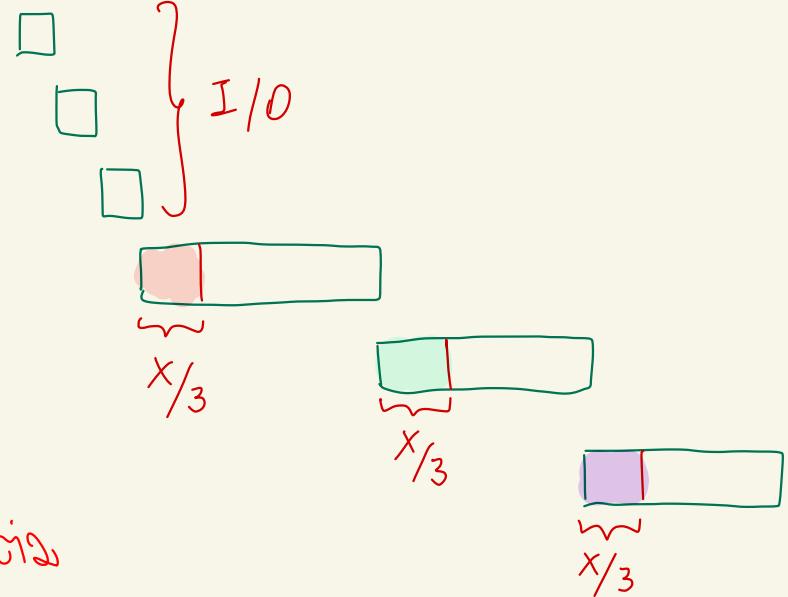


now a) → 9 ឈឺតាមតិចនៅលើលេខ 1 tq នៅក្នុងប្រព័ន្ធ → បានរាយខ្លួន → ដូចនេះបាននៅមី?

1 → 1 time quantum →



โดยสิ่งที่มี I/O ต้อง



Max-Min Fairness

- How do we balance a mixture of repeating tasks:
 - Some I/O bound, need only a little CPU
 - Some compute bound, can use as much CPU as they are assigned
- One approach: maximize the minimum allocation given to a task
 - If any task needs less than an equal share, schedule the smallest of these first
 - Split the remaining time using max-min
 - If all remaining tasks need at least equal share, split evenly

Multi-level Feedback Queue (MFQ)

implement โครงสร้าง

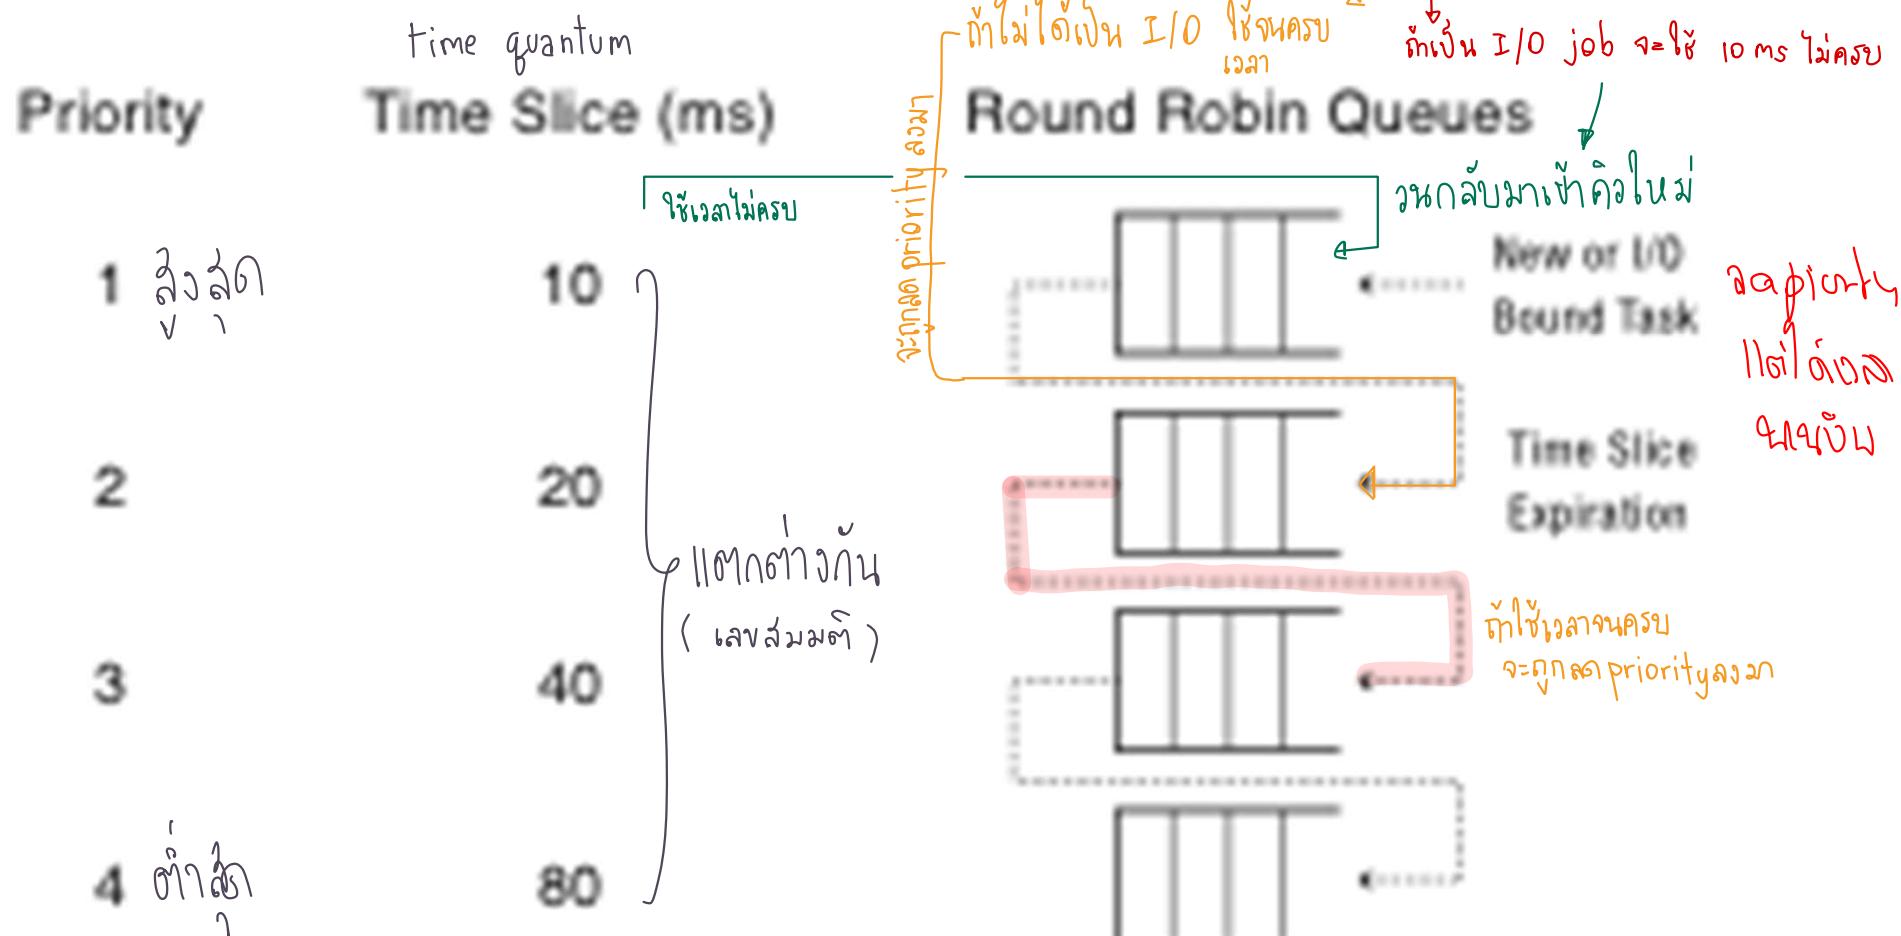
- Goals: ประสิทธิภาพดี
 - Responsiveness ✓
 - Low overhead ✓
 - Starvation freedom ไม่เมื่ง. เมื่งคิ้วนาน job ไม่ได้ทำงาน
 - Some tasks are high/low priority มีก. กำหนด priority
 - Fairness (among equal priority tasks) โปรต์เดย
- Not perfect at any of them! nobody perfect <3
 - Used in Linux (and probably Windows, MacOS)
 - ↓
love the way you are.

MFQ

- Set of Round Robin queues
 - Each queue has a separate priority
- High priority queues have short time slices
 - Low priority queues have long time slices
- Scheduler picks first thread in highest priority queue
- Tasks start in highest priority queue
 - If time slice expires, task drops one level

เมื่อมา job ใหม่ จะต้องว่ามี priority สูงสุด

MFQ



ลักษณะ โถง โดย ใช้ I/O เพิ่มเพื่อให้ job ต้องหงายพื้น prio 1 เลื่อน

แก้ ใช้ parameter ที่กำหนดนับเวลาที่ job พัฒนาใช้ CPU ต้องเวลาที่กำหนดไว้ จะลด priority ลงมา

แต่อาจมี I/O หนัก ที่ใช้เวลามากก็ได้ จะมี parameter อื่นมาแก้ไข

Uniprocessor Summary (1)

- FIFO ง่ายๆ ทำได้
 → FIFO is simple and minimizes overhead.
- If tasks are variable in size, then FIFO can have very poor average response time. ถ้า job มีขนาดใหญ่ FIFO จะช้า
- If tasks are equal in size, FIFO is optimal in terms of average response time. → ขนาดเท่ากัน = optimal FIFO
- Considering only the processor, SJF is optimal in terms of average response time. ที่การคำนวณเวลา avg response time
- SJF is pessimal in terms of variance in response time.
↳ กรณีจัดการจาก
↓ ตัวอย่าง

Uniprocessor Summary (2)

- If tasks are variable in size, Round Robin approximates SJF.
ເຖິງບເຄື່ອງໃຈ
- If tasks are equal in size, Round Robin will have very poor average response time.
- Tasks that intermix processor and I/O benefit from SJF and can do poorly under Round Robin.

↓
ຕີໃນ workload ມີໃຫຍ່ I/O, ຍັງໃຫຍ່ ພສມກັນ

↓
Round Robin ຍົມ work → ຍົມໄວະບະ || ອົງຄົວ
↓
SJF ຂົດໂຄກວ່າ ✓

→ ດີຍມາດຫາກັນນີ້ແລ້ວ → ແລ້ວ

ສູນໄວ່ຈະ → ດັດໄວ່ແລ້ວ

Uniprocessor Summary (3)

- Max-Min fairness can improve response time for I/O-bound tasks.
- Round Robin and Max-Min fairness both avoid starvation.
- By manipulating the assignment of tasks to priority queues, an MFQ scheduler can achieve a balance between responsiveness, low overhead, and fairness.

→ កំណត់ប្រសិទ្ធភាពនៃវា

ដោយ implement → MFQ

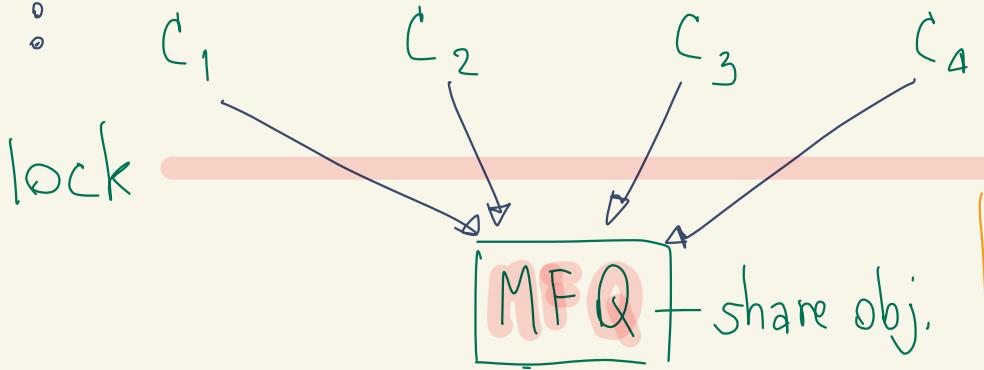
សេបលុយរប. \rightsquigarrow multi core

Multiprocessor Scheduling

- What would happen if we used MFQ on a multiprocessor? \rightarrow ចិត្តឱ្យ
 - Contention for scheduler spinlock
 - Cache slowdown due to ready list data structure pinging from one CPU to another
 - Limited cache reuse: thread's data from last time it ran is often still in its old cache \downarrow switch រាង
អែក្រុង load data នៅក្នុងវា



processor :



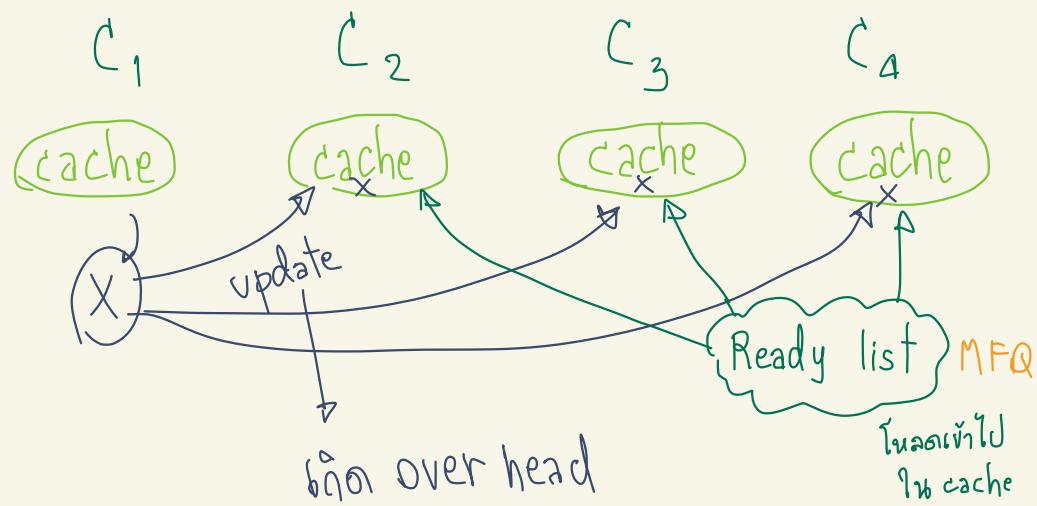
จ่ายงานใน
processor
ให้ล่าสุด

↓
ลักษณะนี้มีอยู่
share obj.

problem: 1. ต้องต่อที่จะ job

2. cache

processor :



|| ต่อ processor

มี cache
ของตัวเอง

ถ้าทุกตัวต้องก.

ทำงานกับตัวไป
เรียงกัน

ในล็อกเป้าไป
ใน cache
(check ค่า X กับฟันปูล. ยัง)

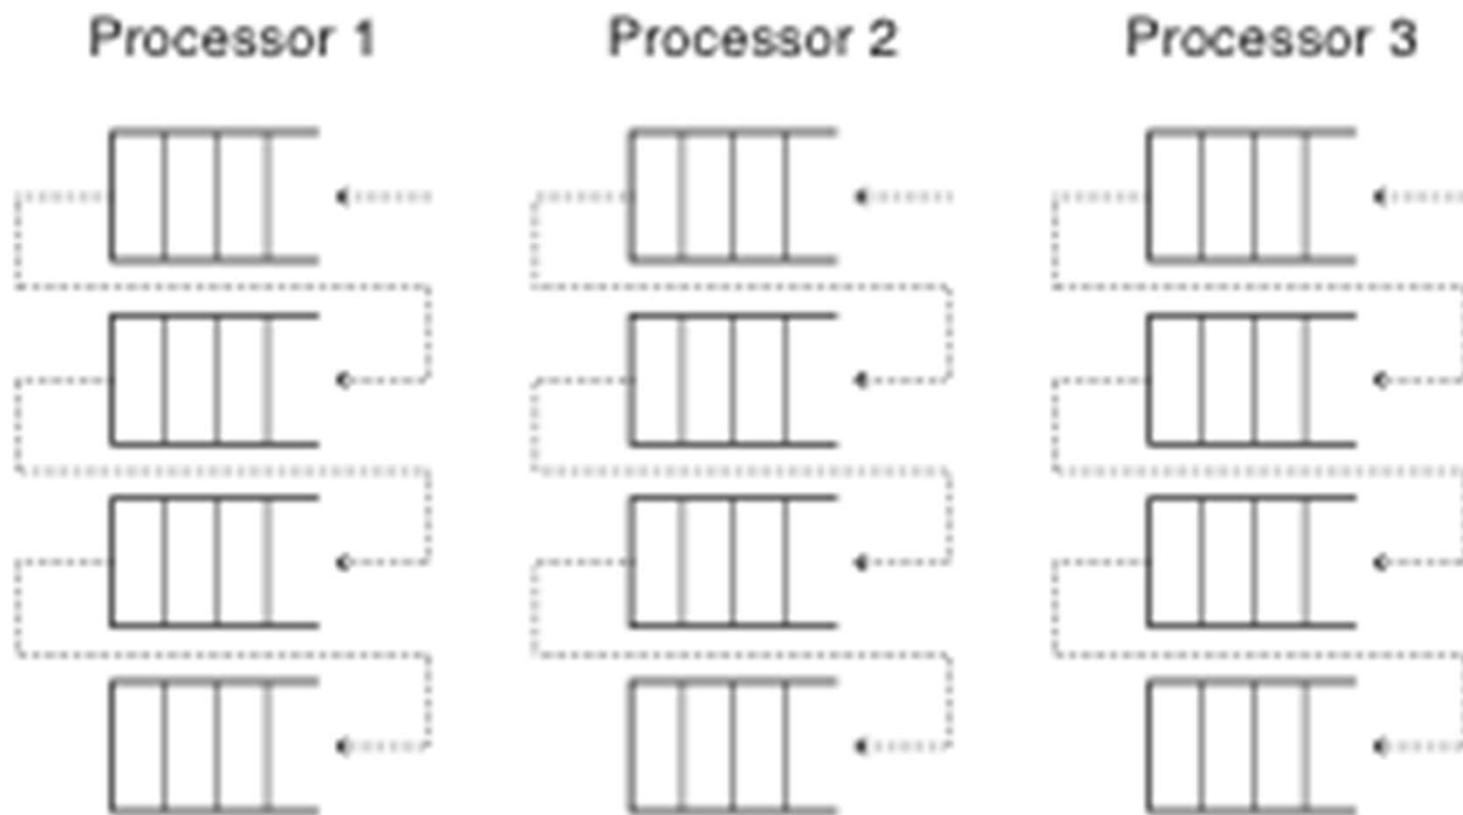
↳ มีผลต่อก. ทำงานของ scheduler

Per-Processor Affinity Scheduling

- Each processor has its own ready list → แต่ละ processor มี ready list (MFQ) เป็นของตัวเอง
- Protected by a per-processor spinlock
- Put threads back on the ready list where it had most recently run
– Ex: when I/O completes, or on Condition->signal
 ↓
 ไม่ต้องไป check
 ก็ update ที่ลุ้นอยู่
- Idle processors can steal work from other processors กรณี processor ที่ว่างงานอยู่ สามารถไปปีนอยู่จาก list ของตัวอื่นได้



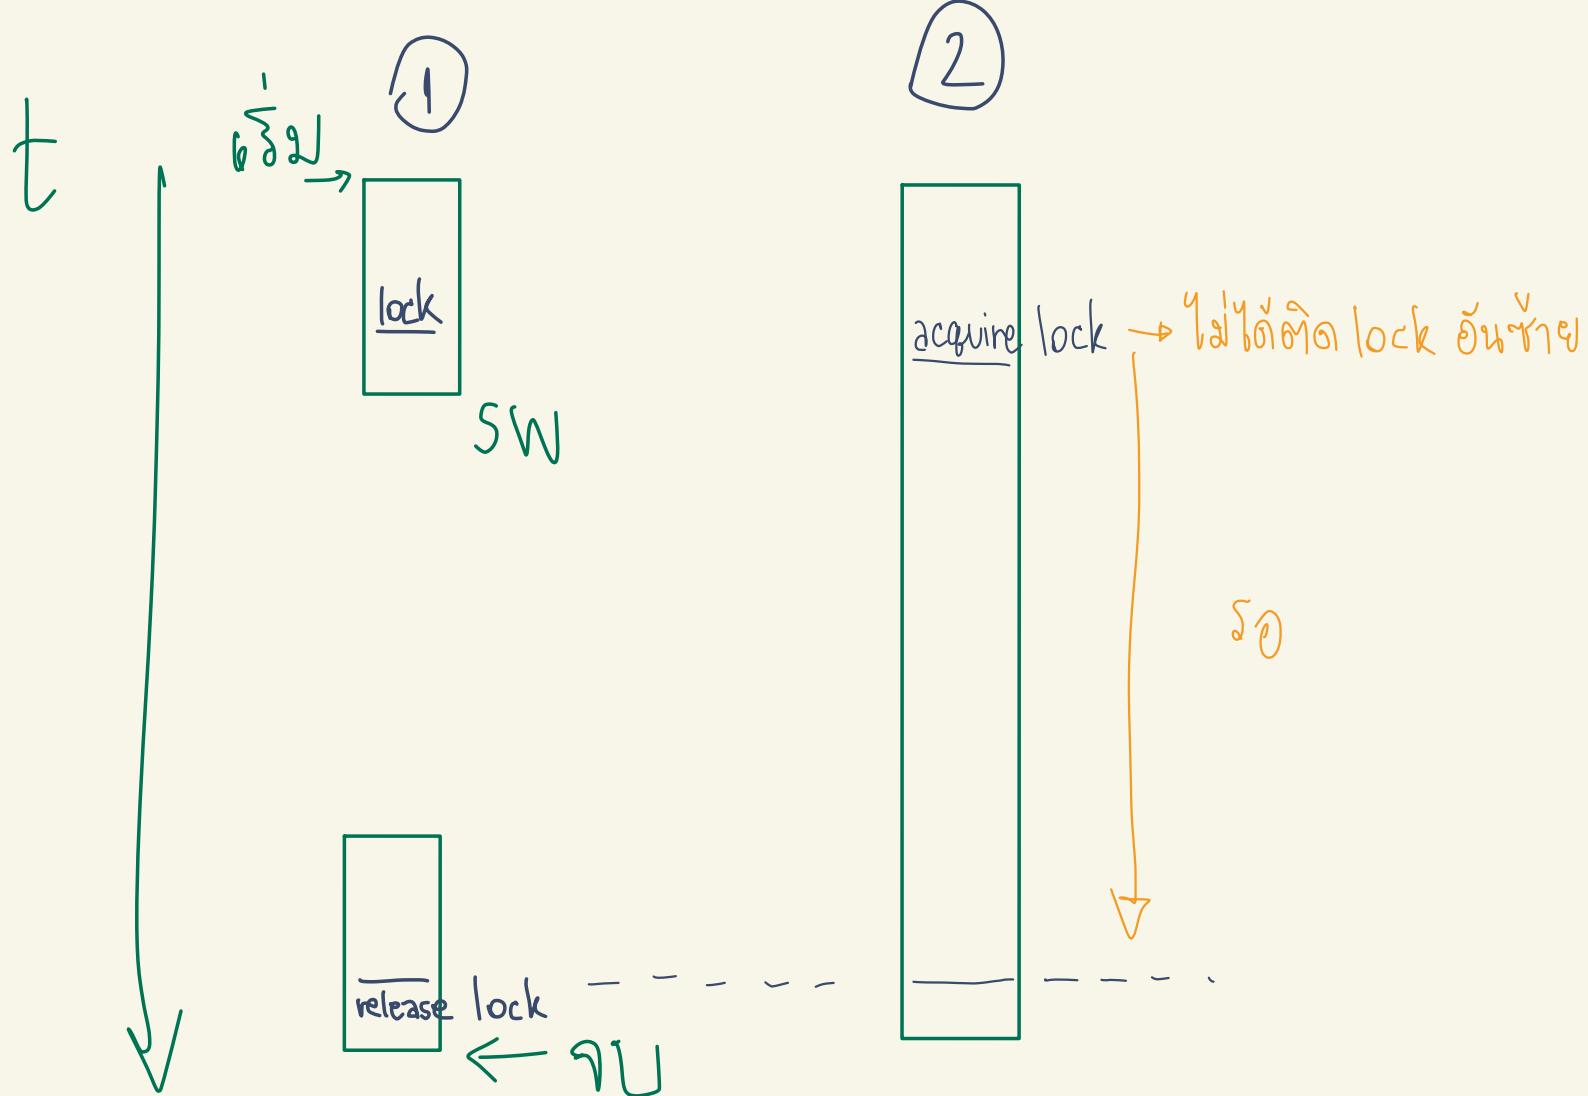
Per-Processor Multi-level Feedback with Affinity Scheduling



ការណែន thread ឲ្យរាយការ ទាន់ភ្លាមៗ ក្នុង

Scheduling Parallel Programs

- What happens if one thread gets time-sliced while other threads from the same program are still running?
 - Assuming program uses locks and condition variables, it will still be correct
 - What about performance?

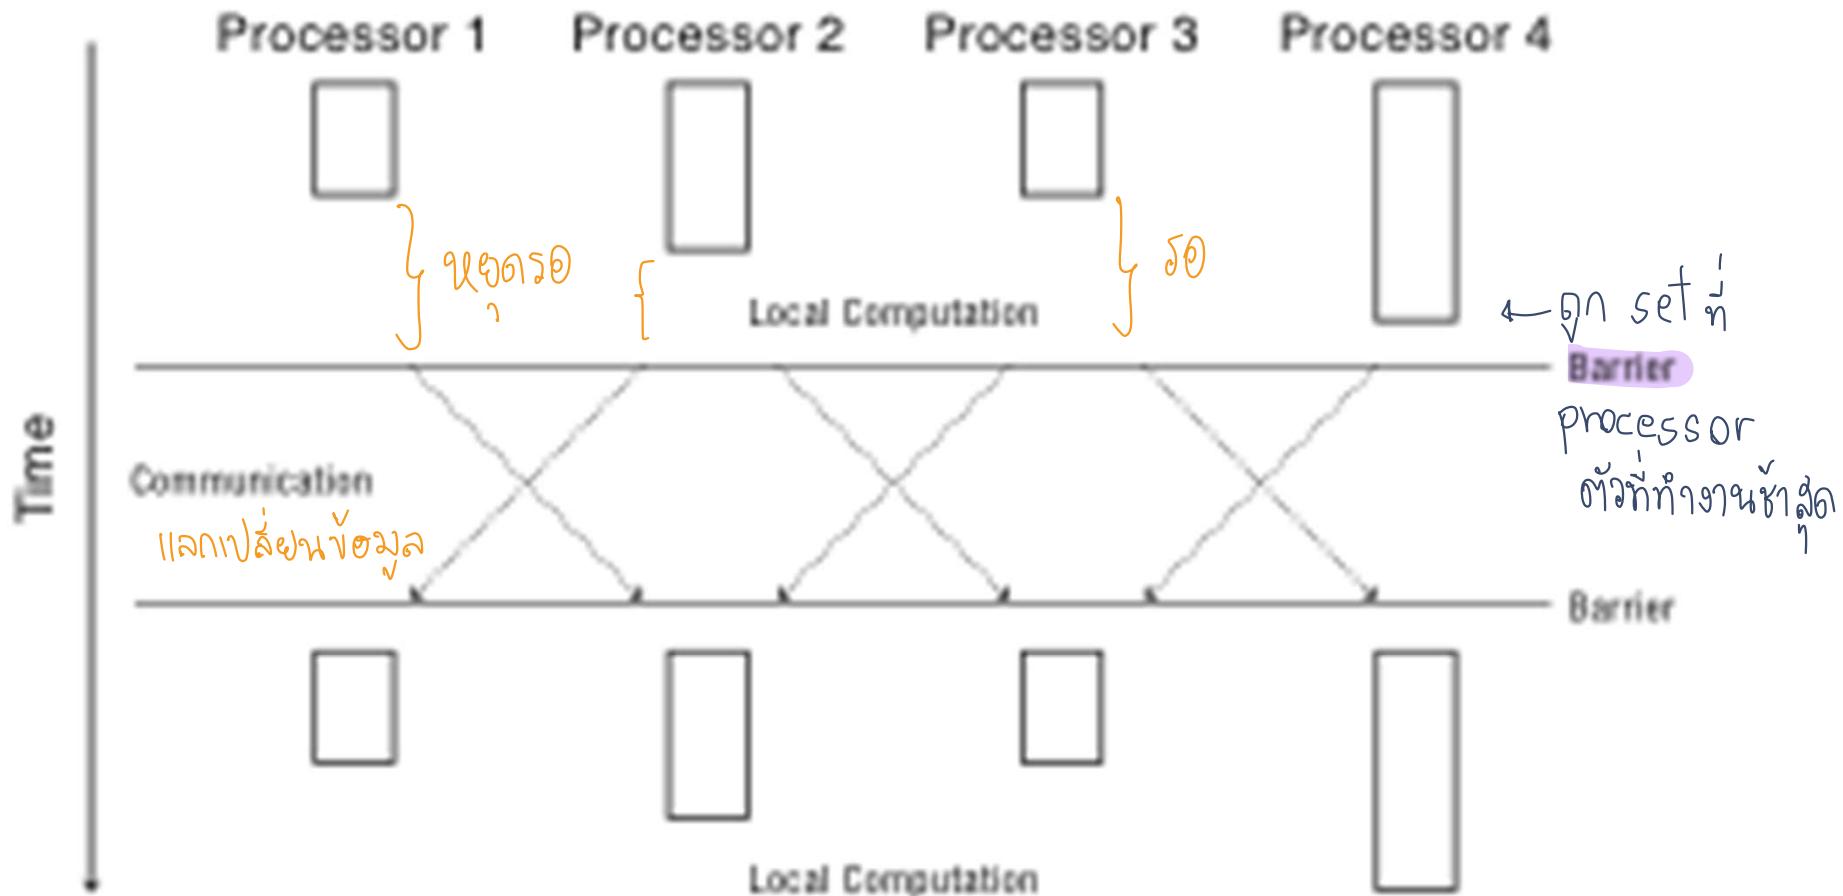


หลักใช้เขียน program แบบ Parallel

Bulk Synchronous Parallelism

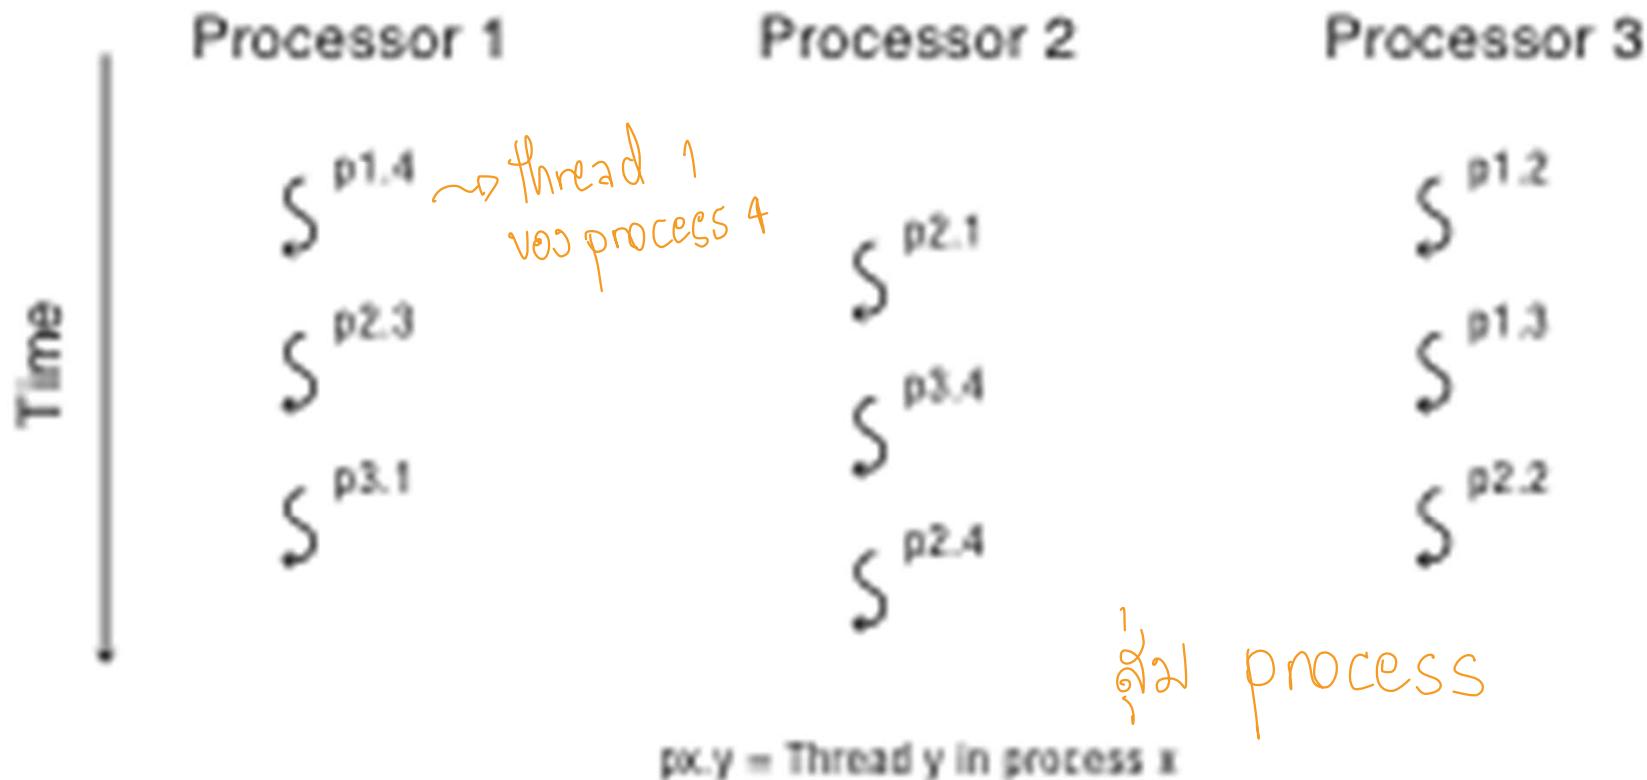
- Loop at each processor:
 - Compute on local data (in parallel) คำนวณ
 - Barrier $\xrightarrow{\text{set}}$ ส่ง data ที่คำนวณแล้ว และเปลี่ยนป้อมูลกัน
 - Send (selected) data to other processors (in parallel)
 - Barrier $\xrightarrow{\text{清除}}$
- Examples:
 - MapReduce
 - Fluid flow over a wing
 - Most parallel algorithms can be recast in BSP
 - Sacrificing a small constant factor in performance

Tail Latency

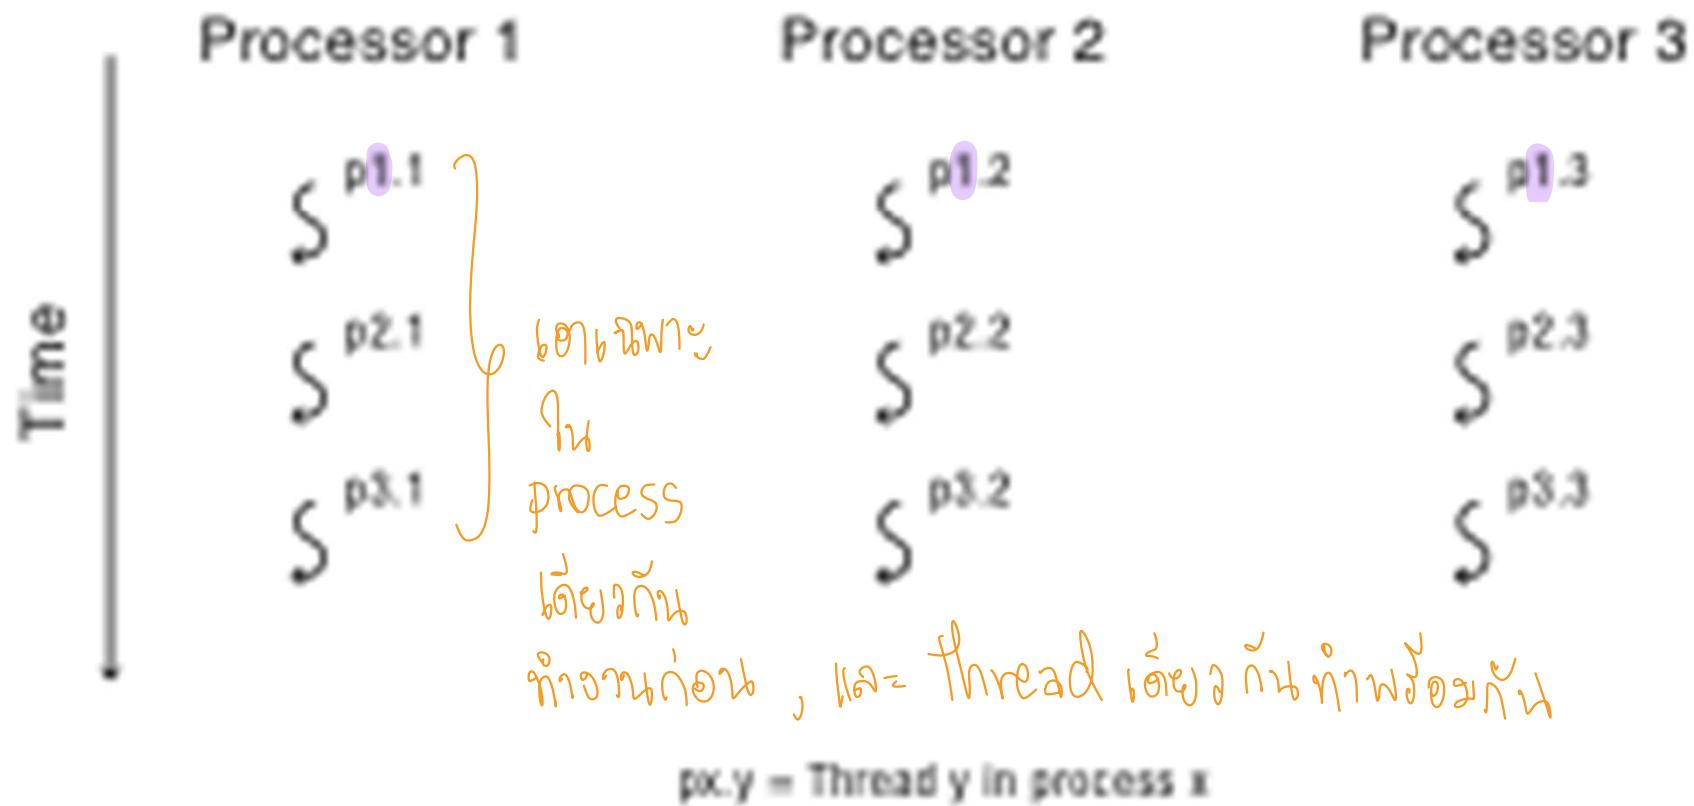


Scheduling Parallel Programs

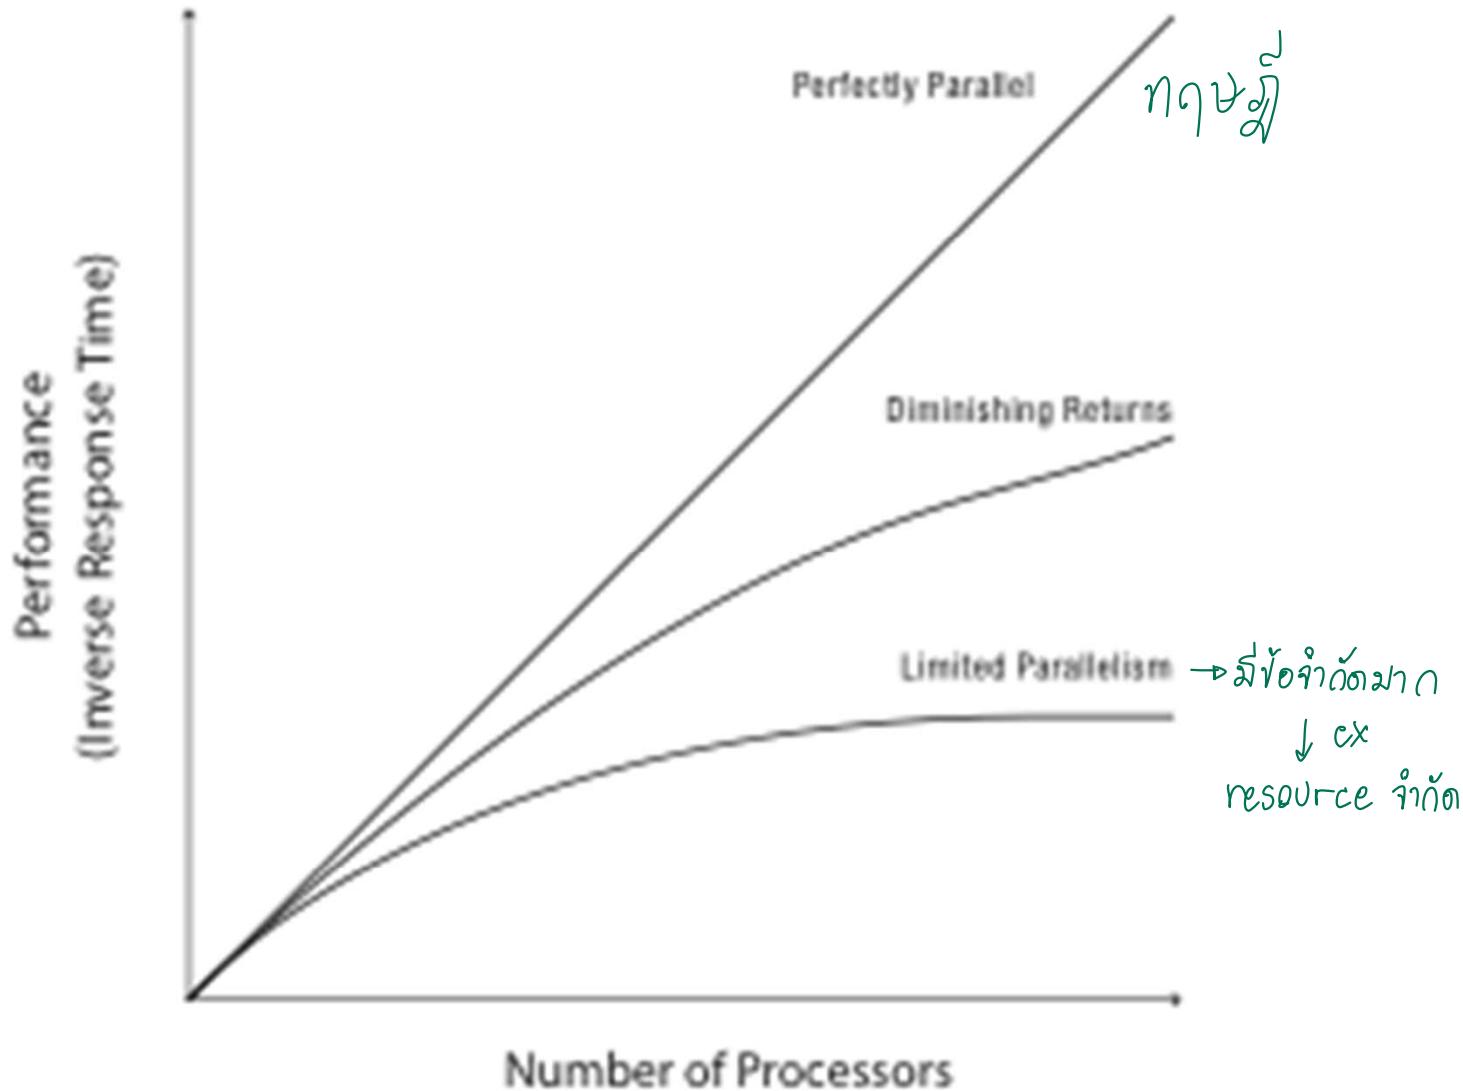
Oblivious: each processor time-slices its ready list independently of the other processors



Gang Scheduling



Parallel Program Speedup

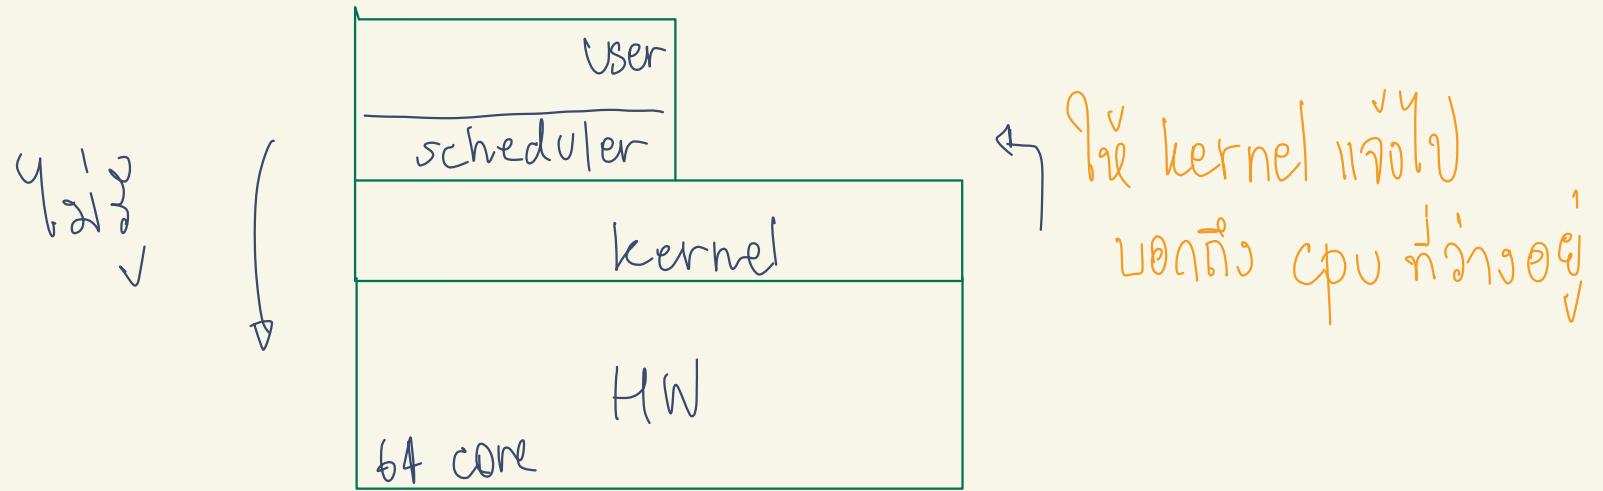


Space Sharing



Scheduler activations: kernel tells each application its # of processors with upcalls every time the assignment changes

▷ ផ្លាស់បន្ថែម User space → អាមេរិកត្រូវតាមរយៈការងារ hardware



មានសំណង

Queueing Theory

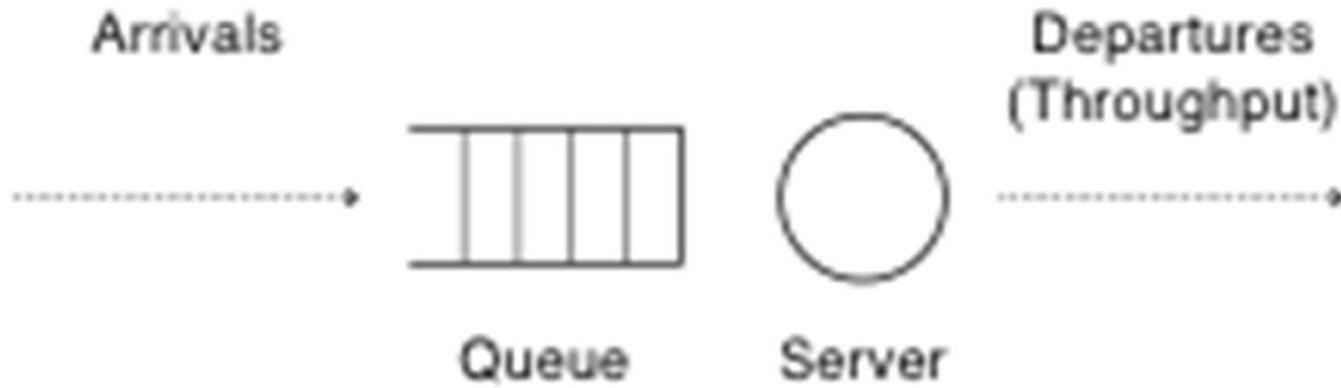
មិនអាចសរុប

- Can we predict what will happen to user performance:
 - If a service becomes more popular?
 - If we buy more hardware?
 - If we change the implementation to provide more features?

បើន stat math → វិធានប្រតិបត្តិកម្ម ទូទៅការចែត នូវការ

គាំងនេះ avg. time នៃលទ្ធផើយទៅ

Queueing Model



Assumption: average performance in a stable system,
where the arrival rate (λ) matches the departure rate (μ)

Definitions

- Queueing delay (W): wait time
 - Number of tasks queued (Q)
- Service time (S): time to service the request
- Response time (R) = queueing delay + service time
- Utilization (U): fraction of time the server is busy
 - Service time * arrival rate (λ)
- Throughput (X): rate of task completions
 - If no overload, throughput = arrival rate

Little's Law

$$N = X * R$$

N: number of tasks in the system

Applies to *any* stable system – where arrivals
match departures.

Question

Suppose a system has throughput (X) = 100 tasks/s,
average response time (R) = 50 ms/task

- How many tasks are in the system on average?
- If the server takes 5 ms/task, what is its utilization?
- What is the average wait time?
- What is the average number of queued tasks?

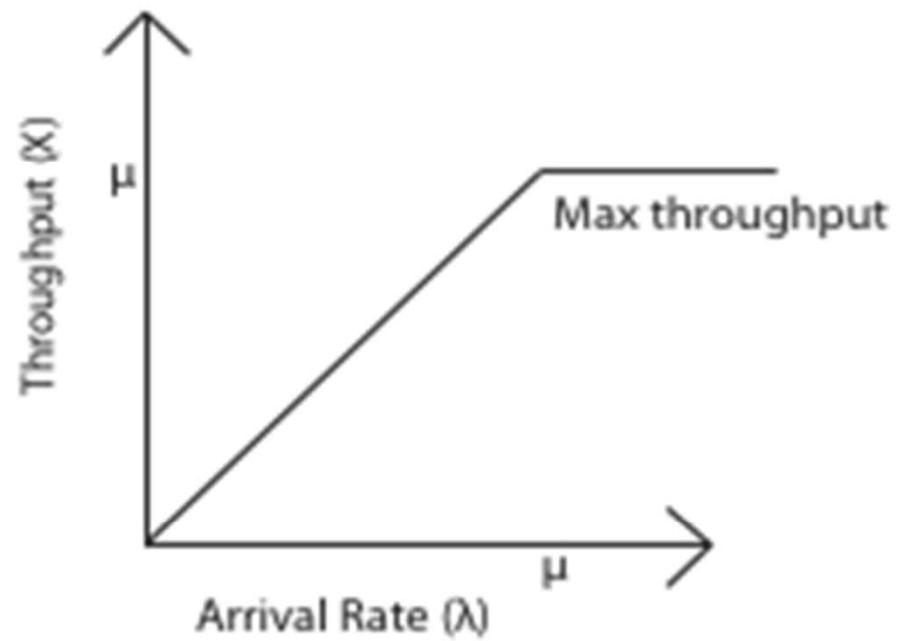
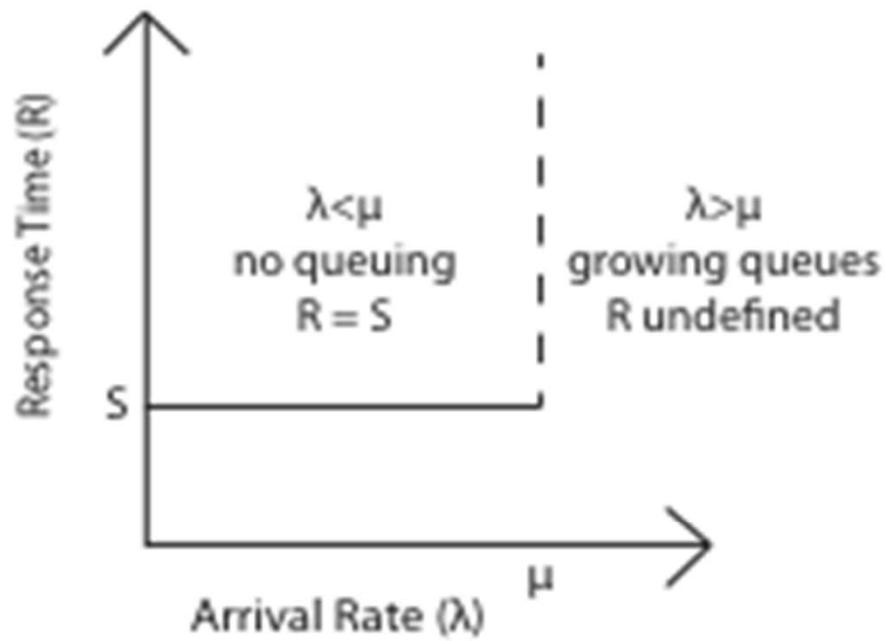
Question

- From example:
 - $X = 100 \text{ task/sec}$
 - $R = 50 \text{ ms/task}$
 - $S = 5 \text{ ms/task}$
 - $W = 45 \text{ ms/task}$
 - $Q = 4.5 \text{ tasks}$
- Why is $W = 45 \text{ ms}$ and not $4.5 * 5 = 22.5 \text{ ms}$?
 - Hint: what if $S = 10\text{ms}$? $S = 1\text{ms}$?

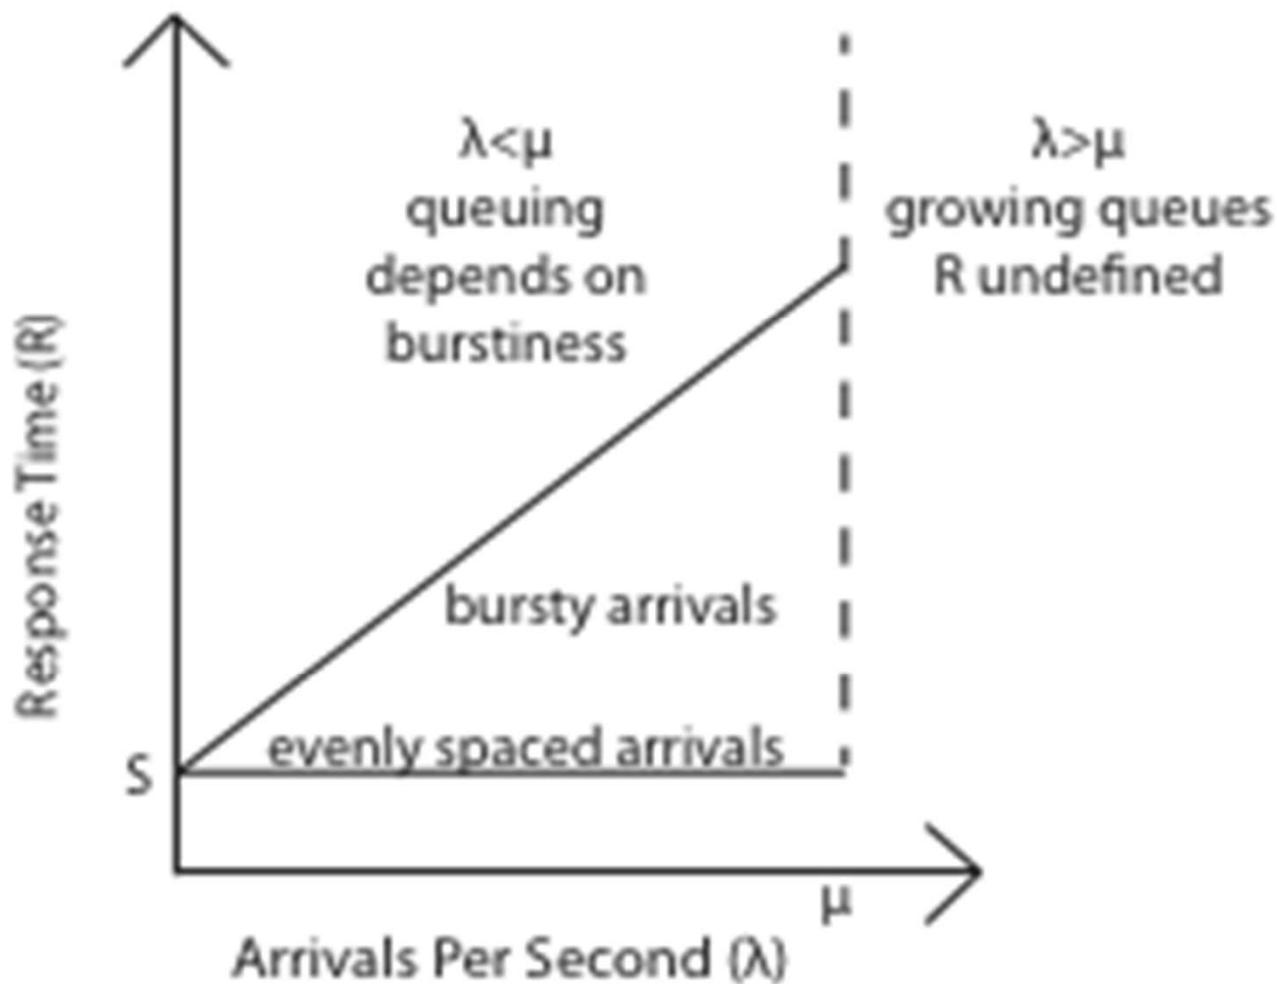
Queueing

- What is the best case scenario for minimizing queueing delay?
 - Keeping arrival rate, service time constant
- What is the worst case scenario?

Queueing: Best Case



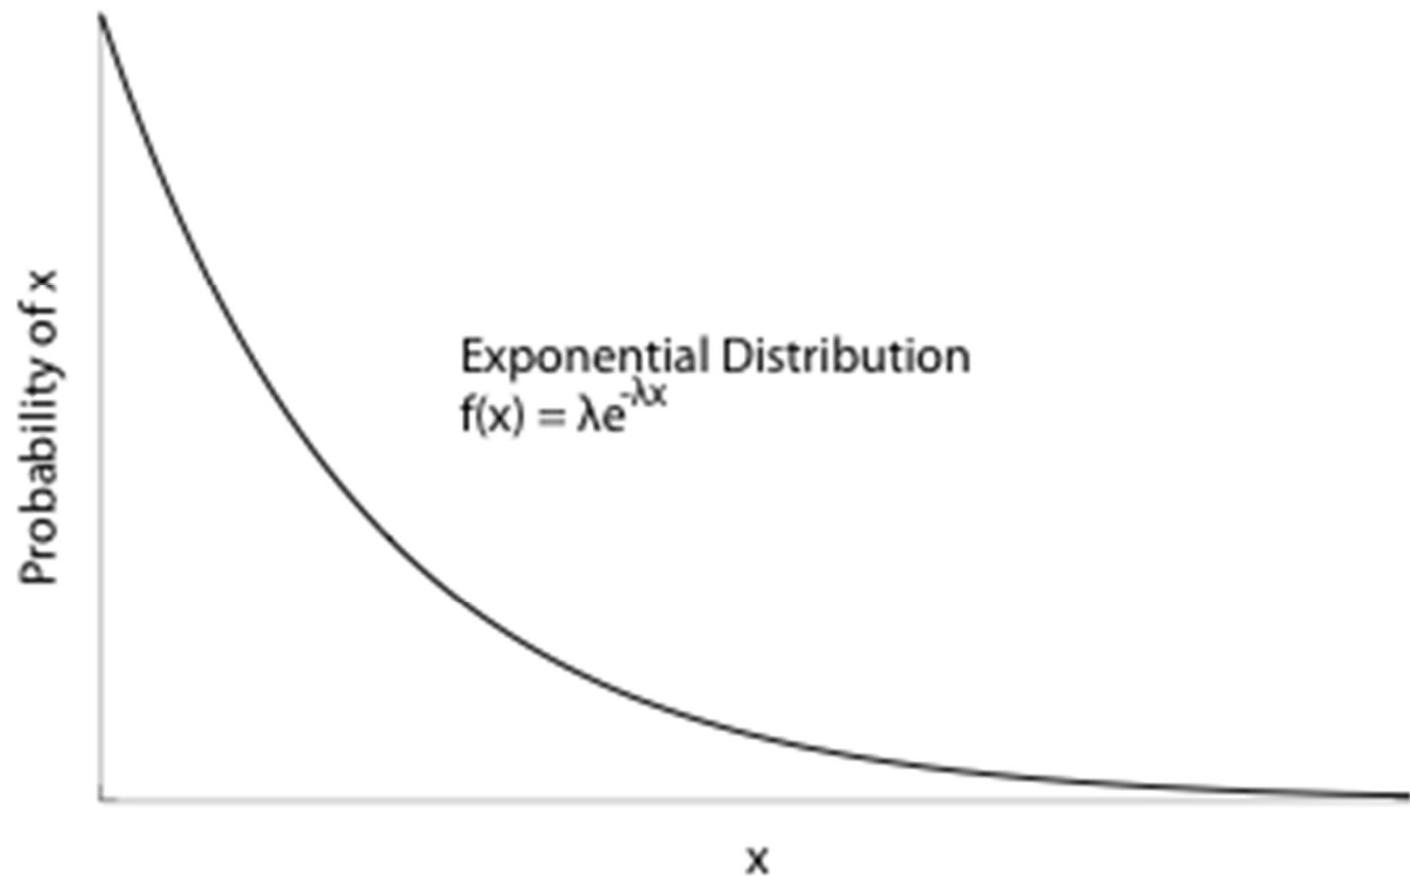
Response Time: Best vs. Worst Case



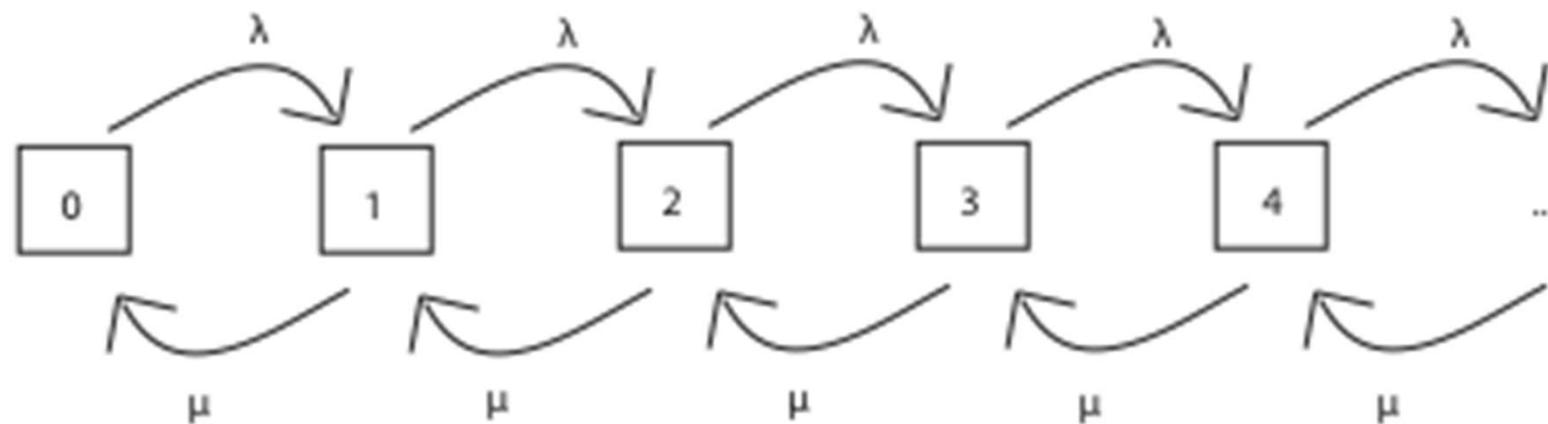
Queueing: Average Case?

- What is average?
 - Gaussian: Arrivals are spread out, around a mean value
 - Exponential: arrivals are memoryless
 - Heavy-tailed: arrivals are bursty
- Can have randomness in both arrivals and service times

Exponential Distribution

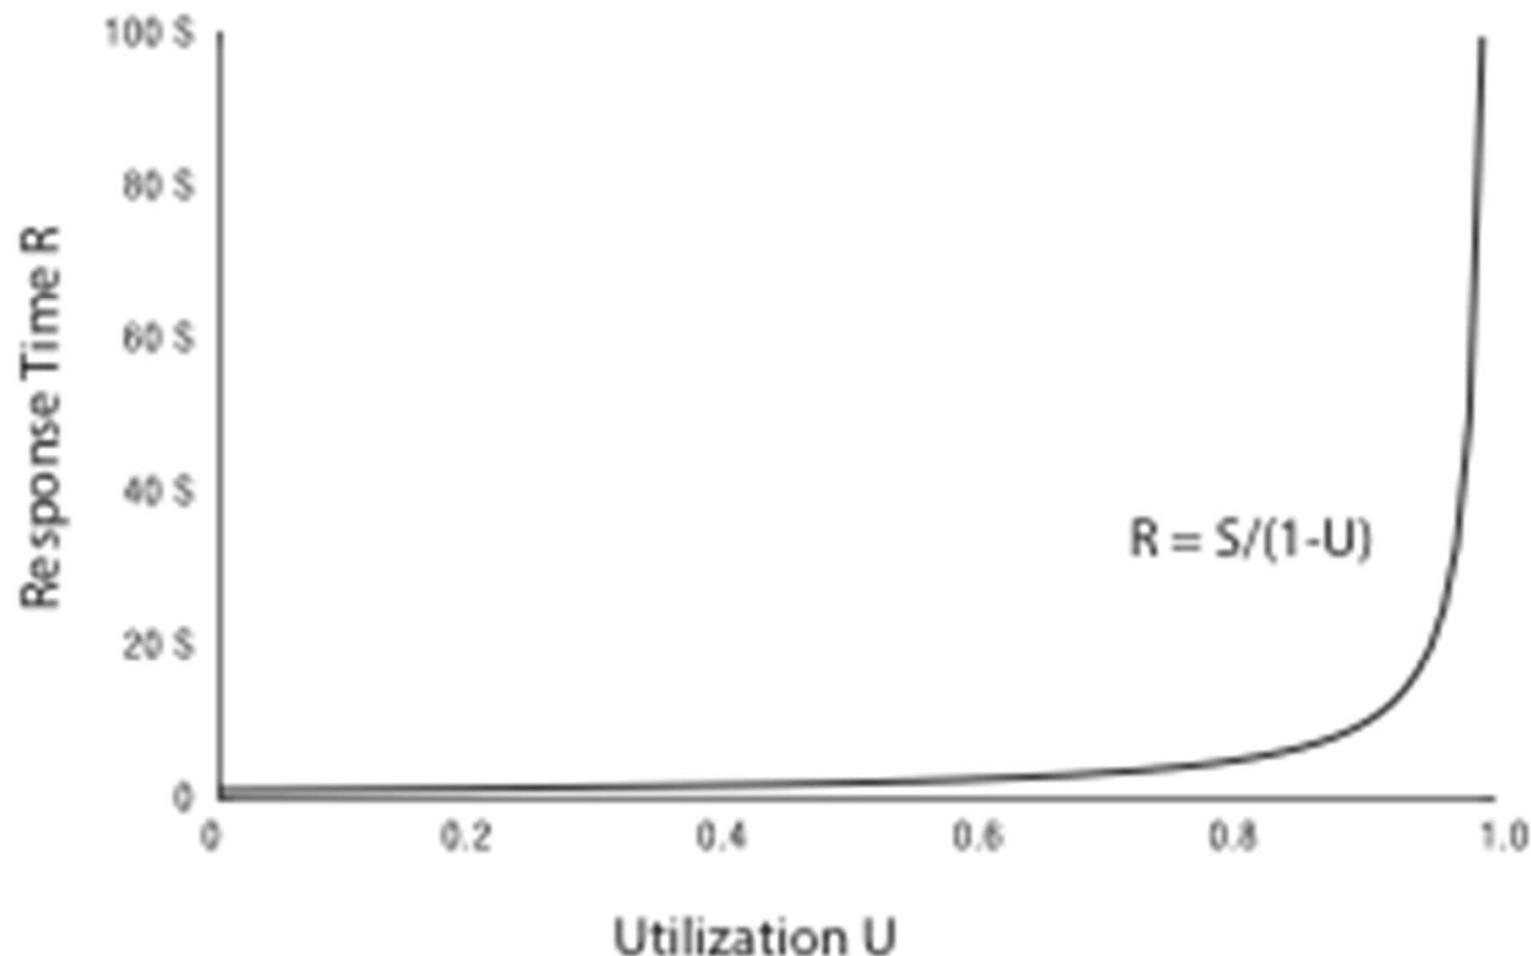


Exponential Distribution



Permits closed form solution to state probabilities,
as function of arrival rate and service rate

Response Time vs. Utilization



Question

- Exponential arrivals: $R = S/(1-U)$
- If system is 20% utilized, and load increases by 5%, how much does response time increase?
- If system is 90% utilized, and load increases by 5%, how much does response time increase?

Variance in Response Time

- Exponential arrivals
 - Variance in R = $S/(1-U)^2$
- What if less bursty than exponential?
- What if more bursty than exponential?

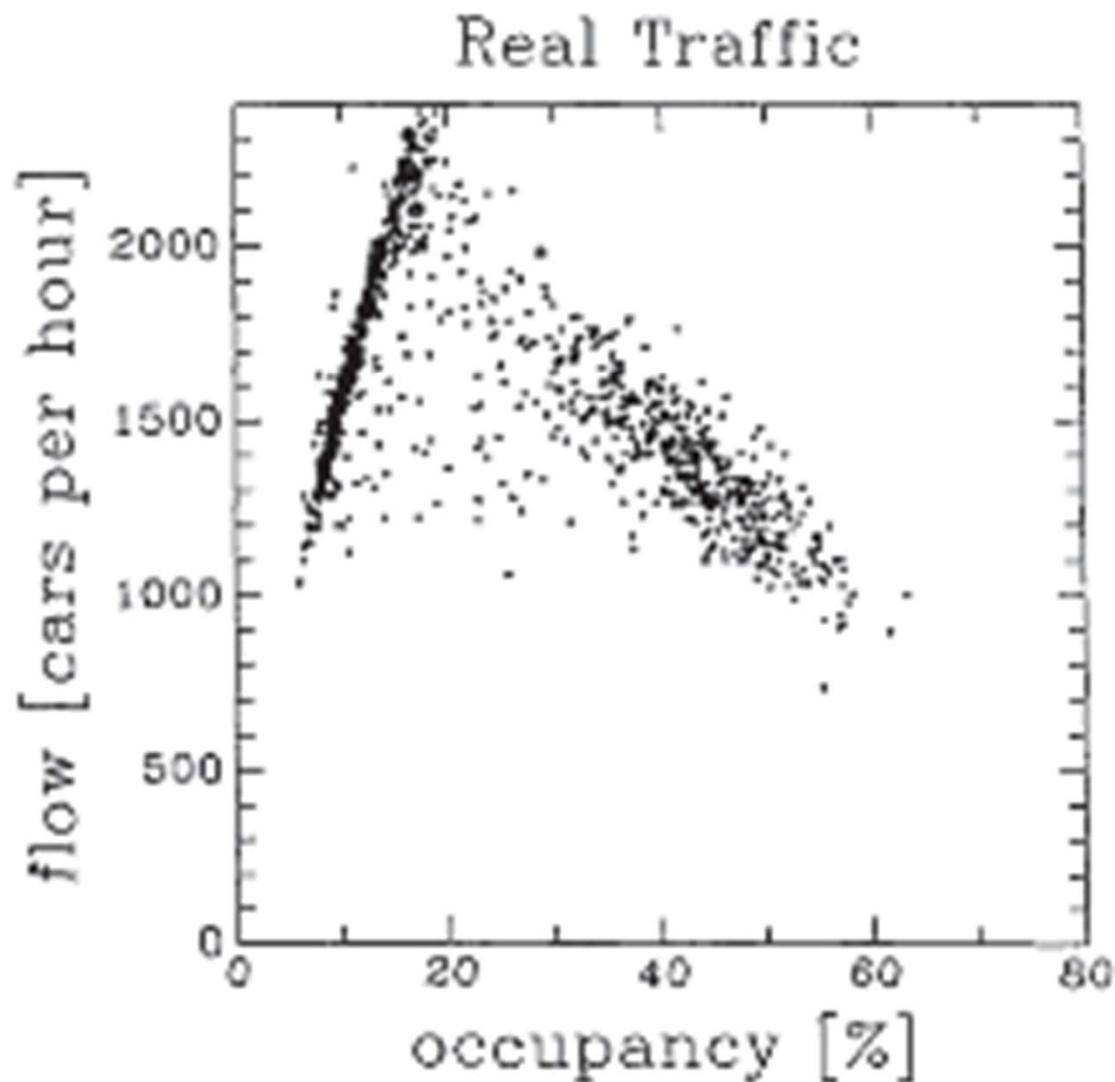
What if Multiple Resources?

- Response time =
Sum over all i
$$\frac{\text{Service time for resource } i}{(1 - \text{Utilization of resource } i)}$$
- Implication
 - If you fix one bottleneck, the next highest utilized resource will limit performance

Overload Management

- What if arrivals occur faster than service can handle them
 - If do nothing, response time will become infinite
- Turn users away?
 - Which ones? Average response time is best if turn away users that have the highest service demand
 - Example: Highway congestion
- Degrade service?
 - Compute result with fewer resources
 - Example: CNN static front page on 9/11

Highway Congestion (measured)



Why Do Metro Buses Cluster?

- Suppose two Metro buses start 15 minutes apart
 - Why might they arrive at the same time?