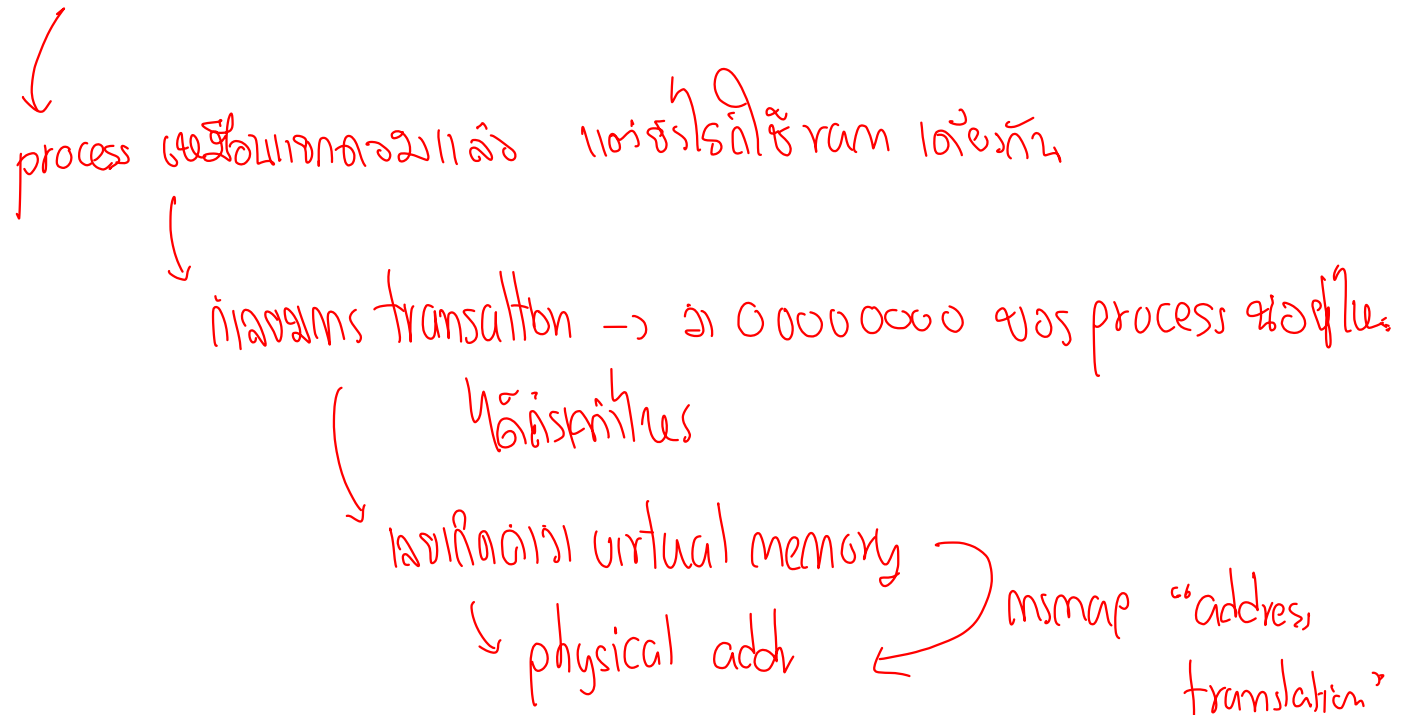


Address Translation



Main Points

- Address Translation Concept
 - How do we convert a virtual address to a physical address?
- Flexible Address Translation
 - Base and bound
 - Segmentation
 - Paging
 - Multilevel translation
- Efficient Address Translation *miss overhead*
 - Translation Lookaside Buffers
 - Virtually and physically addressed caches

Address Translation Goals

- Memory protection → ให้อุป process ใส, แล=จึ้ process
- Memory sharing (analgo) security ให้อุป process ใส addr ตวาวไรที่ 0
↳ check boundary ให้อุป process? ↳ ให้อุป process ใส addr
- Shared libraries, interprocess communication
- Sparse addresses → ให้อุป process ใส addr ใสที่ 0
- Multiple regions of dynamic allocation (heaps/stacks)
- Efficiency
- Memory placement
- Runtime lookup
- Compact translation tables

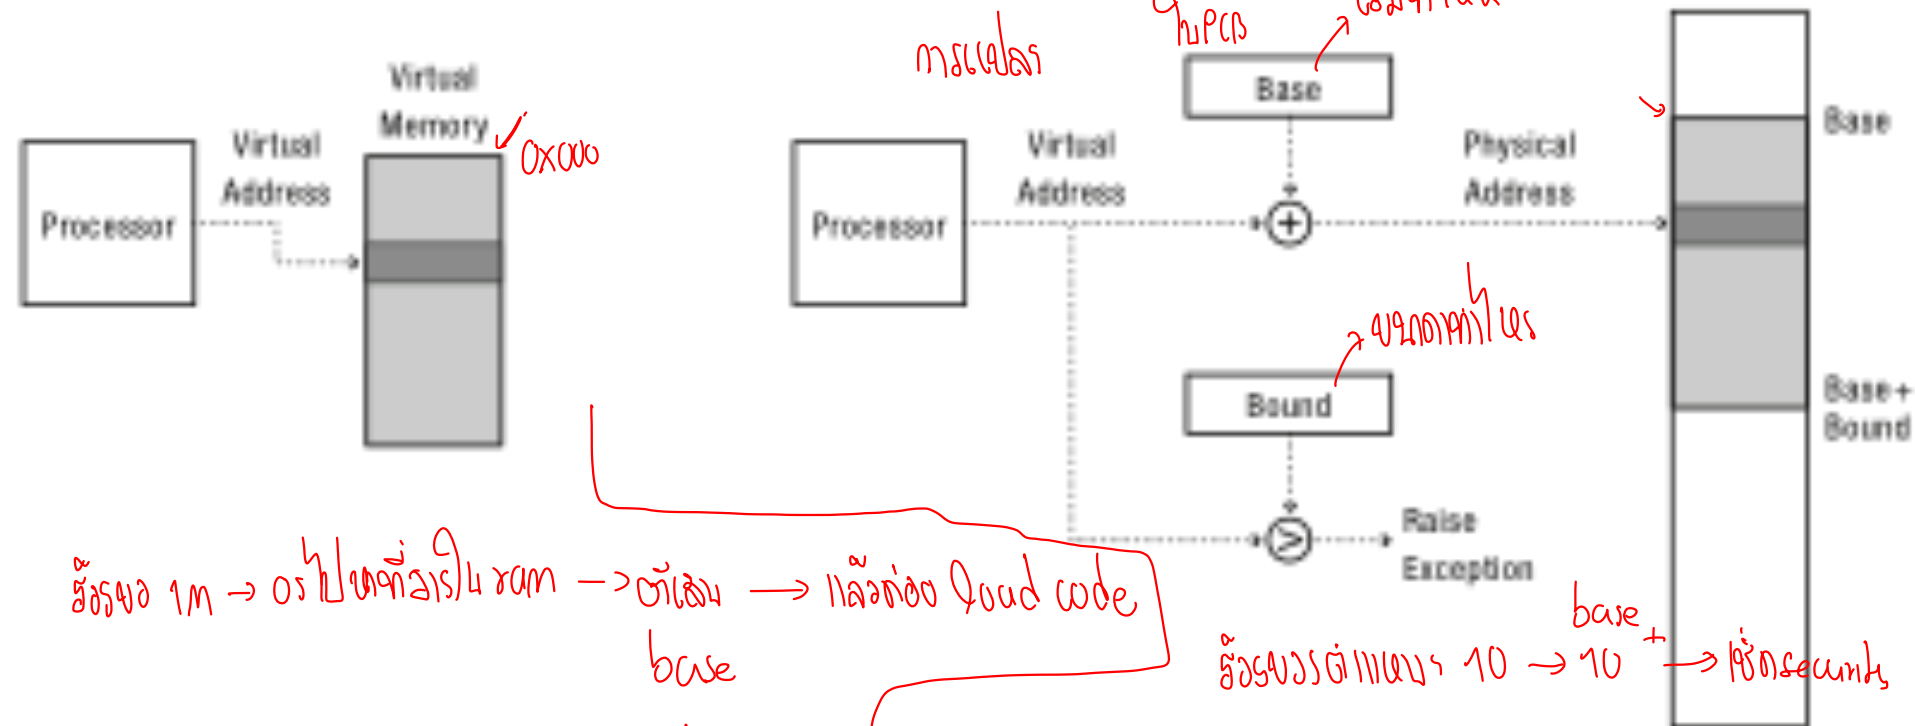
Virtually Addressed Base and Bounds

algorithm 118n

Processor's View

Implementation

Physical Memory



ตัวอย่าง 1M → 0x100000000 → base → แล้วคือ load code

base

base + bound

ตัวอย่าง 10 → 10 + base → security

* วิ่งบนโปรเซสเซอร์ได้ทันที ใน process ขึ้นอยู่กับ

(↳ แตกกัน → 4GB heap stack data code

Virtually Addressed Base and Bounds

- Pros?

- Simple
- Fast (2 registers, adder, comparator)
- Safe
- Can relocate in physical memory without changing process

- Cons?

- Can't keep program from accidentally overwriting its own code
- Can't share code/data with other processes
- Can't grow stack/heap as needed

↓ ถ้าพบ pointer 1 ข้อ → 1 ข้อจบ

↓ แปลงเป็นค่าจริง/บรรทัด



↓ ถ้าพบ 6 ใช้บรรทัดได้
↓ ถ้าได้ base, bound

↓ ถ้าพบ defragment ของมันก็จะ
หาพื้นที่ใหม่

↓ แล้ว copy อีก

↓ ใช้ของ ram มันก็จำกัดแหละ

↓
↓ ถ้าพบค่าจริงไม่ได้ → ขยายไม่ได้
↓ แปลงก่อนแล้วค่อย compile
↓ 1 program > 1 window
↓ ขยายไม่ได้ → ใช้รับพจนานุกรม
↓ ใช้ของระบบเวลาทำงาน

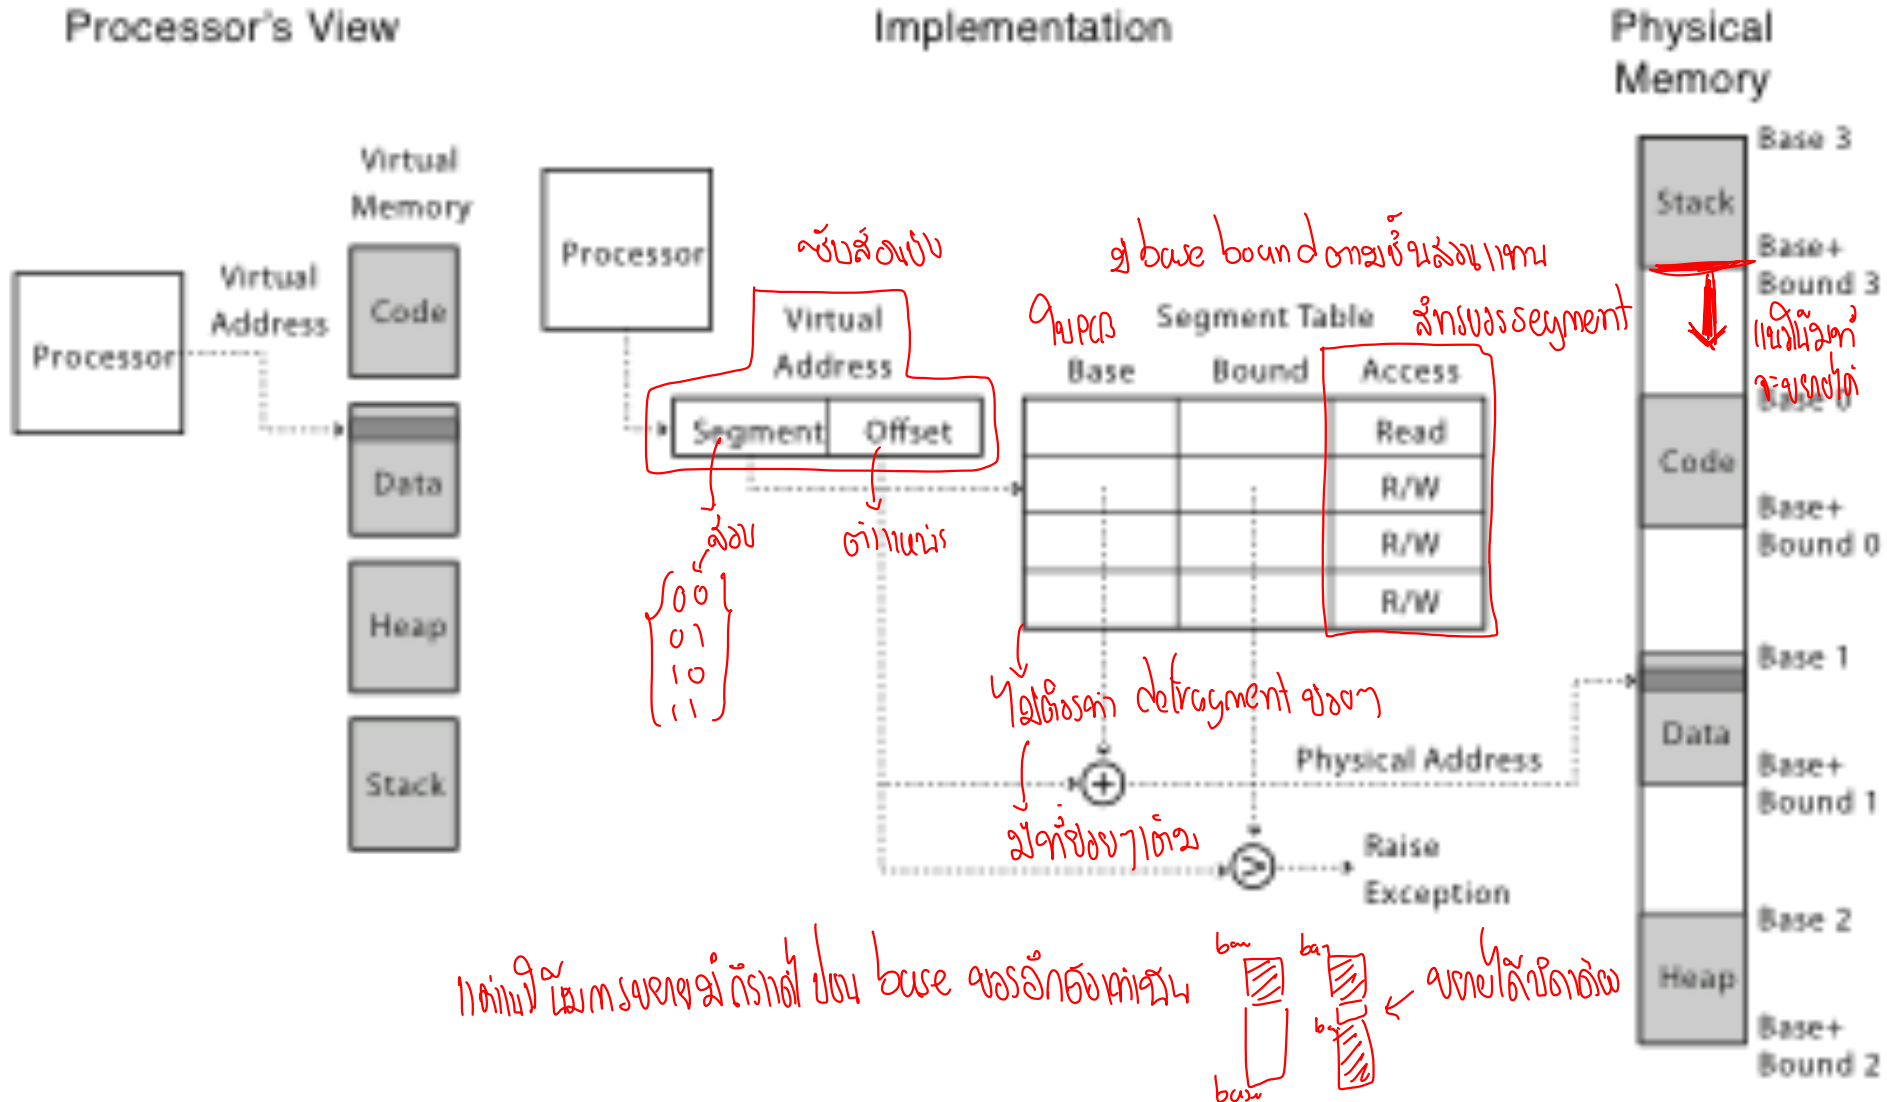
การแบ่ง = 11 บิต ข้อจำกัดในแบบ → physical memory address

↓ 128 boundary 0-184467

Segmentation

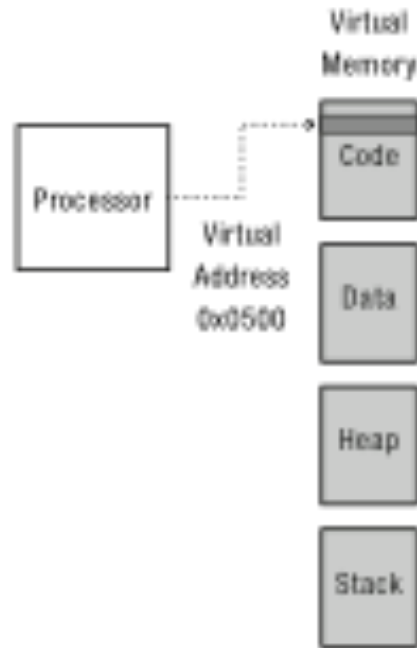
- Segment is a contiguous region of *virtual* memory
- Each process has a segment table (in hardware)
 - Entry in table = segment
- Segment can be located anywhere in physical memory
 - Each segment has: start, length, access permission
- Processes can share segments
 - Same start, length, same/different access permissions

Segmentation



Processor's View

Process 1's View

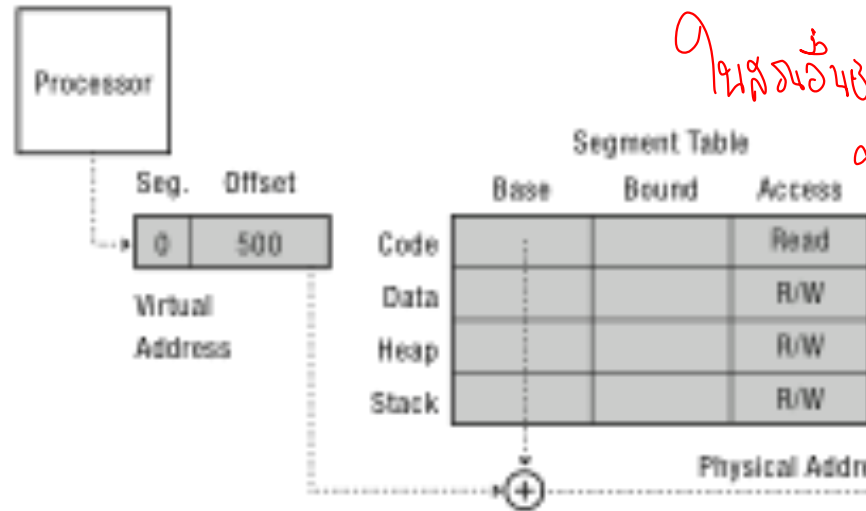


Process 2's View

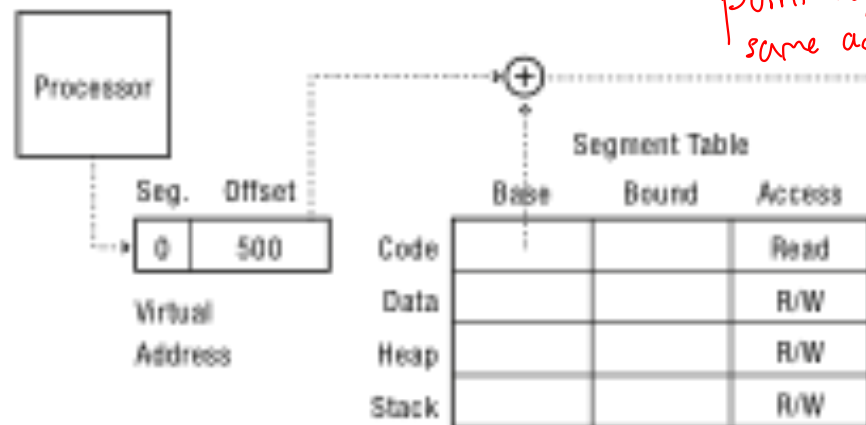


Implementation

การ code sharing ไม่



ในสว่านขุดดิน
ของขี้



point to the
same add

Physical Memory



Segmentation

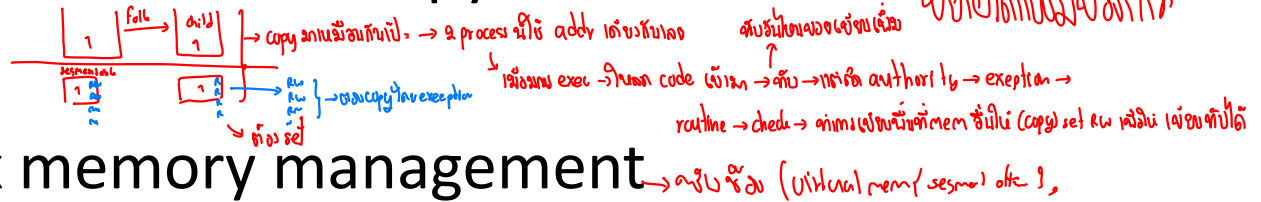
- Cons?

- Complex memory management

- May need to rearrange memory from time to time to ^{base} make room for new segment or growing segment

- External fragmentation: wasted space between chunks

ຂໍຂອບການ ຈົມ ສາມາດຊຸມທັງໝົດ



2) free memory 100%

การจัดกระจายทรัพยากร

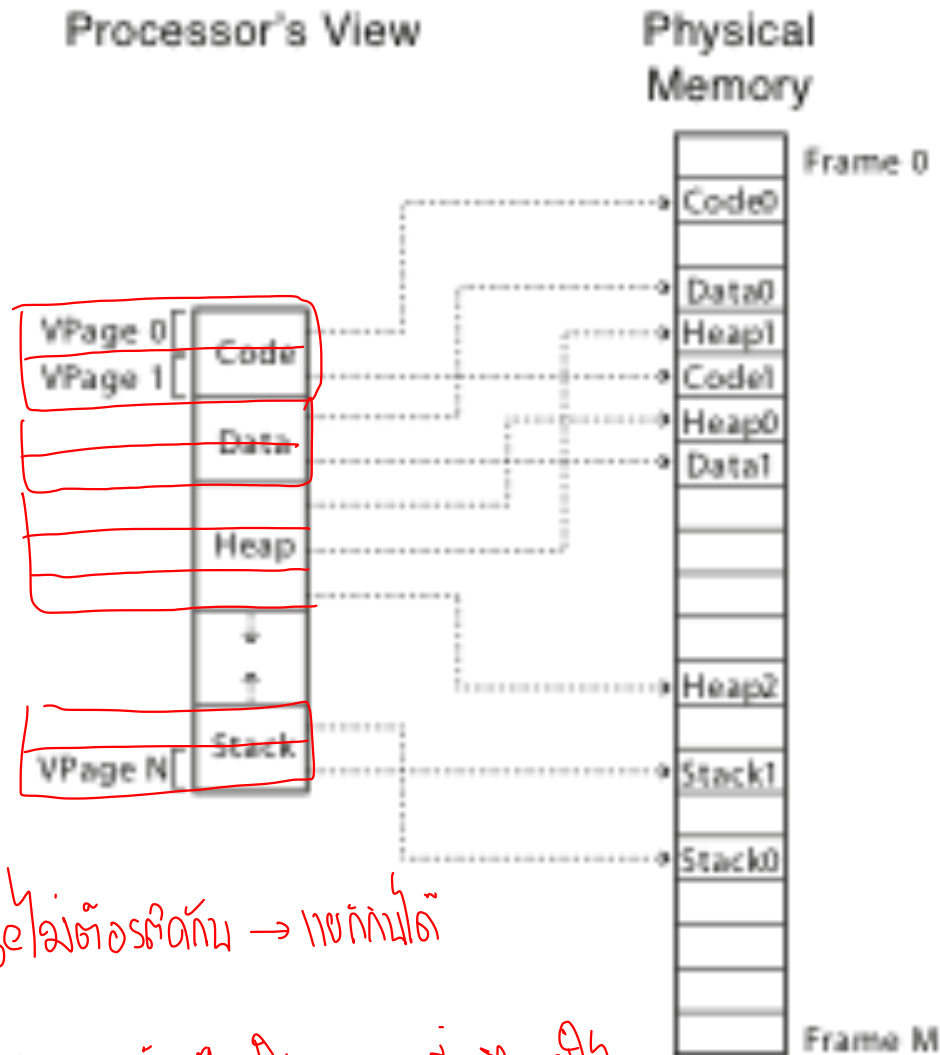
1290s base and boom

Giosani defragment (anwgrn)

120 ram 100 →
 80 ได้ 100 100 100 100 → 100 100 100 100
Paged Translation
 → no defragmentation

- Manage memory in fixed size units, or pages
- Finding a free page is easy
 - Bitmap allocation: 00111111000000001100
 - ↳ 00111111: physical (frame) only
 - 00000000: virtual (page) only
 - 1100: page slot
 - Each bit represents one physical page frame
- Each process has its own page table
 - Stored in physical memory
 - Hardware registers
 - pointer to page table start
 - page table length → address up to virtual mem

Paged Translation (Abstract)



1. 100 page 2. 100 3. 100 → 100

↓
 ๓๖๖ แล้วจึงทำใน page ที่เหลือ อีก

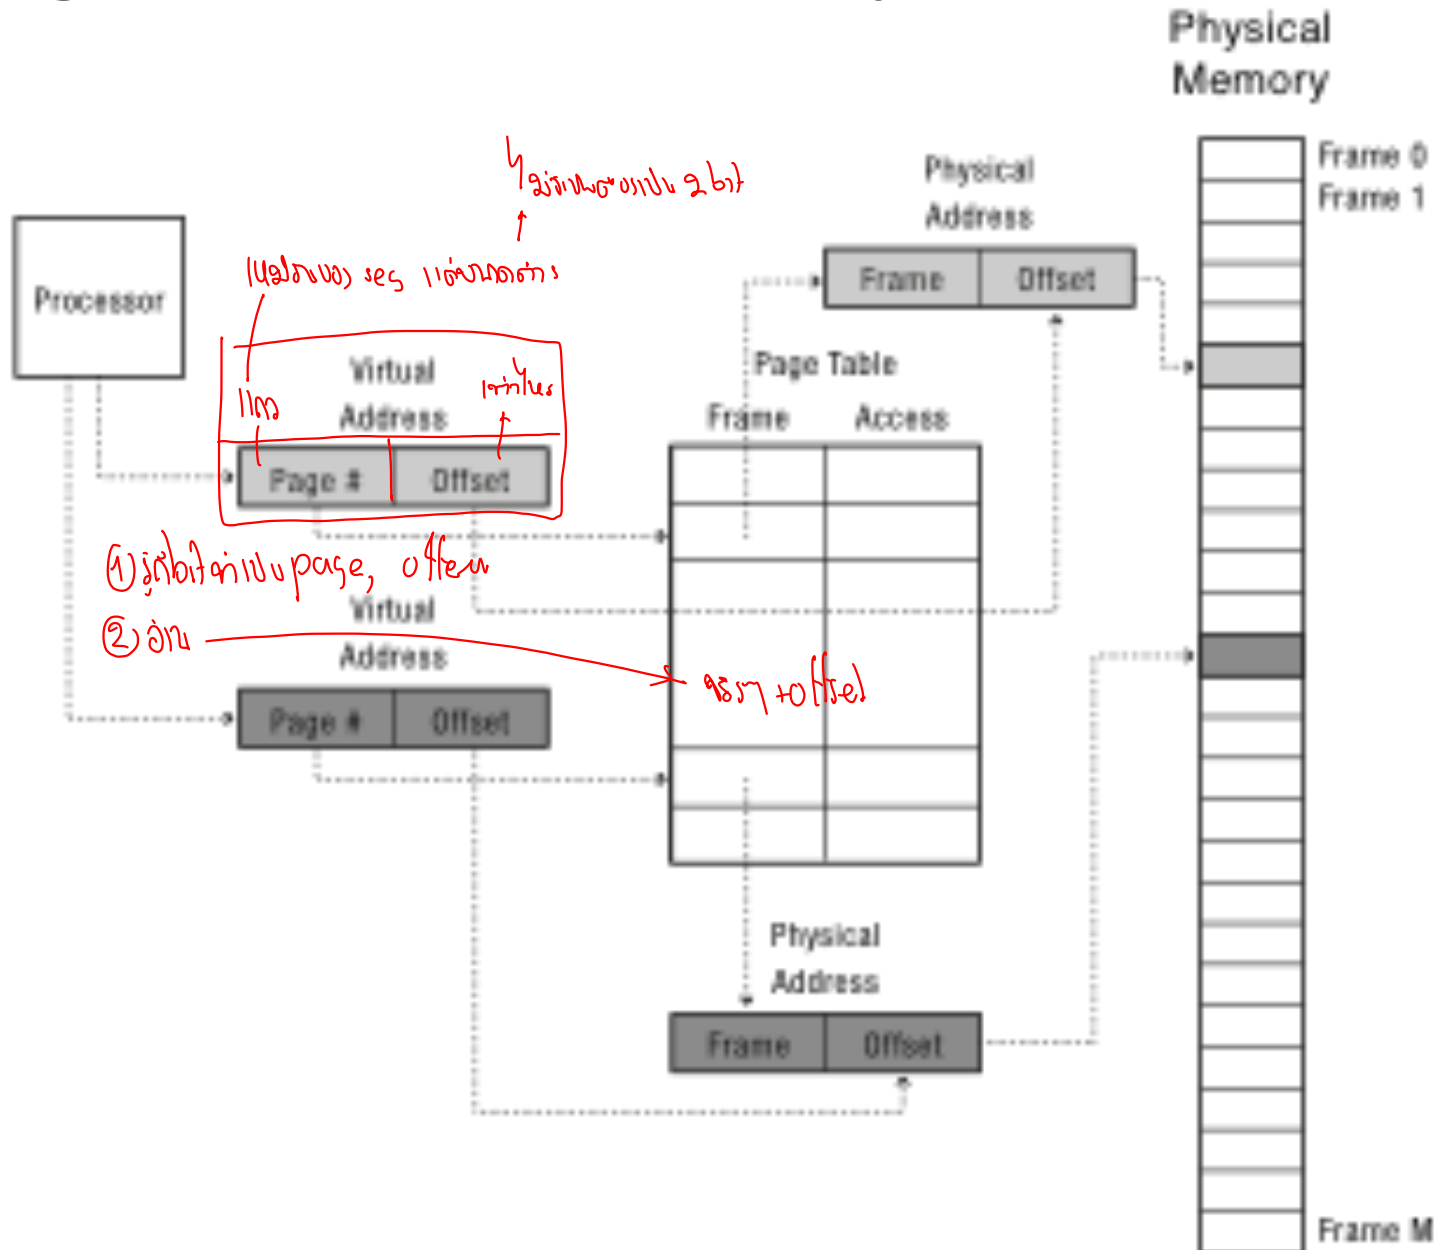
↳ Majoritari memoria ett gamm b₄ b₅, sesun or

အိတ်ကပ်ကပ်

၁၇၆၁၇၇၇၇

19/10/20

Paged Translation (Implementation)



Process View

A
B
C
D
E
F
G
H
I
J
K
L

Physical Memory

I
J
K
L
E
F
G
H
A
B
C
D

Page Table

4	
3	
1	

① frame number
② offset
offset = translation

Sparse Address Spaces

page fragmentation 96% 100% 100% 100%
จัดสรรส่วนต่างๆ ของ code 96% 100% 100% 100%
share code /
proten ✓
memory x bad

- Might want many separate dynamic segments

- Per-processor heaps
- Per-thread stacks
- Memory-mapped files
- Dynamically linked libraries

- What if virtual address space is large?

- 32-bits, 4KB pages => 500K page table entries
- 64-bits => 4 quadrillion page table entries

new algo ←
cachetables
initial allocation map leaf
cache
new algo 96% 100% 100% 100%
cachetables
initial allocation map leaf
cache
new algo 96% 100% 100% 100%

Multi-level Translation

- Tree of translation tables

- Paged segmentation
- Multi-level page tables
- Multi-level paged segmentation

↓
segmentation

- Fixed-size page as lowest level unit of allocation

- Efficient memory allocation (compared to segments)
- Efficient for sparse addresses (compared to paging)
- Efficient disk transfers (fixed size units)
- Easier to build translation lookaside buffers
- Efficient reverse lookup (from physical -> virtual)
- Variable granularity for protection/sharing

allocation map
→ available lanes, fix size frame

↓
frame size
frame bits

↓
translate address

Paged Segmentation

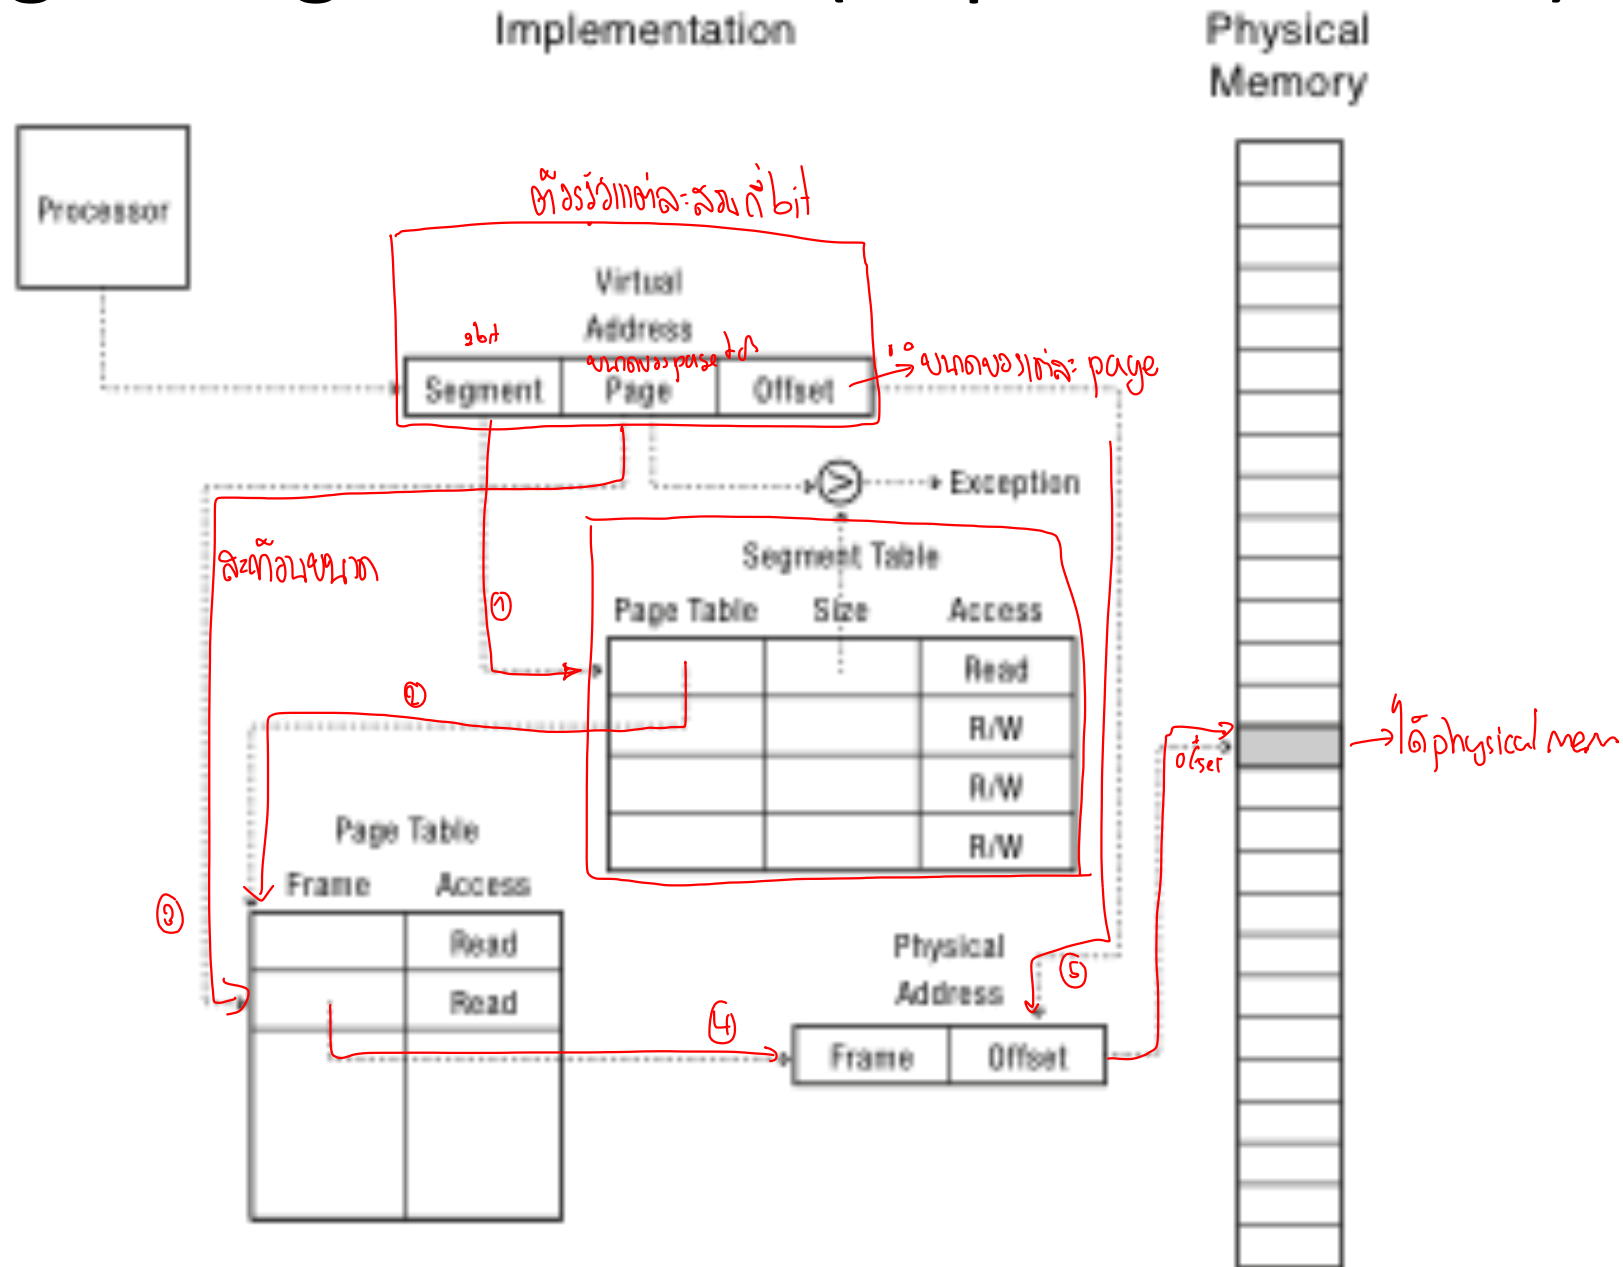
↓ page table + segment table association

- Process memory is segmented
- Segment table entry:
 - Pointer to page table
 - Page table length (# of pages in segment)
 - Access permissions
- Page table entry:
 - Page frame
 - Access permissions
- Share/protection at either page or segment-level

1.4

→ 1 page 1 page

Paged Segmentation (Implementation)

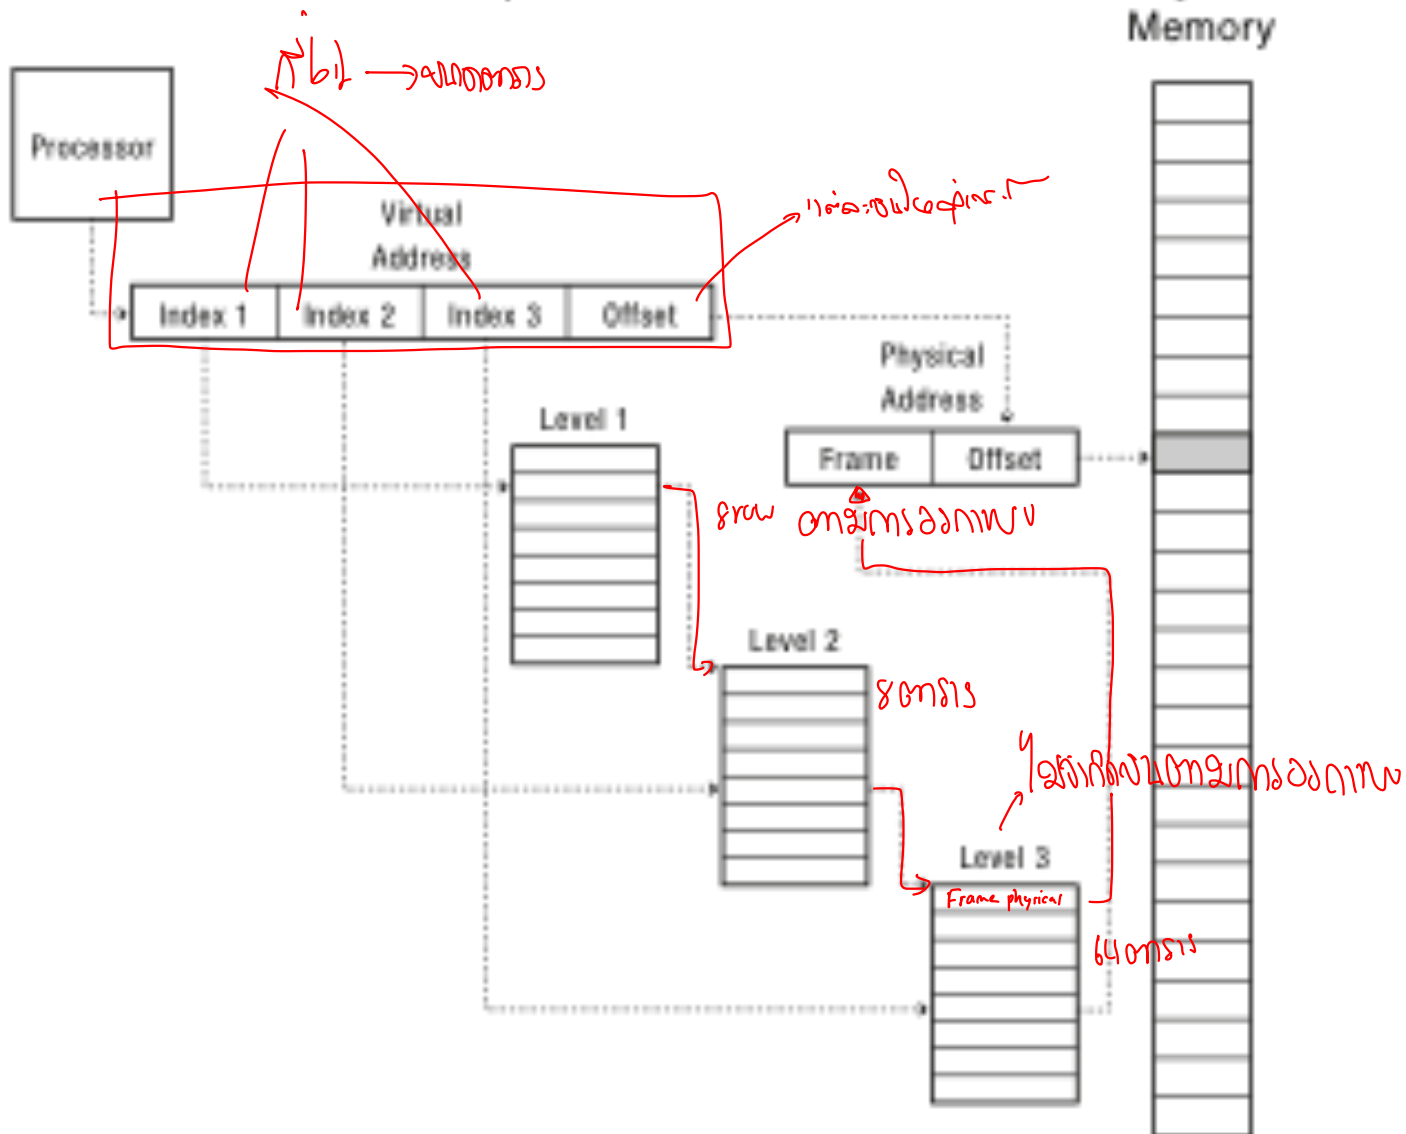


01.58

Multilevel Paging

Implementation

Physical Memory



Multilevel Translation

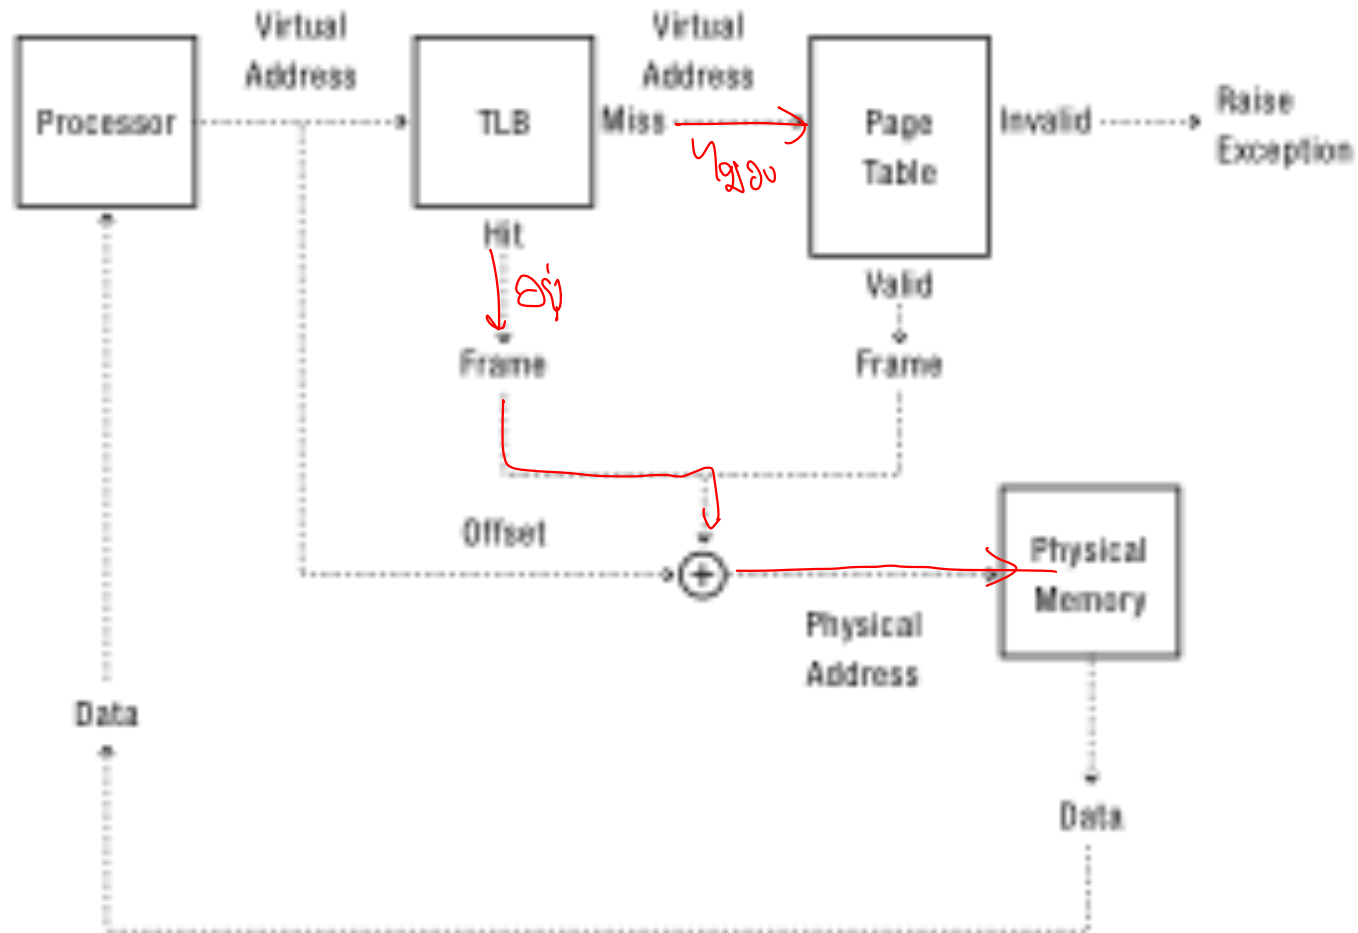
- Pros:
 - Allocate/fill only page table entries that are in use
 - Simple memory allocation
 - Share at segment or page level
- Cons:
 - Space overhead: one pointer per virtual page
 - Two (or more) lookups per memory reference

Efficient Address Translation

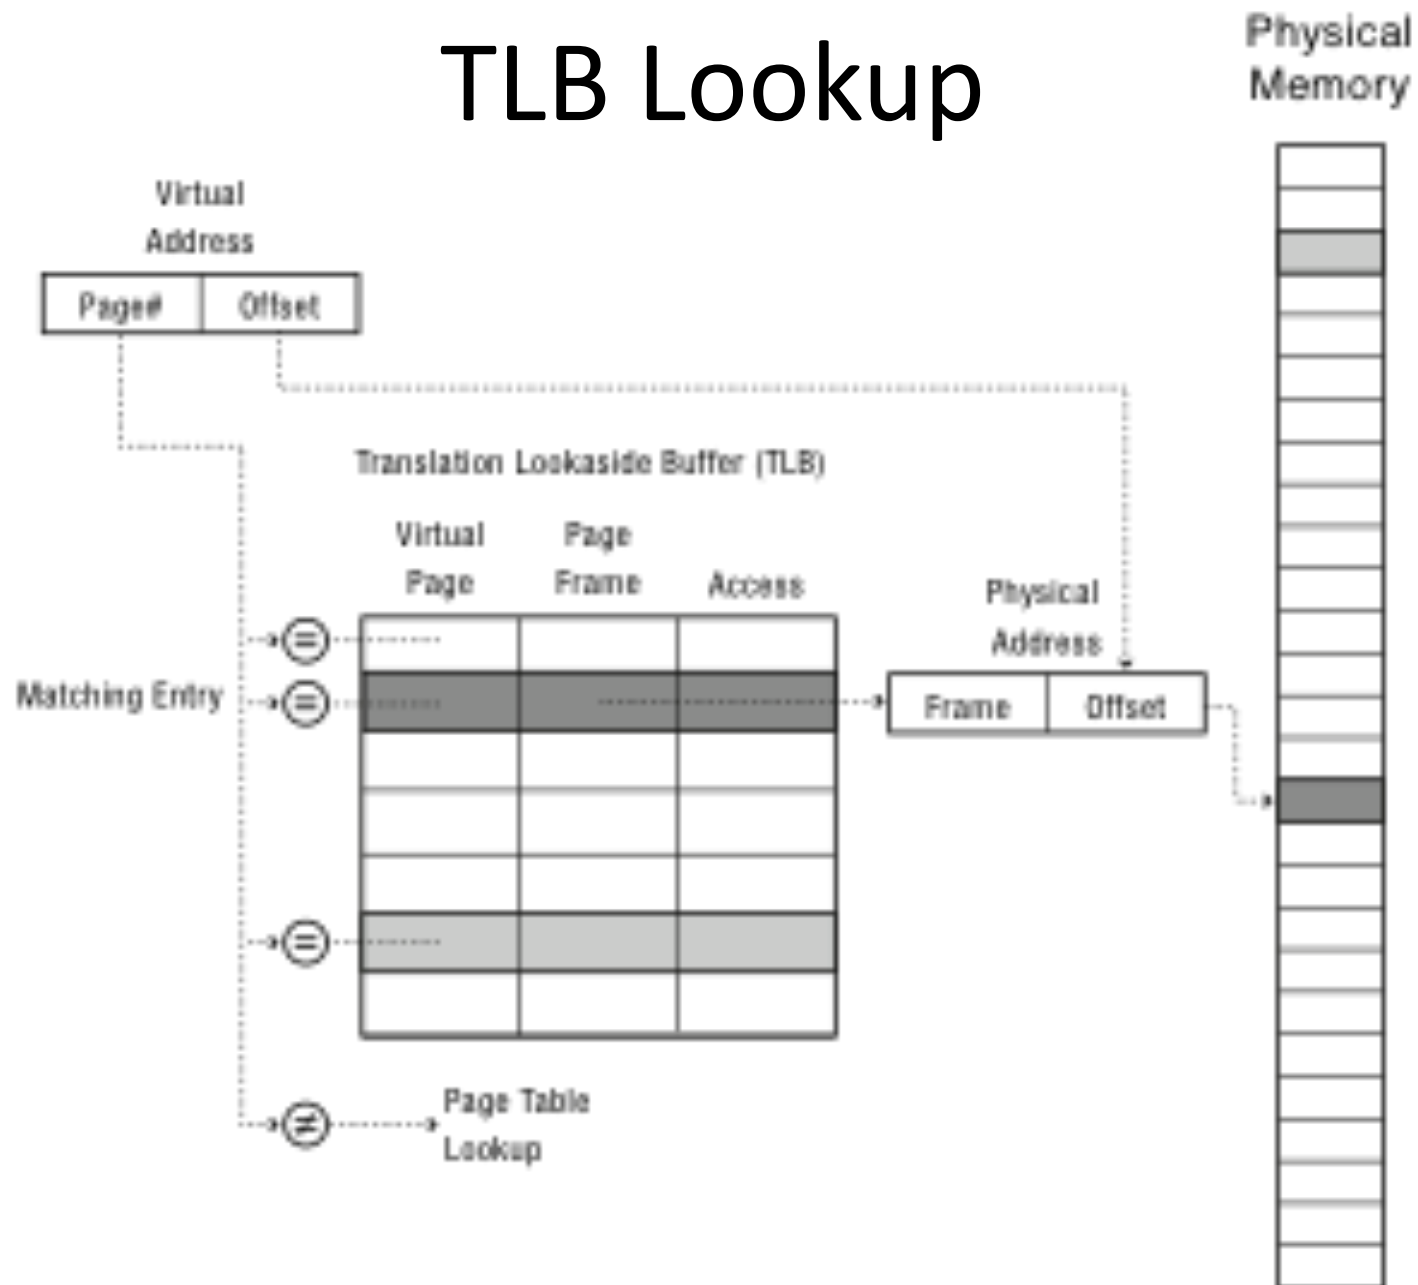
- Translation lookaside buffer (TLB) → *how to optimize translation*
 - Cache of recent virtual page -> physical page translations
↓
จำกัดขนาด
↓
limit size
 - If cache hit, use translation
 - If cache miss, walk multi-level page table
- Cost of translation =
Cost of TLB lookup +
 $\text{Prob}(\text{TLB miss}) * \text{cost of page table lookup}$

TLB and Page Table Translation

ပြန်ကြည့်ရန် TLB



TLB Lookup



Address Translation Uses

- Process isolation
 - Keep a process from touching anyone else's memory, or the kernel's
- Efficient interprocess communication → segmentation គឺជាអវិជ្ជមាន
↓
ឯកសារ
ឯកសារ
- Shared code segments
 - E.g., common libraries used by many different programs
- Program initialization → ^{memory} page 1, ^{stack} page 2, ប្រើប្រាស់ page 1 ហើយ ប្រើប្រាស់ page 2 ផងដែរ
- Dynamic memory allocation ទទួលបានដោយសេរី
– Allocate and initialize stack/heap pages on demand

Address Translation (more)

- Cache management → *cache*
 - Page coloring
- Program debugging → *process ที่ debug* *share* *debugger*
 - Data breakpoints when address is accessed
- Zero-copy I/O → *directly from file* → *print* *overhead* *buffer*
 - Directly from I/O device into/out of user memory
- Memory mapped files *mapped from hdd to ram* *swap file*
 - Access file data using load/store instructions
- Demand-paged virtual memory
 - Illusion of near-infinite memory, backed by disk or memory on other machines

Address Translation (even more)

- Checkpointing/restart
 - Transparently save a copy of a process, without stopping the program while the save happens
- Persistent data structures
 - Implement data structures that can survive system reboots
- Process migration
 - Transparently move processes between machines
- Information flow control
 - Track what data is being shared externally
- Distributed shared memory
 - Illusion of memory that is shared between machines