

Artificial Intelligence

Instructor: Kietikul Jearanaitanakij

Department of Computer Engineering

King Mongkut's Institute of Technology Ladkrabang


Lecture 8

Learning from examples

- Forms of Learning
- Supervised Learning
- Learning Decision Trees
- Evaluating the Hypothesis

- An agent is learning if it improves its performance on future tasks.
→ វា/វាគឺជា agent ដែល វា បាន រៀន បាន កាន់តែ ប្រសើរ ឡើង
- Why would we want an agent to learn? Why wouldn't the designers just program the agent to do the task?
- There are three main reasons.
→ វា មាន ៣ ហេតុ មូលដ្ឋាន ៤
- **First**, the designers cannot anticipate **all possible situations** that the agent might find itself in. For example, a robot designed to navigate mazes must learn the layout of each new maze it encounters.
- **Second**, the designers cannot anticipate **all changes over time**; a program designed to predict tomorrow's stock market prices must learn to adapt when conditions change from boom to bust.
→ វា មាន ហេតុ មូលដ្ឋាន ៥
- **Third**, sometimes human programmers have **no idea how to program** a solution themselves. For example, most people are good at recognizing the faces of family members, but even the best programmers are unable to program a computer to accomplish that task, except by using learning algorithms.
→ គេ មិន មាន algorithm ឡើយ → គេ ត្រូវ ប្រើ ការ រៀន

Forms of learning



1. Unsupervised learning

→ ฝึกให้รู้

- The agent learns patterns in the input even though no explicit label is supplied.
- The most common unsupervised learning task is clustering.
- For example, a taxi agent might gradually develop a concept of “Good traffic days” and “Bad traffic days” without ever being given labeled examples of each by a teacher.

- X1: Rain, BigEvent, WorkDay, MonthEnd
- X2: NonRain, NonBigEvent, NonWorkDay, NonMonthEnd
- X3: Rain, NonBigEvent, WorkDay, NonMonthEnd
- X4: NonRain, BigEvent, NonWorkDay, MonthEnd

A taxi agent may cluster X1, X4 as one cluster and X2, X3 as another cluster.

ฝึกให้รู้

↓ ฝึกให้รู้ input ว่างๆ

↓ ฝึกให้รู้ input ว่างๆ

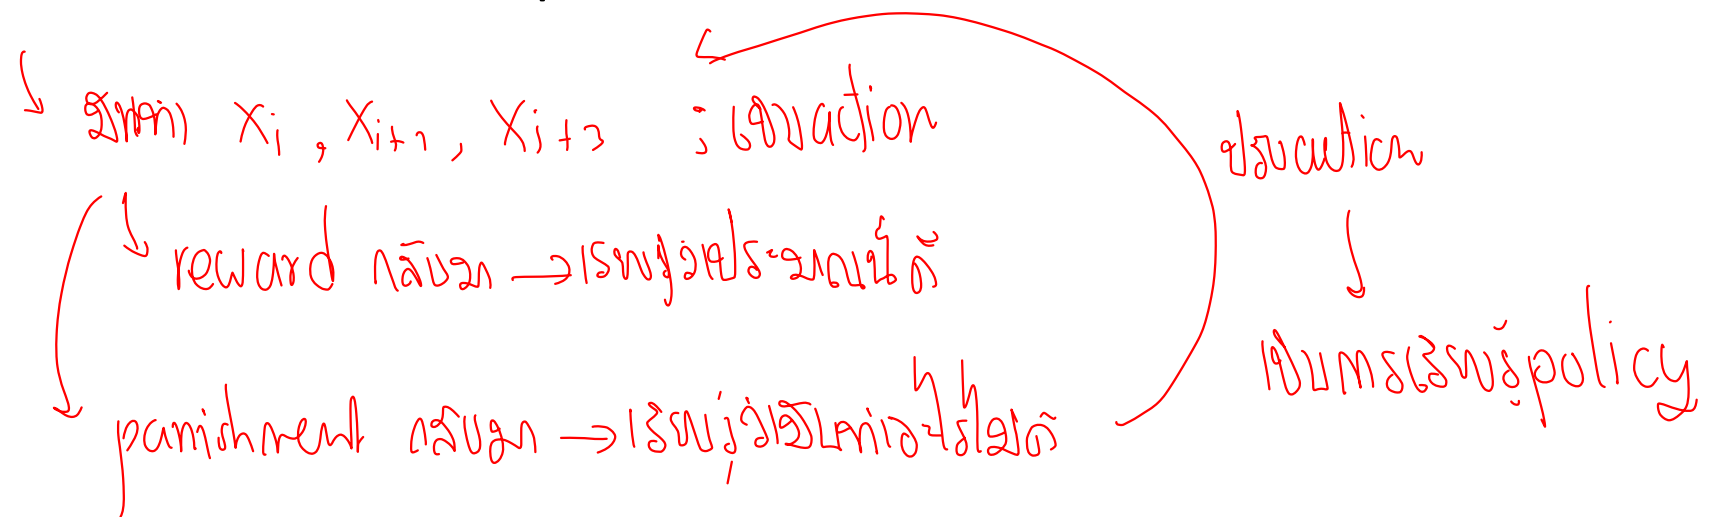
↓ ฝึกให้รู้ input ว่างๆ
good or bad

2. Reinforcement learning

→ ការកំណត់រំលឹកខ្លួន → តែងតែរំលឹកខ្លួន ឬ ឈាត់

per loop get reward or punishment

- The agent learns from a series of reinforcements – rewards or punishments.
- For example, the lack of a tip at the end of the journey gives the taxi agent an indication that it did something wrong.
- It is up to the agent to decide which of the actions prior to the reinforcement were most responsible for it.



3. Supervised learning \rightarrow \tilde{x} \tilde{y}

- The agent observes some example input-output pairs and learns a function that maps from input to output.
- For example, from a taxi agent problem. By using given labeled “input-output” examples, a taxi agent will learn a function that maps from input to the outputs “Good traffic days” and “Bad traffic days”.

- X1: Rain, BigEvent, WorkDay, MonthEnd, “Bad traffic days”
- X2: NonRain, NonBigEvent, NonWorkDay, NonMonthEnd, “Good traffic days”
- X3: Rain, NonBigEvent, WorkDay, NonMonthEnd, “Good traffic days”
- X4: NonRain, BigEvent, NonWorkDay, MonthEnd, “Bad traffic days”

Supervised Learning

The task of supervised learning is this:

Given a **training set** of N example input–output pairs

$$(x_1, y_1), (x_2, y_2), \dots (x_N, y_N) ,$$

where each y_j was generated by an unknown function $y = f(x)$,
discover a function h that approximates the true function f .

ပြန်လည်လေ့လာမှုဆိုင်ရာ အချက်အလက်များကို အသုံးပြု၍ h ကို f နှင့် ချိန်ညှိခြင်း

- Learning is a search through the space of possible hypotheses for one that will perform well, even on new examples beyond the training set.
- To measure the accuracy of a hypothesis we give it a test set of examples that are distinct from the training set.
- We say a hypothesis **generalizes** well if it correctly predicts the value of y for novel examples.

↓ အသစ်အဆန်း အချက်အလက်များအပေါ်တွင် ချိန်ညှိမှုမရှိဘဲ ချောမွေ့စွာ ခန့်မှန်းနိုင်ခြင်း

- When the output **y** is one of a **finite set of values** (such as sunny, cloudy or rainy), the learning problem is called **classification**, and is called Boolean or binary classification if there are only two values.
- When **y** is a **real number** (such as tomorrow's temperature), the learning problem is called **regression**.

Overfitting

- Overfitting occurs when a statistical learning model describes random errors or noises instead of the underlying relationship.
 - A learning model that has been **overfit** will generally have poor predictive performance, as it can exaggerate minor fluctuations in the data.
 - A good learning model should be **generalize** to both unseen and training data.

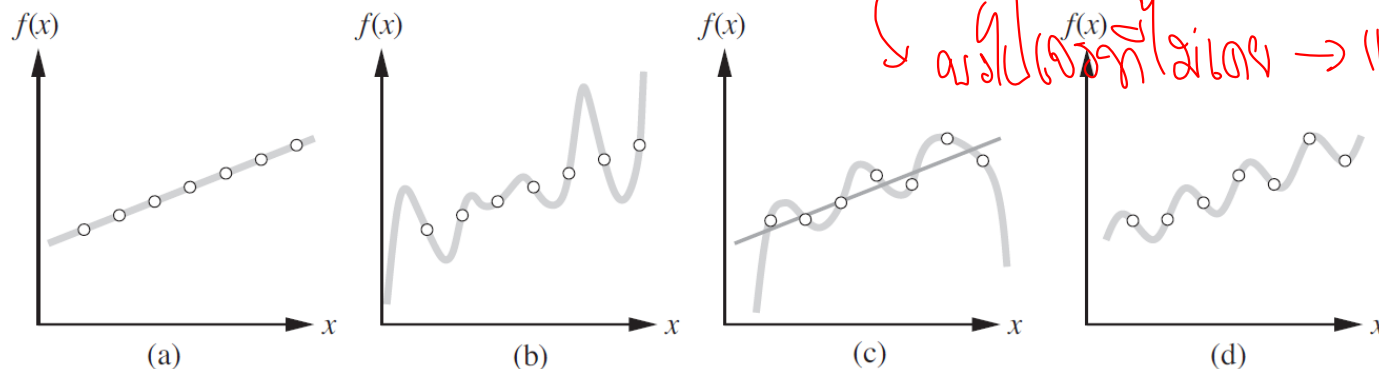


Fig. a : some data with an exact fit by a straight line (the polynomial $0.4x + 3$)

Fig. b : a high-degree polynomial that is also consistent with the same data.

Fig. c : A straight line that is not consistent with any of the data points, but might generalize fairly well for unseen values of x .

Fig. d : data in can be fitted exactly by a function of the form $ax + b + c \sin(x)$. \Rightarrow tend to overfitting.

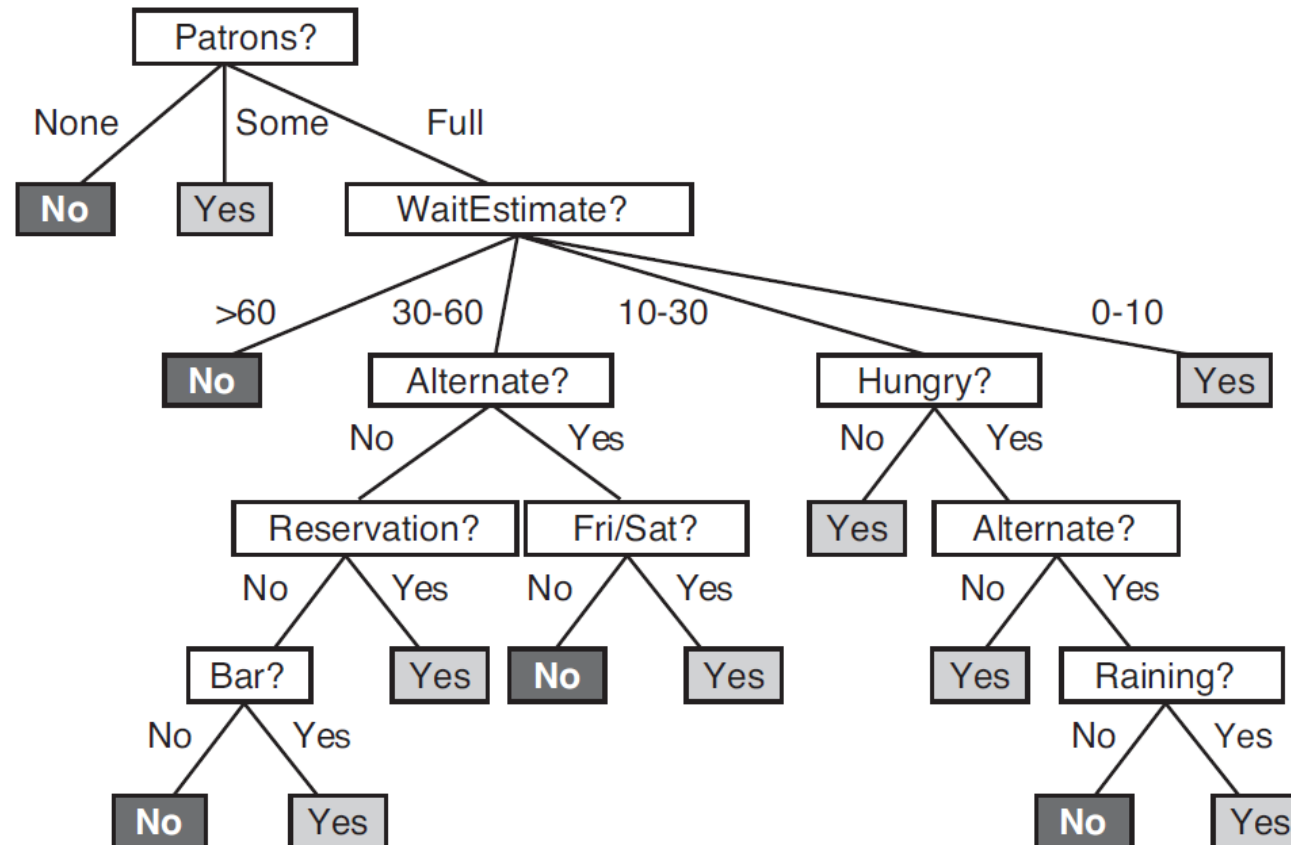
Learning Decision Trees (ID3 algorithm)

- A decision tree represents a function that takes as input a vector of attribute values and returns a “decision”—a single output value.
- As an example, we will build a decision tree to decide whether to wait for a table at a restaurant. The aim here is to learn a definition for the goal predicate `WillWait`.
- First we list the attributes that we will consider as part of the input
 1. **Alternate**: whether there is a suitable alternative restaurant nearby.
 2. **Bar** : whether the restaurant has a comfortable bar area to wait in.
 3. **Fri/Sat**: true on Fridays and Saturdays.
 4. **Hungry**: whether we are hungry.
 5. **Patrons**: how many people are in the restaurant (values are None, Some, and Full).
 6. **Price**: the restaurant’s price range (\$, \$\$, \$\$\$).
 7. **Raining**: whether it is raining outside.
 8. **Reservation**: whether there is a reservation.
 9. **Type**: the kind of restaurant (French, Italian, Thai, or burger).
 10. **WaitEstimate**: the wait estimated by the host (0–10 minutes, 10–30, 30–60, or >60).

Examples for the restaurant domain.

Example	Input Attributes										Goal
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
\mathbf{x}_1	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>0–10</i>	$y_1 = \text{Yes}$
\mathbf{x}_2	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>30–60</i>	$y_2 = \text{No}$
\mathbf{x}_3	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Some</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	$y_3 = \text{Yes}$
\mathbf{x}_4	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Thai</i>	<i>10–30</i>	$y_4 = \text{Yes}$
\mathbf{x}_5	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>>60</i>	$y_5 = \text{No}$
\mathbf{x}_6	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Italian</i>	<i>0–10</i>	$y_6 = \text{Yes}$
\mathbf{x}_7	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	$y_7 = \text{No}$
\mathbf{x}_8	<i>No</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Thai</i>	<i>0–10</i>	$y_8 = \text{Yes}$
\mathbf{x}_9	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>>60</i>	$y_9 = \text{No}$
\mathbf{x}_{10}	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>Italian</i>	<i>10–30</i>	$y_{10} = \text{No}$
\mathbf{x}_{11}	<i>No</i>	<i>No</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>0–10</i>	$y_{11} = \text{No}$
\mathbf{x}_{12}	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>30–60</i>	$y_{12} = \text{Yes}$

A decision tree for deciding whether to wait for a table.



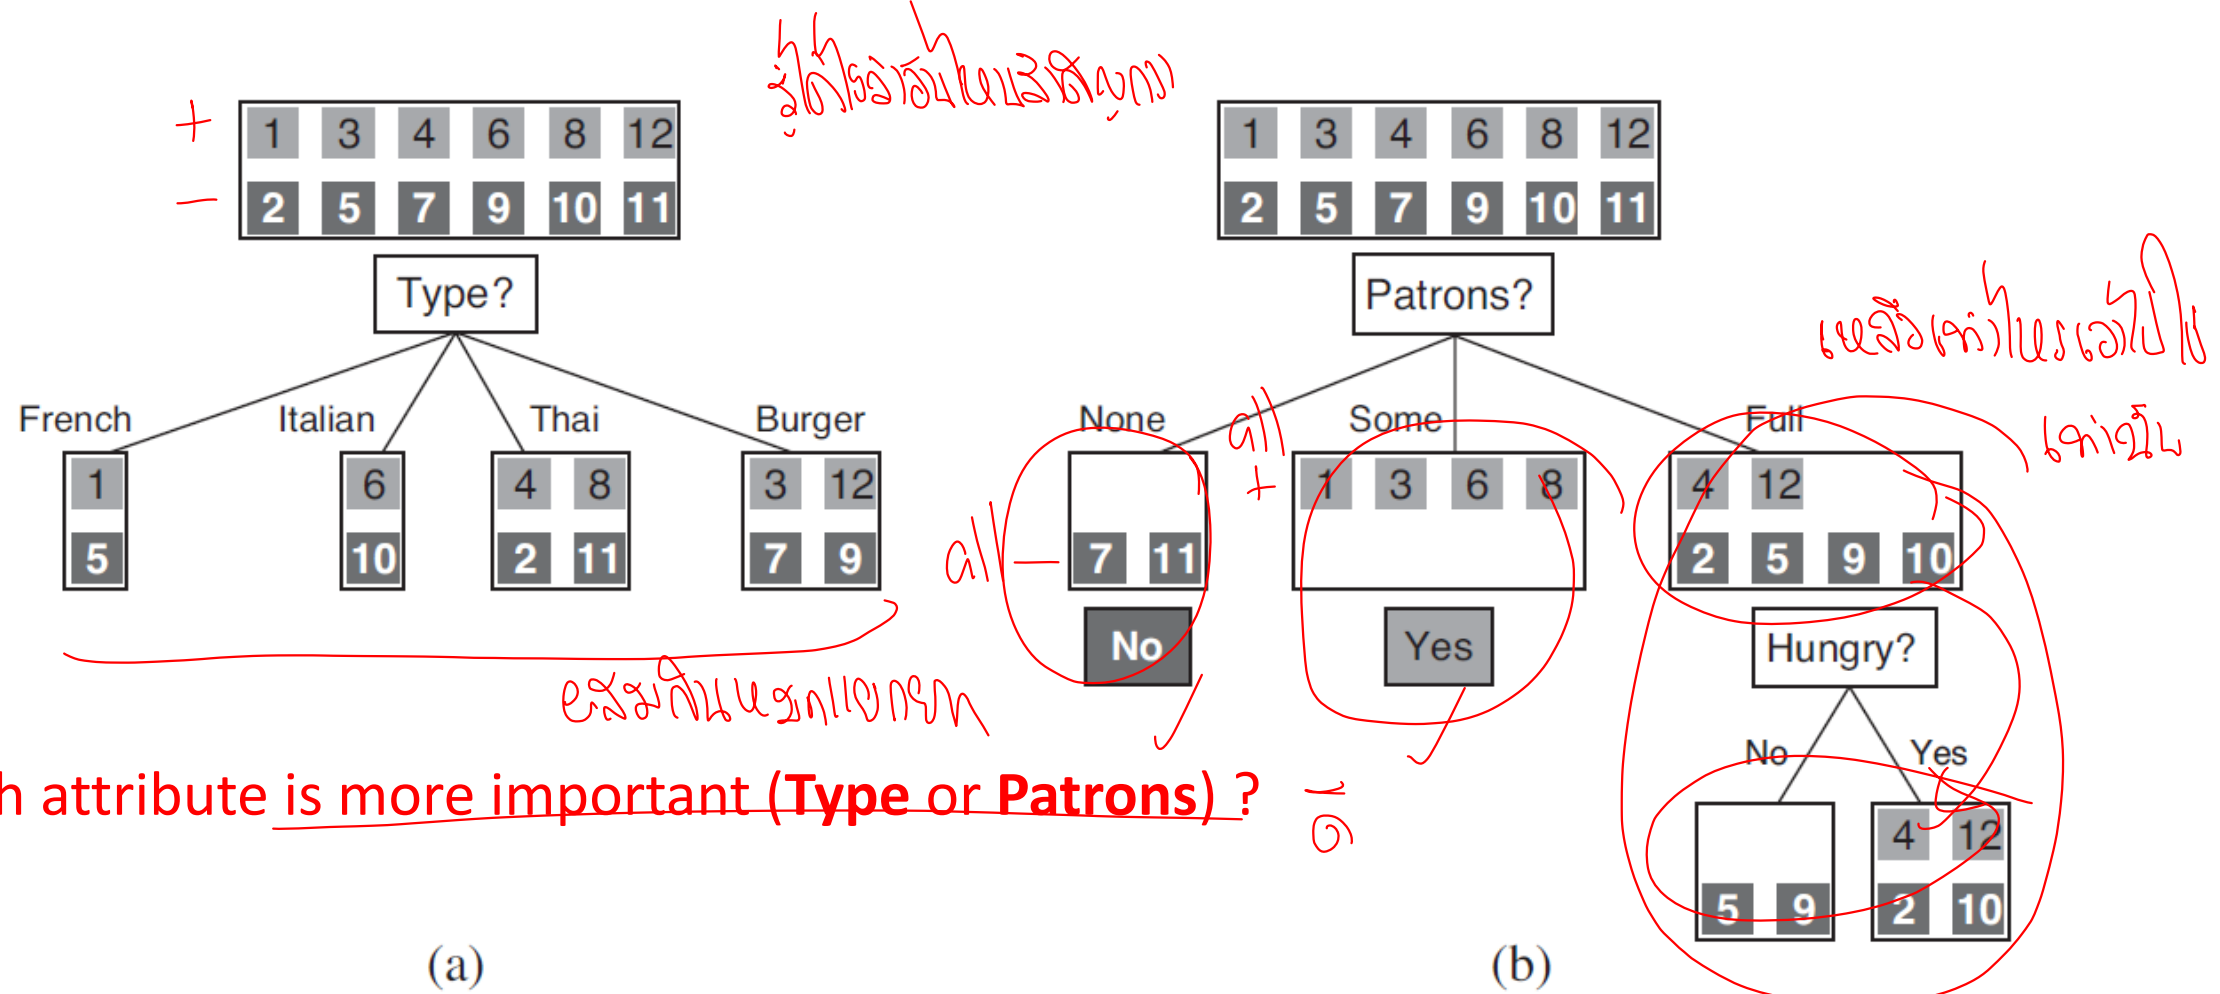
However, this decision tree is not optimal (we will see later).

again

Creating the optimal decision tree

- We hope to get to the correct classification with a small number of tests, meaning that all paths in the tree will be short and the tree as a whole will be shallow. Hence, the optimal decision tree.
- The decision-tree-learning algorithm adopts a greedy divide-and-conquer strategy: always test the most important attribute first.
- By “most important attribute,” we mean the one that makes the most difference to the classification of an example.

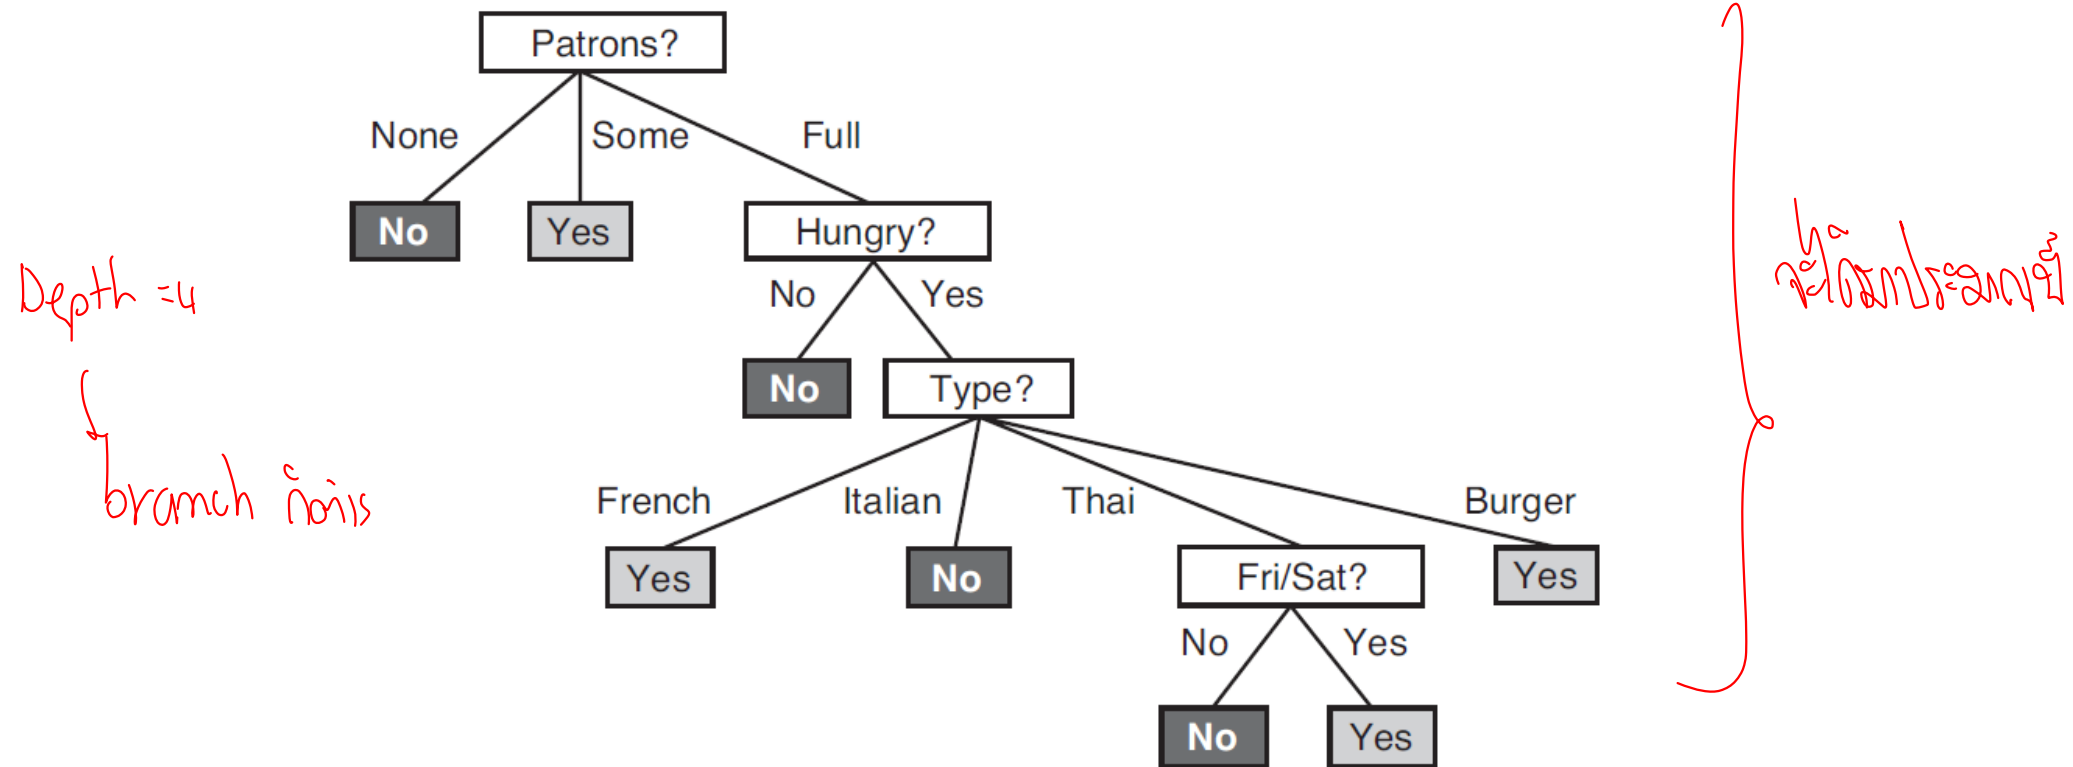
↓ node (150) တစ်ခုတည်းပေးနေတာ



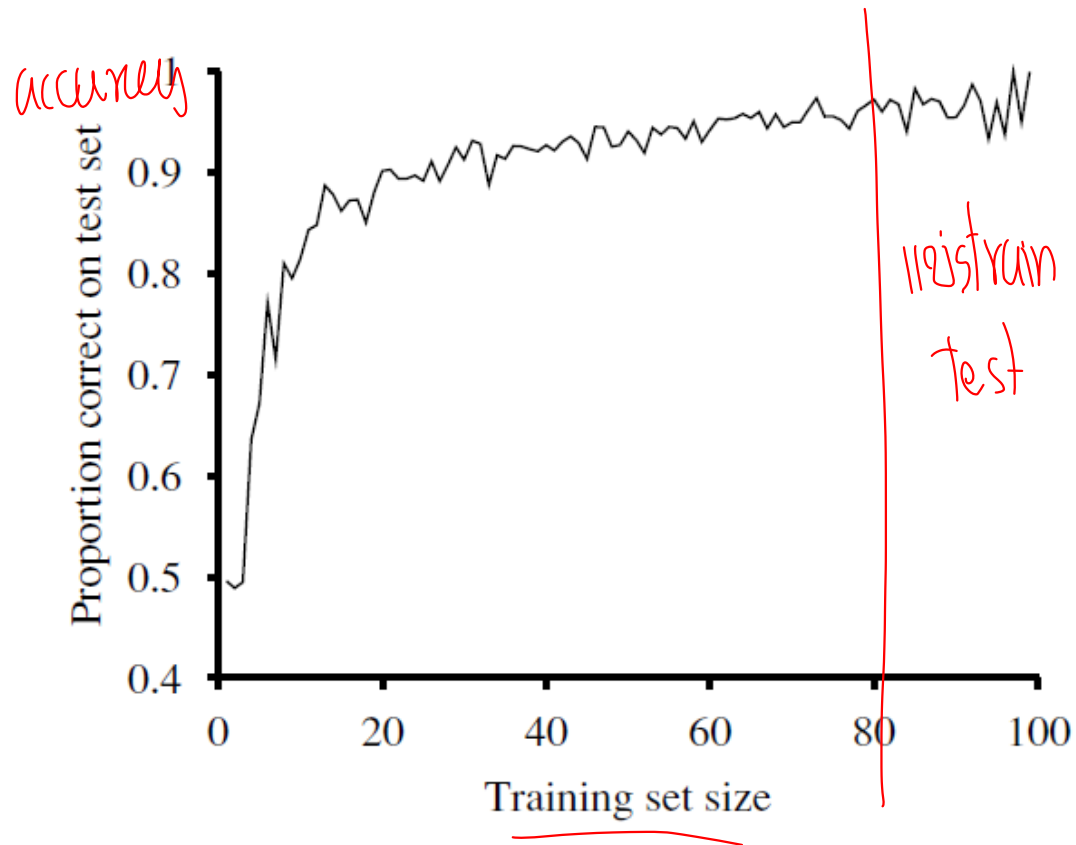
Which attribute is more important (**Type** or **Patrons**) ?

Figure 18.4 Splitting the examples by testing on attributes. At each node we show the positive (light boxes) and negative (dark boxes) examples remaining. (a) Splitting on *Type* brings us no nearer to distinguishing between positive and negative examples. (b) Splitting on *Patrons* does a good job of separating positive and negative examples. After splitting on *Patrons*, *Hungry* is a fairly good second test.

The optimal decision tree induced from the 12-example training set.



- We can evaluate the accuracy of a learning algorithm with a **learning curve**. We have 100 examples at our disposal, which we split into a training set and a test set.



- We learn a hypothesis h with the training set and measure its accuracy with the test set.
- We do this starting with a training set of size 1 and increasing one at a time up to size 99. For each size we actually repeat the process of randomly splitting 20 times, and average the results of the 20 trials.
- The curve shows that as the training set size grows, the accuracy increases.

In this graph, we reach 95% accuracy, and it looks like the curve might continue to increase with more data.

Information gain and Entropy

- A perfect attribute divides the examples into sets, each of which are all positive or all negative and thus will be leaves of the tree.
- The **Patrons** attribute is not perfect, but it is **fairly good**. → မရဘဲမှန်တယ်
- A really **useless attribute**, such as **Type**, leaves the example sets with roughly the same proportion of positive and negative examples as the original set.
- We will use the notion of **information gain**, which is defined in terms of **entropy** (Shannon and Weaver, 1949), to define a formal measure of “**fairly good**” and “**really useless**”
→ လိုအပ်ချက်မရှိဘဲ
- **Entropy** is a measure of the uncertainty of a random variable; acquisition of information corresponds to a reduction in entropy.

(0~1) → အတိအကျမရှိဘဲမှန်တယ် → ပိုမိုအတိအကျသွားတာ → loss ပိုနည်း

- A random variable with only **one value**—a coin that always comes up heads—has no uncertainty and thus its entropy is defined as zero; thus, we gain no information by observing its value.

↓ ၆၀၇၂၀၀ x ၂၀၇၂၀၀ ၂၀၇၂၀၀၇၇-၆၀၇၂၀၀ y → entropy သိရတာပဲ

- If a training set contains **p positive** examples and **n negative** examples, then the entropy of the whole set is

Entropy: $I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) = -\frac{p}{p+n} \log_2 \left(\frac{p}{p+n}\right) - \frac{n}{p+n} \log_2 \left(\frac{n}{p+n}\right)$

information

- We can check that the entropy of a fair coin flip is indeed 1:

$$I\left(\frac{1}{1+1}, \frac{1}{1+1}\right) = -\frac{1}{2} \log_2 \frac{1}{2} - \frac{1}{2} \log_2 \frac{1}{2} = 1.$$

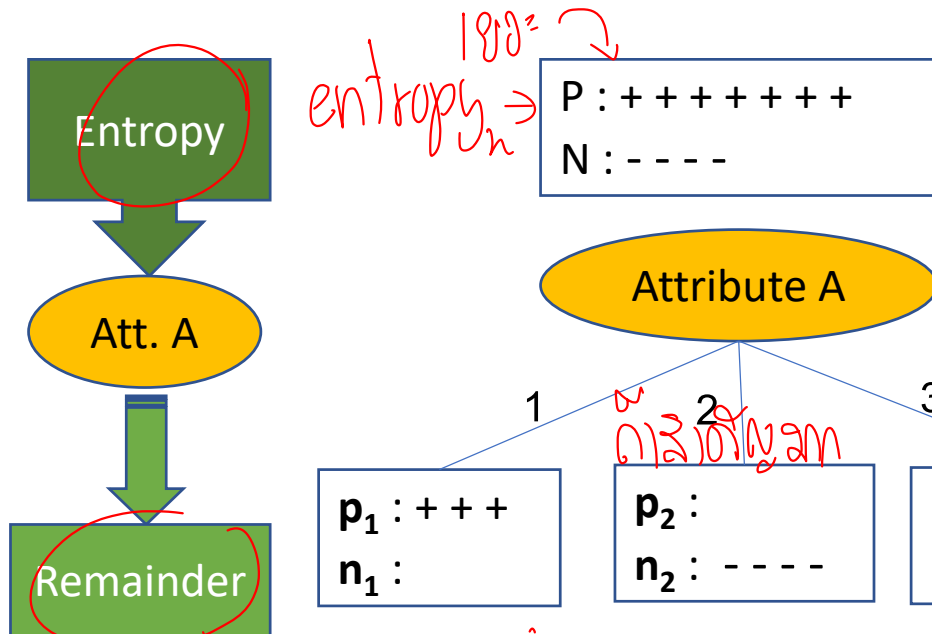
- A flip of a fair coin is equally likely to come up heads or tails, 0 or 1, and this counts as “1 bit” of entropy.

mathematical

- If the coin is loaded to give 99% heads, we get

$$I\left(\frac{99}{99+1}, \frac{1}{99+1}\right) = -\frac{99}{100} \log_2 \frac{99}{100} - \frac{1}{100} \log_2 \frac{1}{100} = 0.08.$$

၆၀၇၂၀၀



- We can measure how much the entropy remaining after testing the attribute A by

$$Remainder(A) = \sum_{i=1}^v \frac{p_i + n_i}{p + n} \cdot I\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$

- The **information gain** from the attribute A is defined as

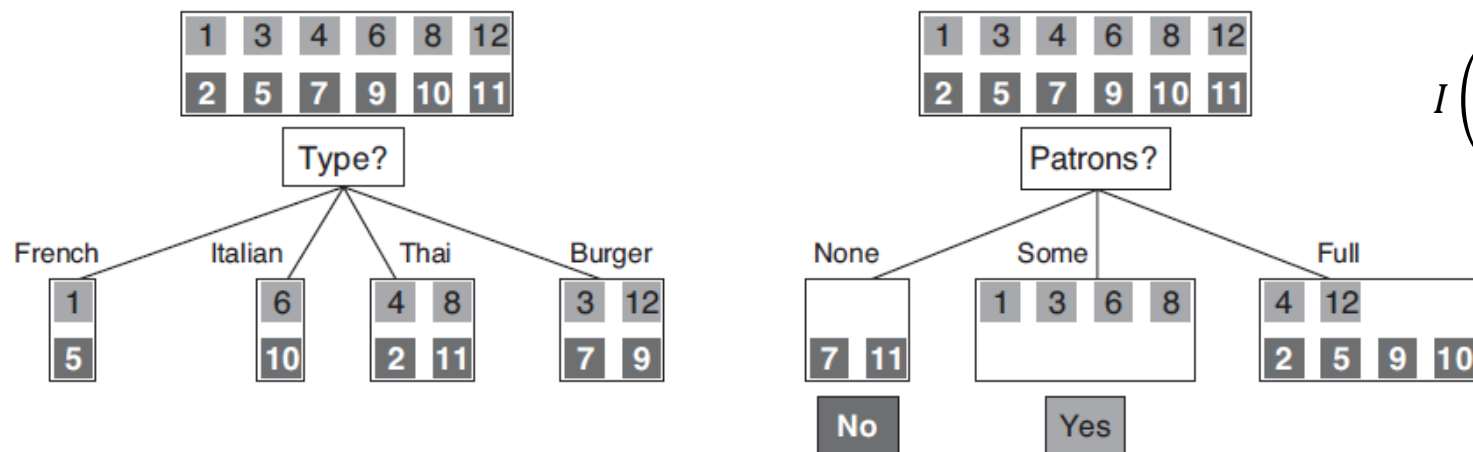
$$Gain(A) = I\left(\frac{p}{p + n}, \frac{n}{p + n}\right) - Remainder(A)$$

(max) (min)

- The more value of the information gain of the attribute A, the more important it is.

Tell us how much information we still need to classify the remaining examples.

- Now, let us calculate the information of both Patrons and Type



$$I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) = -\frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}$$

$$Remainder(A) = \sum_{i=1}^v \frac{p_i + n_i}{p+n} \cdot I\left(\frac{p_i}{p_i+n_i}, \frac{n_i}{p_i+n_i}\right)$$

$$Gain(A) = I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) - Remainder(A)$$

$$Gain(Type) = I\left(\frac{6}{6+6}, \frac{6}{6+6}\right) - \left[\frac{2}{12} \cdot I\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{2}{12} \cdot I\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{4}{12} \cdot I\left(\frac{2}{4}, \frac{2}{4}\right) + \frac{4}{12} \cdot I\left(\frac{2}{4}, \frac{2}{4}\right) \right]$$

= 0 function useless

$$Gain(Patrons) = I\left(\frac{6}{6+6}, \frac{6}{6+6}\right) - \left[\frac{2}{12} \cdot I(0,1) + \frac{4}{12} \cdot I(1,0) + \frac{6}{12} \cdot I\left(\frac{2}{6}, \frac{4}{6}\right) \right]$$

$$= 1 - [0 + 0 + 0.4594]$$

$$= \underline{0.5406}$$

Therefore, Patrons is more important than Type.

↳ Wasserstraßen 1101/500505-02

↓ ឧទាហរណ៍ → ឆ្នាំចេញសៀវភៅ

- Besides ID3 algorithm, there is another simple classification which does not require any training parameter.
- **K nearest neighbors** is a simple algorithm that stores all available examples and classifies new examples based on a similarity measure (e.g., distance functions).
- KNN has been used in statistical estimation and pattern recognition already in the beginning of 1970's as a **non-parametric** technique.

KNN Algorithm

- An example is classified by a **majority vote of its neighbors**.
- An **unseen** example is **assigned to the class most common among its K nearest neighbors** measured by a distance function.
- If **K = 1**, then an **unseen** example is simply assigned to the class of the nearest neighbor, only **one neighbor** which is nearest to an unseen example.

Distance functions

Dimensions

Euclidean

$$\sqrt{\sum_{i=1}^k (x_i - y_i)^2}$$

Manhattan

$$\sum_{i=1}^k |x_i - y_i|$$

Minkowski

$$\left(\sum_{i=1}^k (|x_i - y_i|)^q \right)^{1/q}$$

- It should also be noted that all three distance measures are only valid for continuous variables. In the instance of **categorical variables** the **Hamming distance** must be used.
- Historically, the **optimal K** for most datasets has been between **3-10**. That produces much better results than 1NN.

သိရန်

Hamming Distance

$$x = y \Rightarrow D = 0$$

$$x \neq y \Rightarrow D = 1$$

X	Y	Distance
Male	Male	0
Male	Female	1

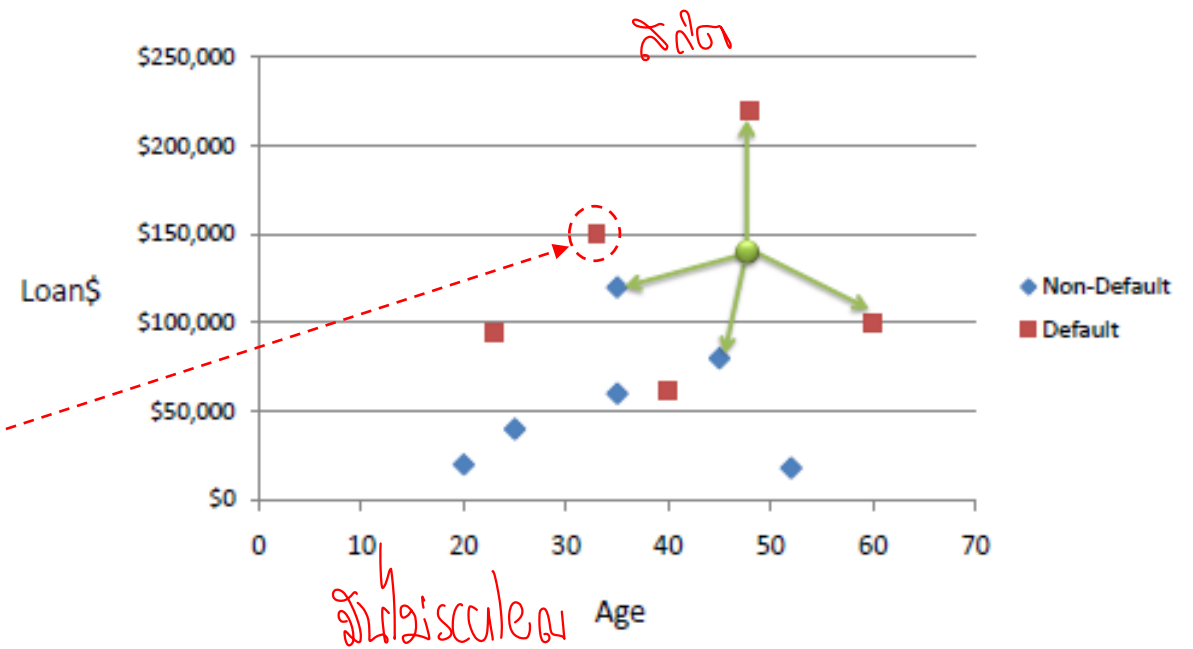
- **Example:** Consider the following data concerning **credit default**. **Age** and **Loan** are two numerical attributes and **Default** is the target.
- We can now use the training set to classify an **unknown case** (Age=48 and Loan=\$142,000) using **Euclidean distance**. If **K=1** then the nearest neighbor is the last case in the training set with **Default=Y**.

$$D = \text{Sqrt}[(48-33)^2 + (142,000-150,000)^2] = 8,000.01 \gg \text{Default=Y}$$

Age	Loan	Default	Distance
25	\$40,000	N	102000
35	\$60,000	N	82000
45	\$80,000	N	62000
20	\$20,000	N	122000
35	\$120,000	N	22000
52	\$18,000	N	124000
23	\$95,000	Y	47000
40	\$62,000	Y	80000
60	\$100,000	Y	42000
48	\$220,000	Y	78000
33	\$150,000	Y	8000
48	\$142,000	?	

Euclidean Distance

$$D = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}$$



- With **K=3**, there are two Default=Y and one Default=N out of three closest neighbors. The prediction for the unknown case is again **Default=Y**.

Standardized Distance

- One major drawback in calculating distance measures directly from the training set is in the case where variables have **different measurement scales** or there is a mixture of **numerical and categorical variables**.
- For example, if one variable is based on annual **income in dollars**, and the other is based on **age in years** then **income will have a much higher influence** on the distance calculated.
- One solution is to **standardize the training set**.
- Using the standardized distance on the same training set, the unknown example returned a different neighbor.

Age	Loan	Default	Distance
0.125	0.11	N	0.7652
0.375	0.21	N	0.5200
0.625	0.31	N	0.3160
0	0.01	N	0.9245
0.375	0.50	N	0.3428
0.8	0.00	N	0.6220
0.075	0.38	Y	0.6669
0.5	0.22	Y	0.4437
1	0.41	Y	0.3650
0.7	1.00	Y	0.3861
0.325	0.65	Y	0.3771
0.7	0.61	?	

Standardized Variable

$$X_s = \frac{X - \text{Min}}{\text{Max} - \text{Min}}$$

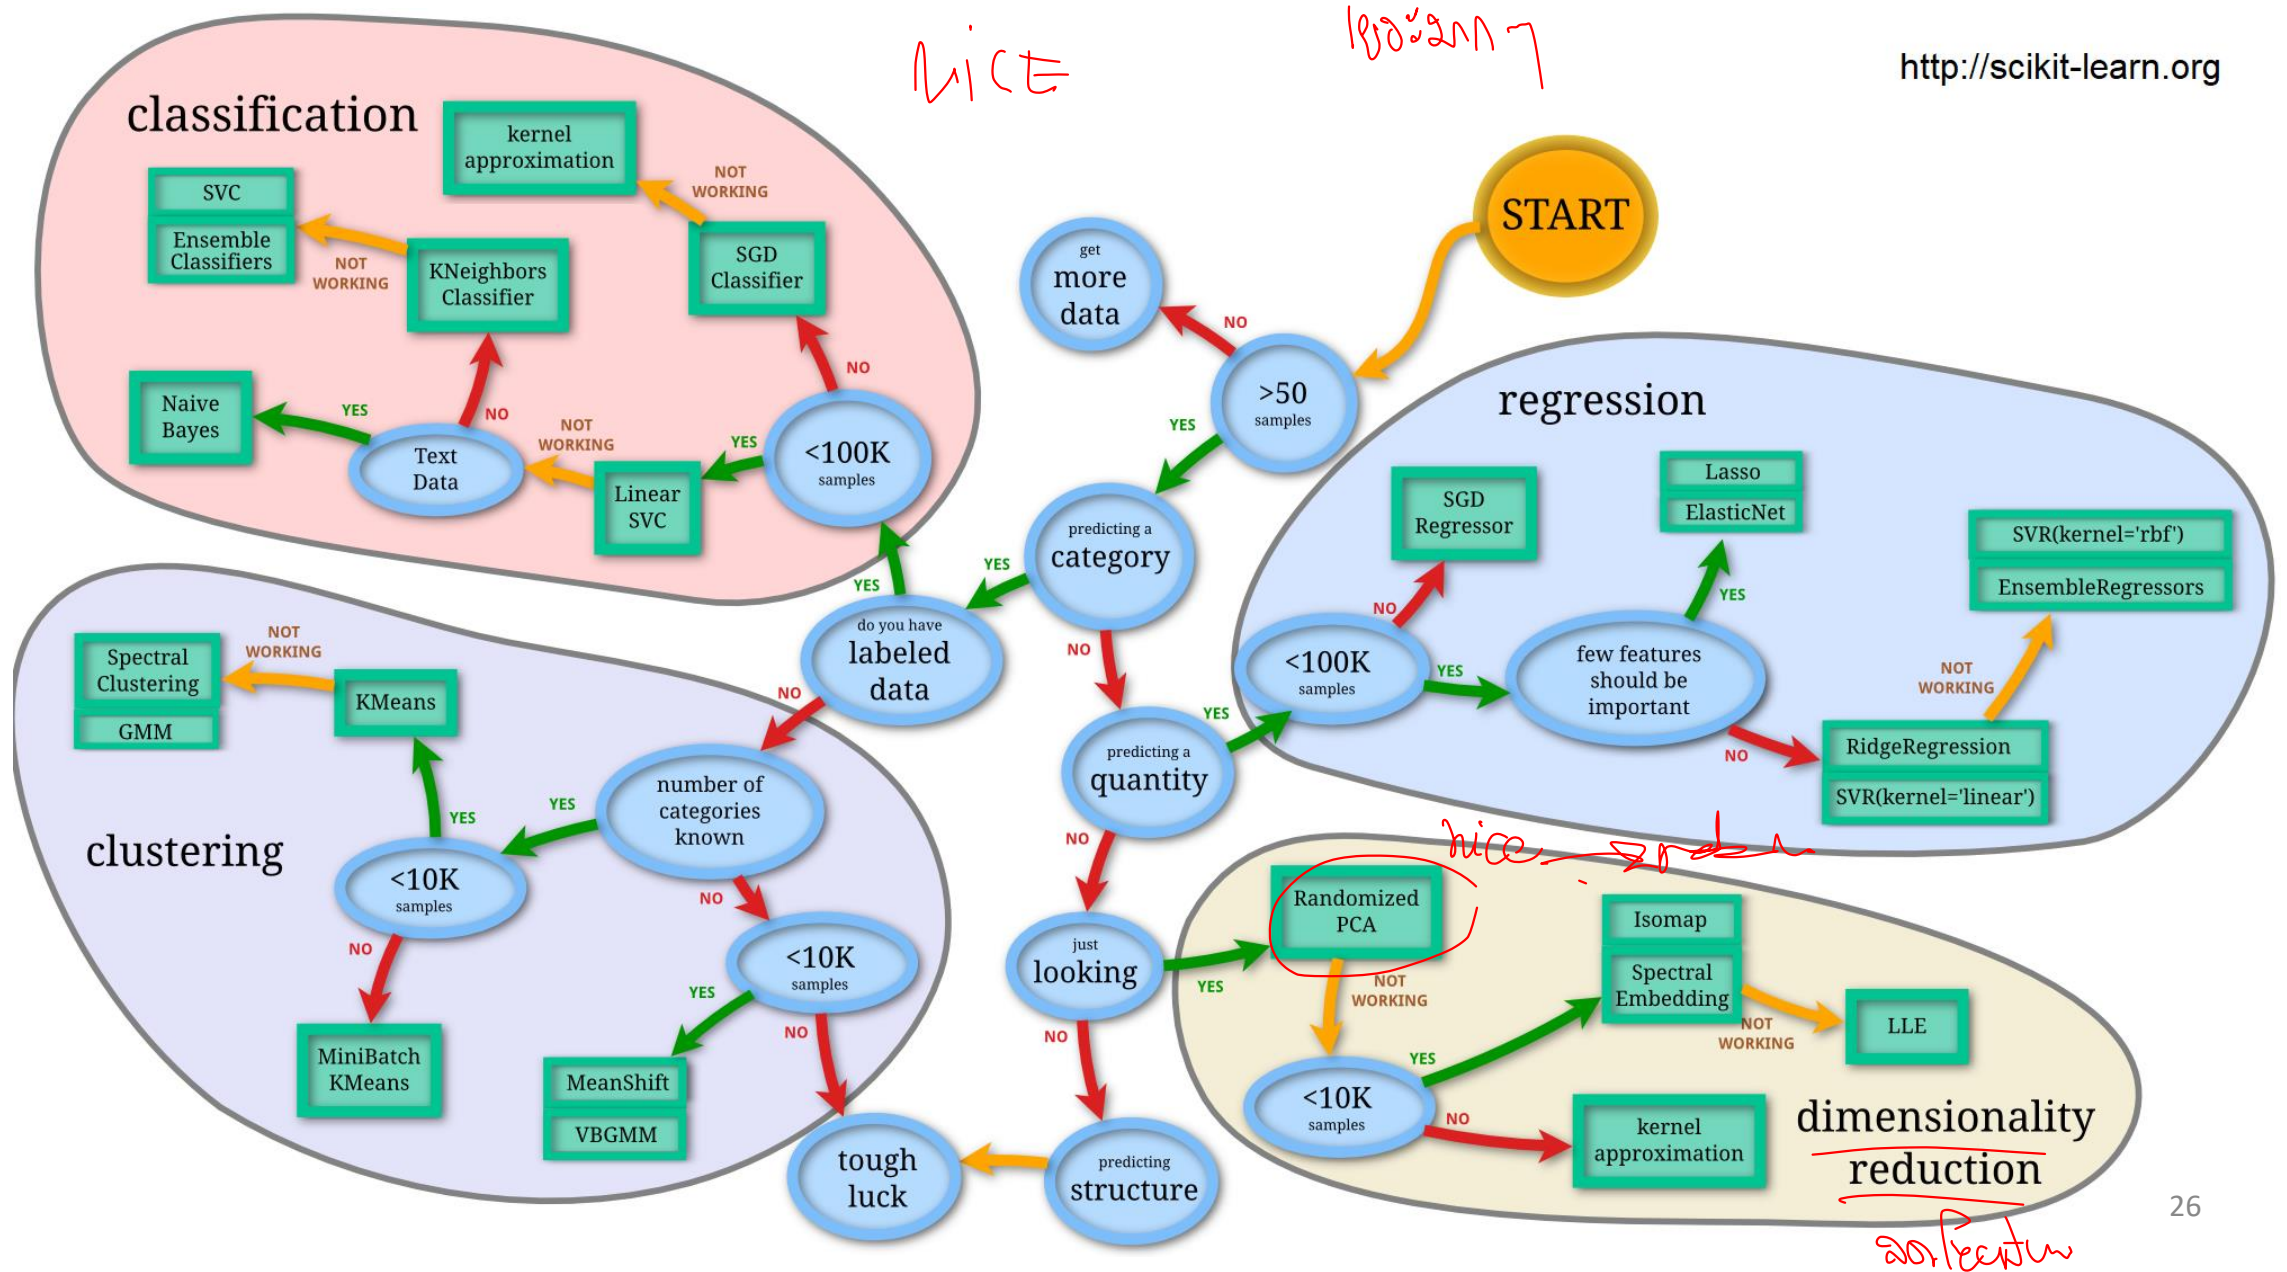
rescale



more no

• Other machine learning algorithms

<http://scikit-learn.org>



Evaluating the hypothesis

- We can simply use the error rate to check the quality of the hypothesis.

100% - accuracy

- Error rate =
$$\frac{\text{Number of examples which are wrong predicted}}{\text{Total number of examples}}$$

- However, a hypothesis h has a low error rate on the training set does not mean that it will generalize well.
- To get the more accurate evaluation, we need to test the hypothesis on an unseen set (test set) of examples.


- **Holdout cross-validation :**


Randomly split the available data into a training set (e.g. 75%) and a test set (e.g. 25%).

- **K-fold cross-validation:**

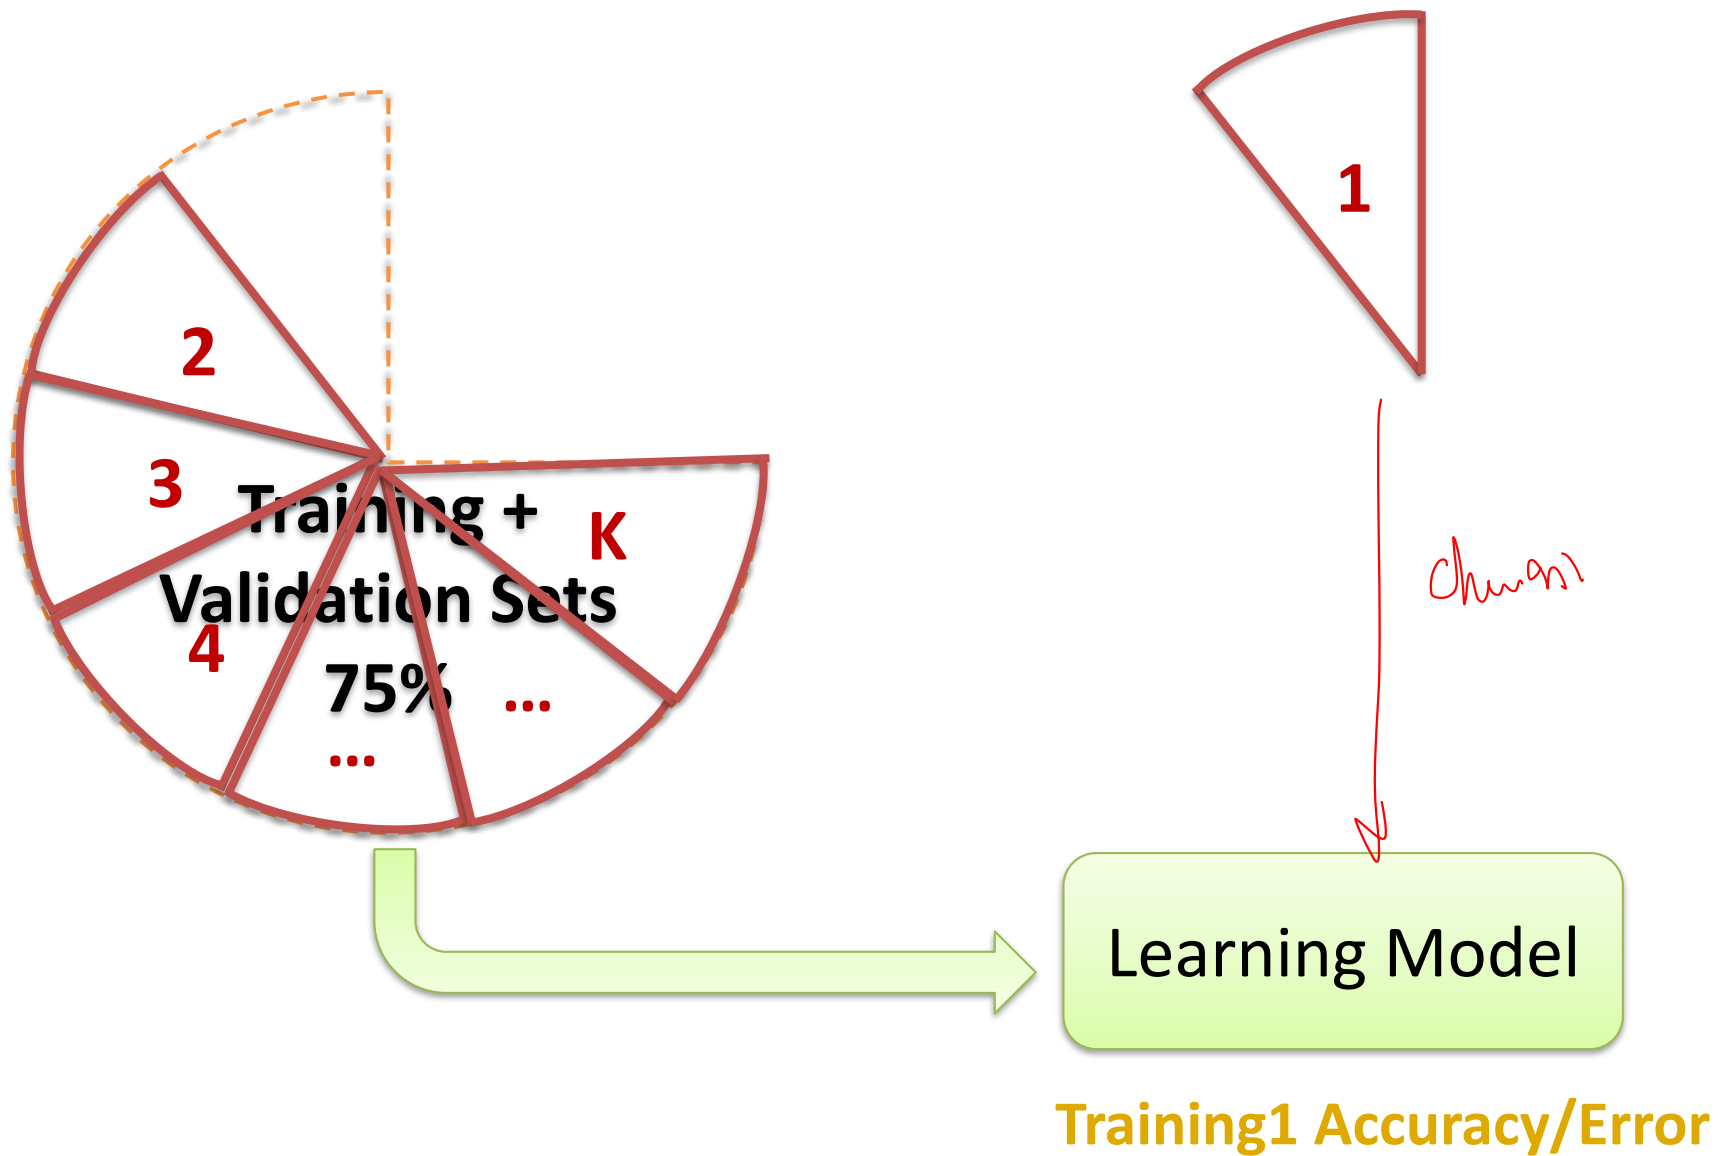
- Split the training data into k equal subsets.
- Then perform k rounds of learning. On each round, $1/k$ of the data is held out as a validation set and the remaining examples are used as training data.
- The average error rate of the k rounds is finally calculated. Popular values for k are 5 and 10. The extreme is $k=n$ (the total number of examples), also known as **leave-one-out cross-validation** or **LOOCV**.



K-Fold Cross Validation





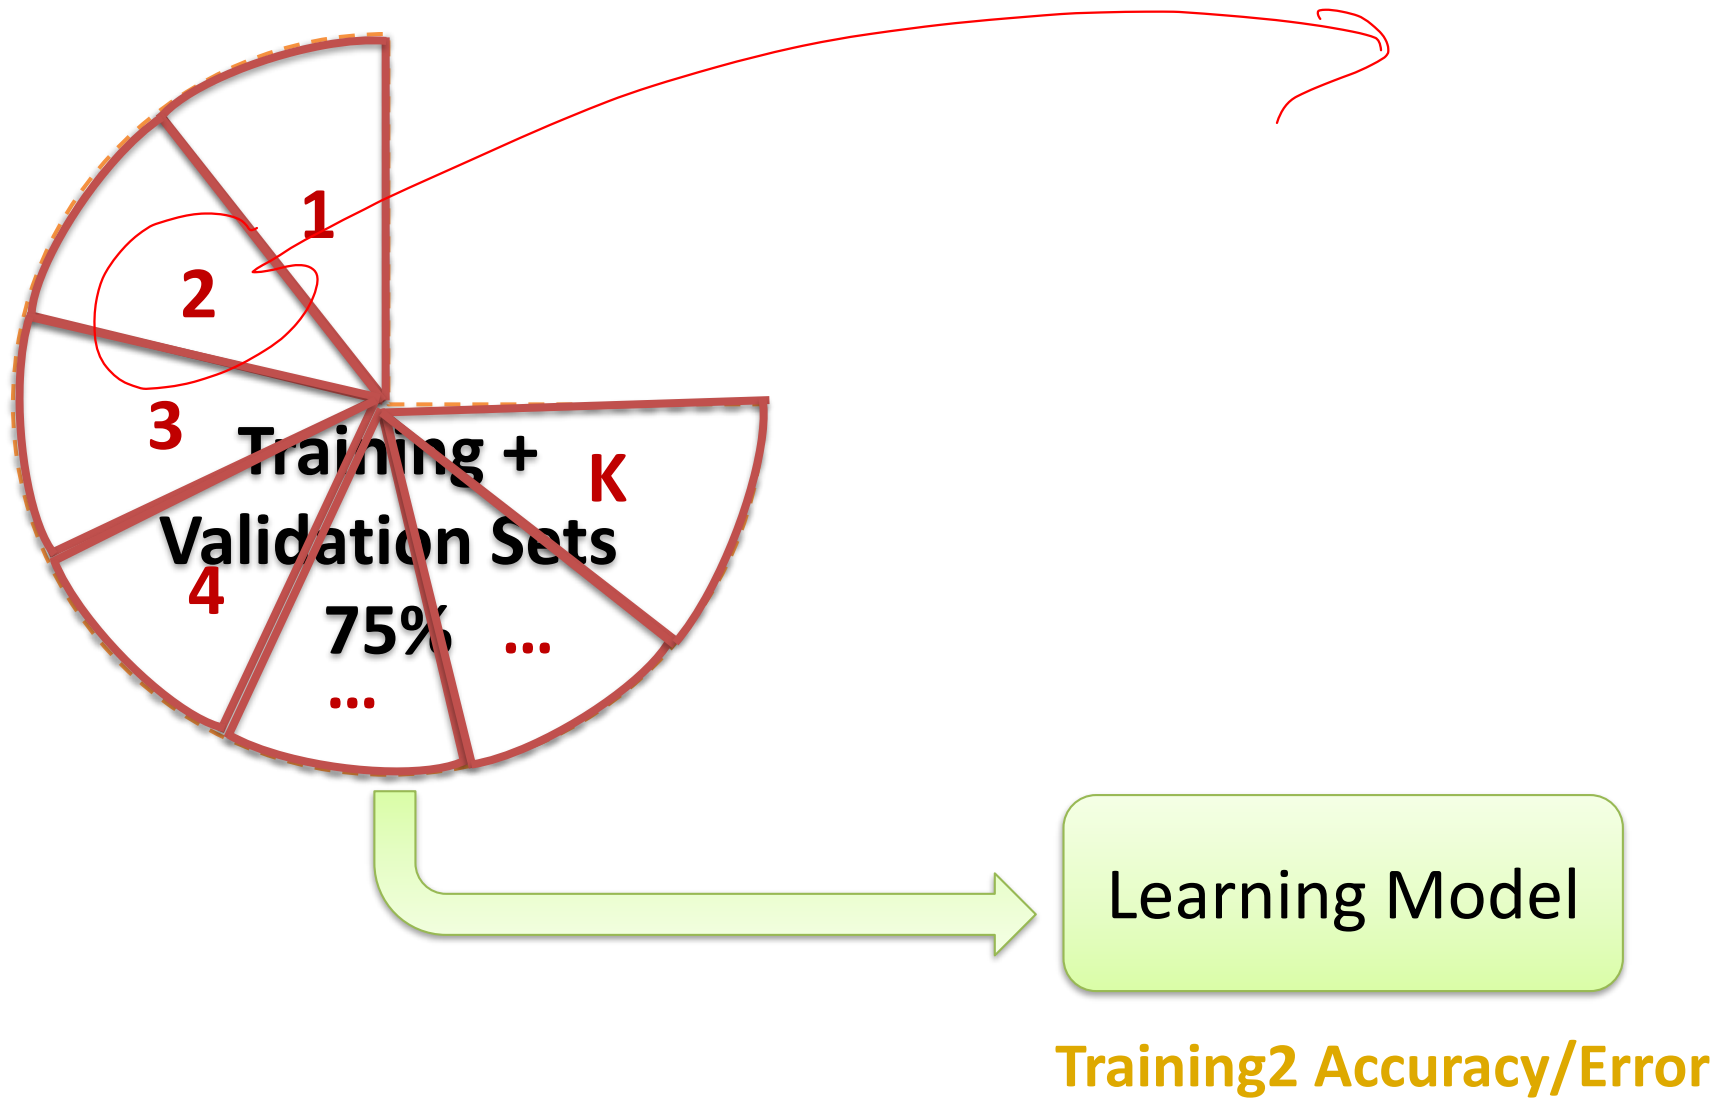
Learning Model







Learning Model


Validation1 Accuracy/Error






Learning Model


Validation2 Accuracy/Error




Learning Model




Learning Model



Learning Model

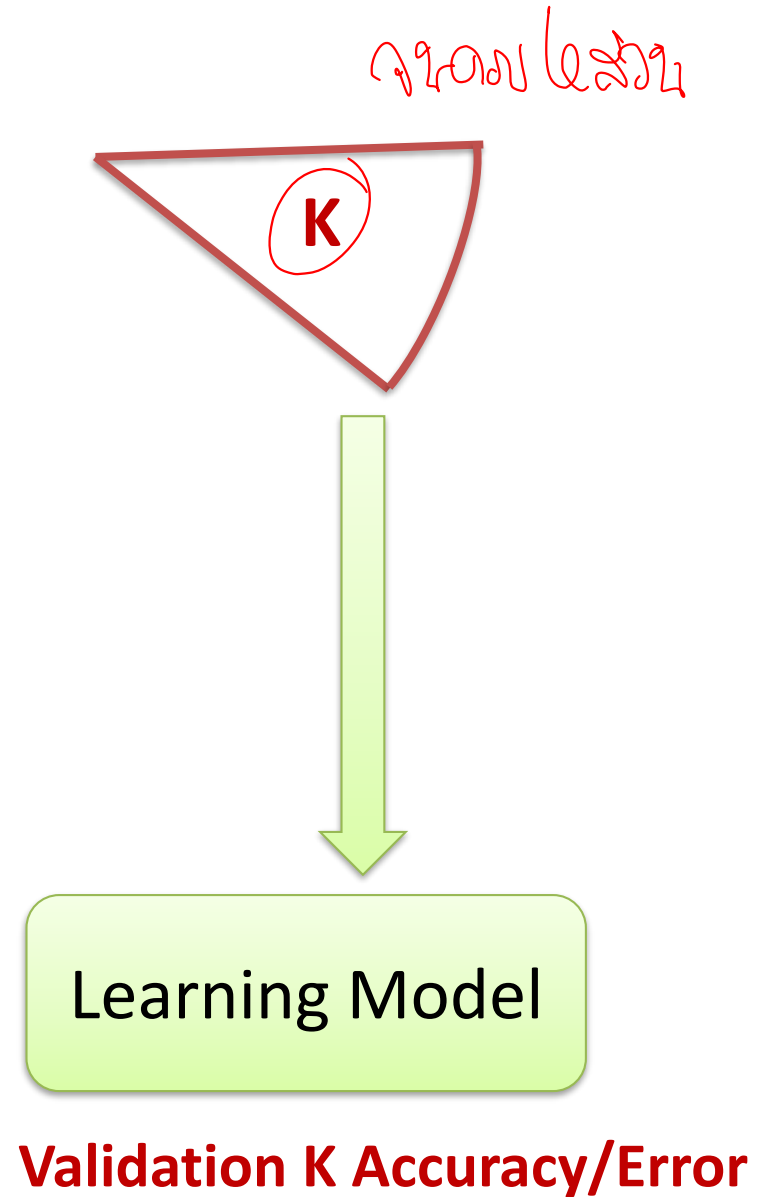




Learning Model



Learning Model


Training K Accuracy/Error

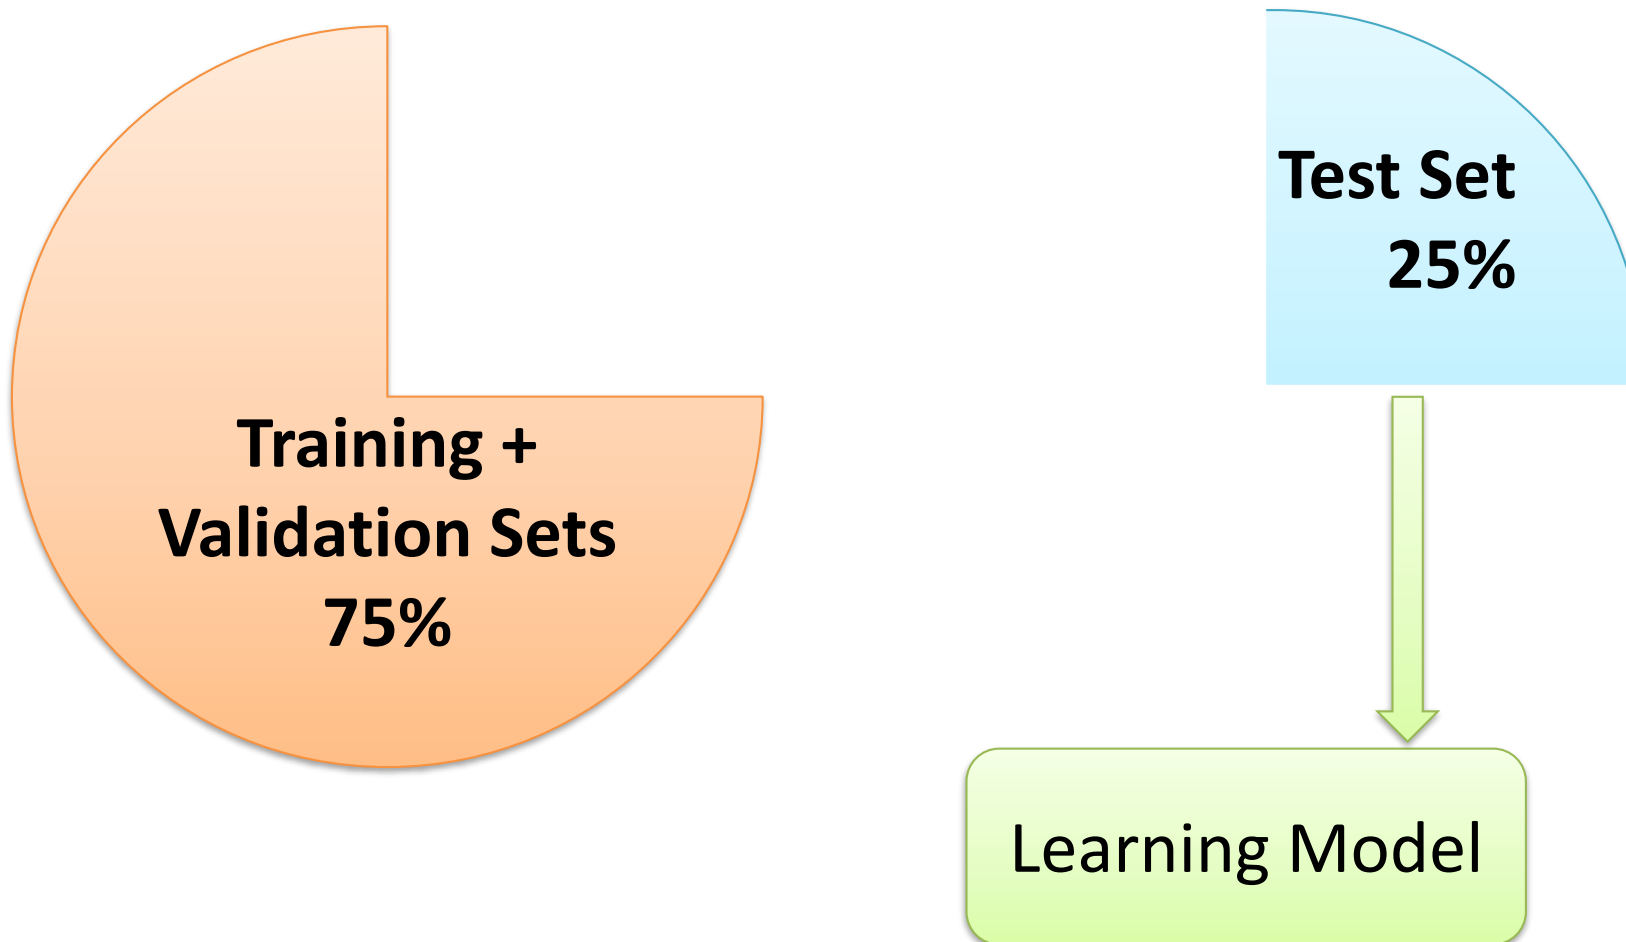




$$\text{Training Accuracy} = (\sum_{i=1}^K \text{Train}_i\text{Acc})/K$$

$$\text{Validation Accuracy} = (\sum_{i=1}^K \text{Valid}_i\text{Acc})/K$$





Note: Percentage of test set can be varied, depends on how many total of instances we have.

Test Accuracy/Error